



UP & RUNNING WITH

MICROSOFT POWER BI DESKTOP



COURSE STRUCTURE

This is a **project-based course** designed for students looking for a practical, hands-on, and highly engaging approach to learning Power BI Desktop for business intelligence

Course resources include:

-  **Downloadable PDF eBook** (200+ pages) containing all course slides, assignments and reference materials
-  **Quizzes and Assignments** to reinforce key concepts and simulate real-world scenarios, with step-by-step solution videos
-  Complete **Bonus Project** to test your abilities and apply the skills developed throughout the course to a brand-new data set

COURSE OUTLINE

1	Introducing Power BI Desktop	<i>Installing Power BI Desktop, exploring the Power BI workflow, comparing Power BI vs. Excel, etc.</i>
2	Connecting & Shaping Data	<i>Connecting to data, shaping & transforming tables, using profiling tools, editing, merging & appending queries, etc.</i>
3	Creating a Data Model	<i>Building relational models, creating table relationships, understanding cardinality and filter flow, etc.</i>
4	Calculating Measures with DAX	<i>Understanding DAX syntax, adding calculated columns and measures, writing common formulas and functions, etc.</i>
5	Visualizing Data with Dashboards	<i>Inserting charts and visuals, customizing formats, editing interactions, applying filters and bookmarks, etc.</i>
6	Optimizing Power BI Performance	<i>Exploring common Power BI optimization tools within the Optimize and External tools menus</i>

COURSE PROJECT

THE SITUATION

You've just been hired as a Business Intelligence Analyst by **AdventureWorks***, a global manufacturing company that produces cycling equipment and accessories

THE BRIEF

The management team needs a way to **track KPIs** (*sales, revenue, profit, returns*), **compare regional performance**, **analyze product-level trends**, and **identify high-value customers**.

All you've been given is a **folder of raw csv files**, which contain information about transactions, returns, products, customers, and sales territories.

THE OBJECTIVE

Use Power BI Desktop to:

- Connect and transform the raw data
- Build a relational data model
- Create calculated columns and measures with DAX
- Design an interactive dashboard to visualize the data



SETTING EXPECTATIONS

1 What you see on your screen **may not always match mine**

- *Power BI Desktop features are updated frequently, with new versions released each month*
- **NOTE:** Power BI is currently only compatible with PC/Windows (not available for Mac)

2 This course is designed to help you build **foundational skills**

- *Our goal is to help you build a deep foundational understanding of the Power BI desktop workflow; some topics may be simplified, and we won't cover some advanced tools (M code, advanced DAX, R/Python visuals, etc.)*

3 This is a **hands-on** and **project-based** learning experience

- *You will get the most value out of this course if you follow along closely with the demos and assignments; we'll be working through the entire BI workflow to create a professional-quality dashboard from scratch*

4 We will not cover **Power BI Service** as part of this course

- *This course focuses on Power BI Desktop specifically; online sharing and collaboration features (app.powerbi.com) require a separate account and are covered in-depth in a separate course*

INTRODUCING POWER BI



MEET POWER BI



In this section we'll **introduce Power BI Desktop**, review the download and installation process, adjust default settings, and explore the Power BI interface and workflow

TOPICS WE'LL COVER:

Introducing Power BI

Power BI vs. Excel

Installation Options

Adjusting Settings

Interface & Workflow

Helpful Resources

GOALS FOR THIS SECTION:

- Download and install Power BI Desktop, and adjust the settings for our course project
- Understand the role that Power BI plays within the broader Microsoft ecosystem
- Explore core components of the Power BI Desktop interface
- Review the business intelligence workflow that we'll follow as we build our course project

MEET POWER BI



Microsoft Power BI is a self-service business intelligence platform, which includes both desktop and web-based applications for connecting, modeling, and visualizing data

Learn more at powerbi.microsoft.com



Figure 1: Magic Quadrant for Analytics and Business Intelligence Platforms



© Gartner, Inc.

WHY POWER BI?



Connect, transform and load millions of rows of data

- *Access data from virtually anywhere (database tables, flat files, web, cloud services, folders, etc.), and create fully automated workflows to extract, transform and load data for analysis*



Build relational models to blend data from multiple sources

- *Create table relationships to analyze holistic performance across an entire relational data model*



Define complex calculations using Data Analysis Expressions (DAX)

- *Enhance datasets and enable advanced analytics with powerful and portable DAX expressions*



Bring data to life with interactive reports and dashboards

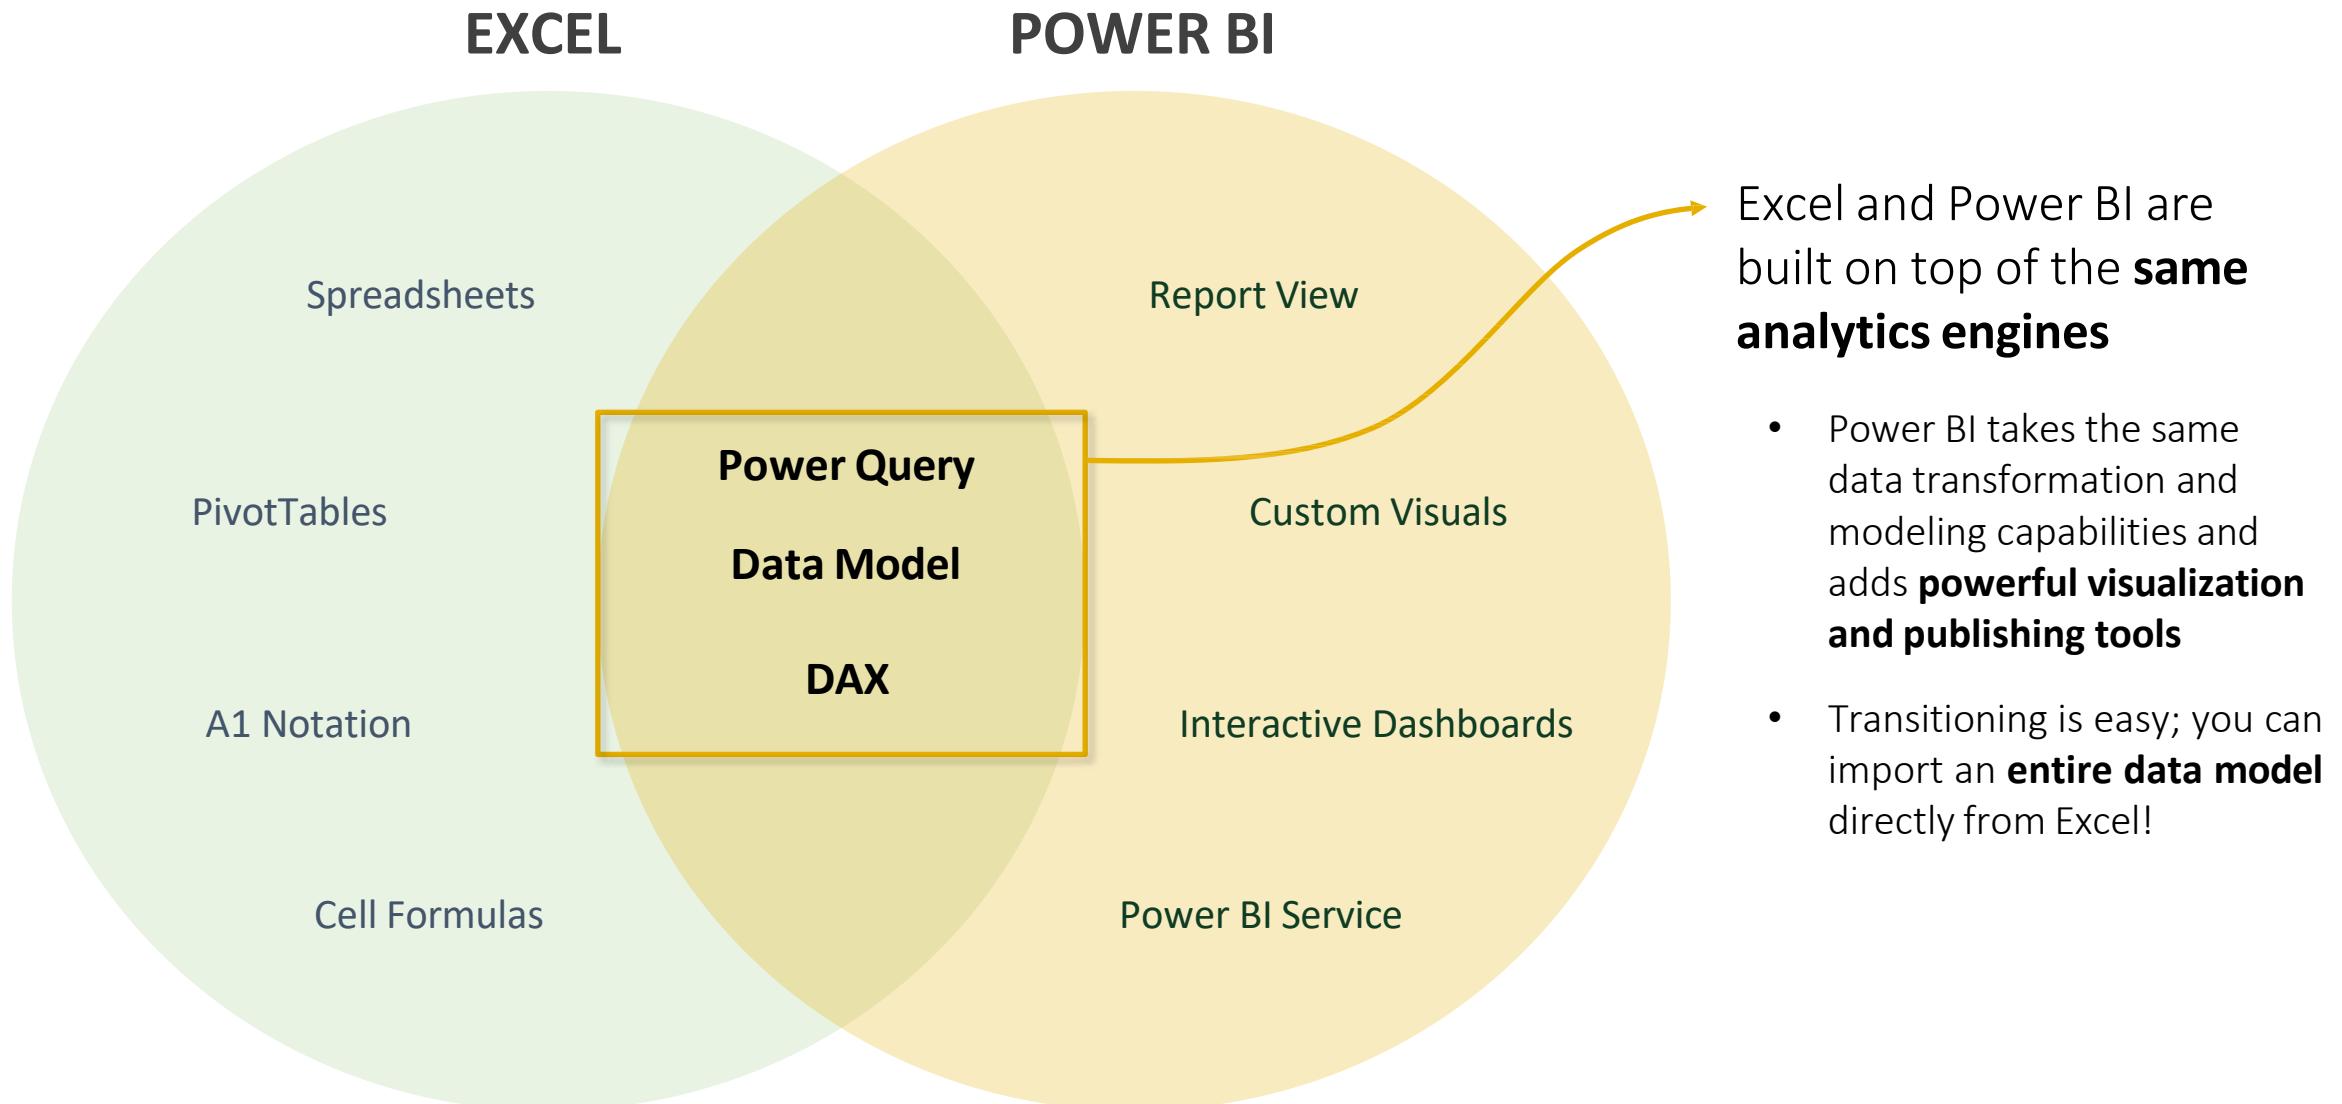
- *Build professional-quality reports and dashboards with best-in-class visualization tools*



Develop a versatile, in-demand skill set

- *Power BI is the industry leader in self-service BI, and the skills you build in this course will be highly transferrable*

EXCEL VS. POWER BI



INSTALLING POWER BI DESKTOP

1) Download from Microsoft store

apps.microsoft.com



- Windows handles **automatic updates**
- Updates only elements that have been changed
- Doesn't require administrator access

2) Download manually from web

powerbi.microsoft.com/downloads



- **No automatic updates** (allows version control)
- Downloads an executable installation file
- Administrator access may be required

3) Install as part of Microsoft 365

microsoft.com/en-us/microsoft-365



- Power BI Desktop is included as part of select enterprise Office/Microsoft 365 subscriptions
- If your company uses a compatible version of Microsoft 365, talk to an admin about getting access to Power BI

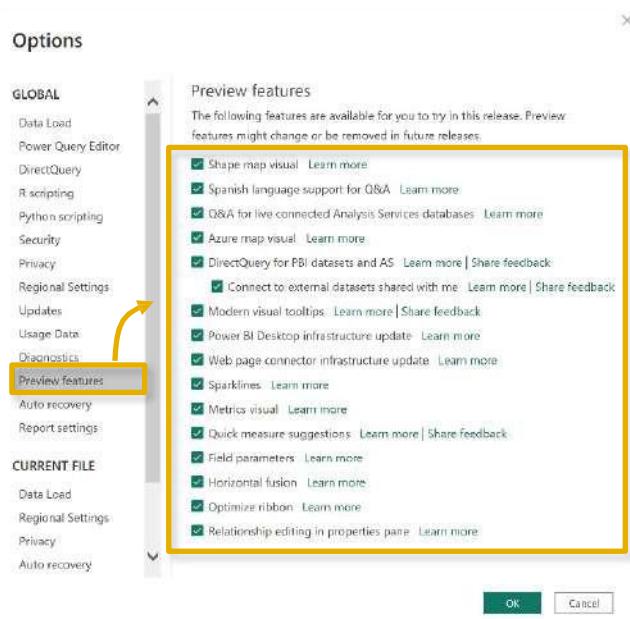


HEY THIS IS IMPORTANT!

You do **NOT** need to register for a Power BI Pro account to access Power BI Desktop

POWER BI SETTINGS

Global > Preview Features

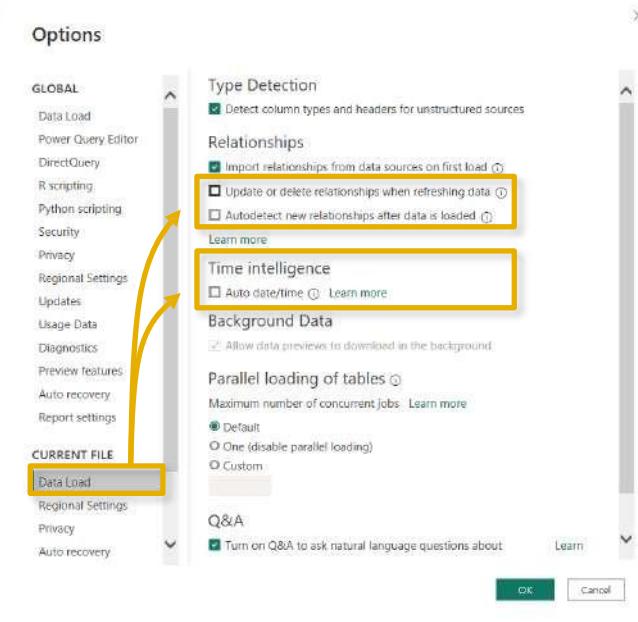


Select all available preview features by default (these change with each monthly release)

HEY THIS IS IMPORTANT!

Options under **CURRENT FILE** need to be adjusted **every time you open a new Power BI workbook** (these settings do not persist across new .pbix files)

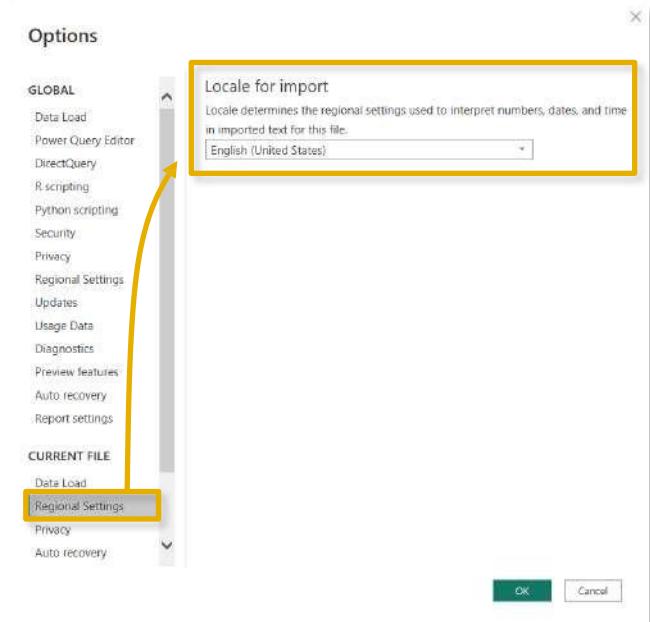
Current File > Data Load



Make sure the following options are **NOT selected**:

- Update or delete relationships when refreshing data*
- Autodetect new relationships after data is loaded*
- Time Intelligence > Auto date/time*

Current File > Regional Settings



Select "**English (United States)**" from the dropdown menu (this will align with the data in course project files)

POWER BI WORKFLOW

Raw data is extracted and transformed in the **Power Query editor**, then loaded to the Power BI “front-end”

The screenshot shows the Power Query Editor interface. The ribbon at the top has tabs for Home, Transform, Add Column, View, Tools, Help, and a dropdown for Data Sources. The main area displays a table with 11 columns and 293 rows, with the first few rows visible. The table includes columns like Product Key, Product Subcategory Key, Product SKU, Product Name, Model Name, and Product Description. A large orange arrow points from this screen to the Power BI desktop interface.



Power Query Editor

Power BI “Back-End”

The screenshot shows the Power BI desktop interface with a dashboard titled "ADVENTUREWORKS". The dashboard features several visualizations: a line chart for "Weekly Revenue" showing revenue trends from Jan 2020 to Jan 2022; a bar chart for "Top 10 Products" listing items like Water Bottle - 30 oz, Patch Kit/8 Patches, and Mountain Bike Tube; and a card for "Most Ordered Product Type" showing Tires and Tubes. The ribbon at the top includes tabs for File, Home, Insert, Modeling, View, Optimize, Help, and External tools, along with various icons for data sources, queries, and calculations.



Model View



Data View



Report View

Power BI “Front-End”

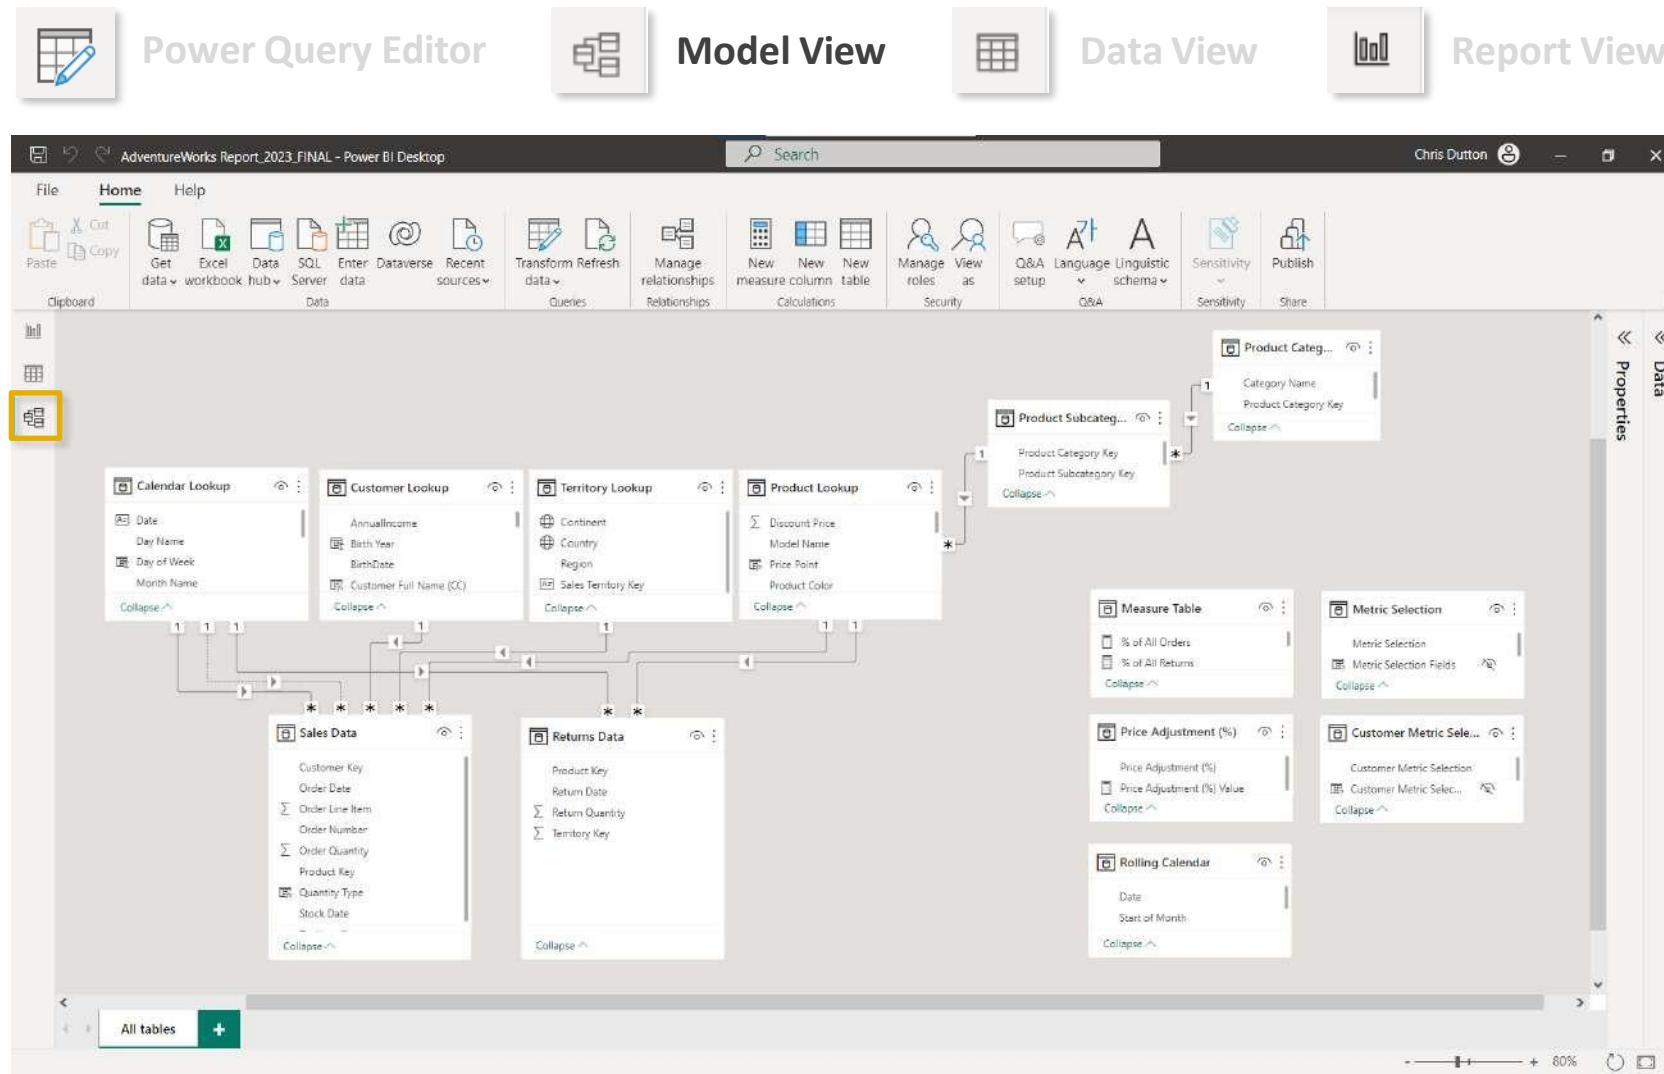
POWER BI WORKFLOW

The screenshot shows the Microsoft Power Query Editor interface. At the top, there are four tabs: Power Query Editor (selected), Model View, Data View, and Report View. The main area displays a table with 29 rows and 11 columns. The columns are labeled: Product Key, Product Subcategory Key, Product SKU, Product Name, Model Name, and Product Description. The table contains various product details such as Patch Kit/8 Patches, Road Tire Tube, Water Bottle - 30 oz., etc. On the left side, a sidebar titled 'Queries [14]' lists several queries including 'Transform File from Sales Data...', 'Helper Queries [3]', 'Other Queries [10]', and 'Product Lookup'. The 'Product Lookup' query is currently selected. The bottom of the screen shows a status bar with the text '11 COLUMNS, 293 ROWS Column profiling based on top 1000 rows.'

1

Data is loaded & transformed in the **Power Query Editor**

POWER BI WORKFLOW



1 Data is loaded & transformed in the **Power Query Editor**

2 Data models are configured in the **Model View**

POWER BI WORKFLOW

Power Query Editor

Model View

Data View

Report View

AdventureWorks Report_2023_FINAL - Power BI Desktop

File Home Help Table tools

Name Customer Lookup

Structure

Table tools

Mark as date table Calendars

Manage relationships Relationships

New measure Quick New column

New table Calculations

Search

Chris Dutton

Prefix Gender Occupation CustomerKey FirstName LastName BirthDate MaritalStatus EmailAddress AnnualIncome TotalChildren Education

Mr. M Professional 11206 Blake Flores Friday, September 24, 1948 M blake60@adventure-works.com \$60,000 2 Partie

Mr. M Professional 12093 Levi Chandra Monday, November 08, 1948 M levi1@adventure-works.com \$60,000 2 Partie

Mr. M Professional 12108 James Walker Thursday, February 23, 1950 M james96@adventure-works.com \$60,000 2 Partie

Mr. M Professional 12117 Cameron Yang Friday, August 03, 1951 M cameron23@adventure-works.com \$60,000 2 Partie

Mr. M Professional 12233 Richard Coleman Sunday, December 05, 1948 M richard61@adventure-works.com \$60,000 2 Partie

Mr. M Professional 12235 Jonathan Robinson Wednesday, September 22, 1948 M jonathan72@adventure-works.com \$60,000 2 Partie

Mr. M Professional 12236 Robert Wang Tuesday, May 25, 1948 M robert36@adventure-works.com \$60,000 2 Partie

Mr. M Professional 13370 Blake Ross Monday, March 22, 1948 M blake51@adventure-works.com \$60,000 2 Partie

Mr. M Professional 13372 Isaac Edwards Sunday, November 28, 1948 M isaac24@adventure-works.com \$60,000 2 Partie

Mr. M Professional 13375 Jason Kumar Sunday, April 11, 1948 M jason26@adventure-works.com \$60,000 2 Partie

Mr. M Professional 13376 Jerry Rai Friday, September 03, 1948 M jerry19@adventure-works.com \$60,000 2 Partie

Mr. M Professional 13397 Adam Mitchell Sunday, April 03, 1951 M adam41@adventure-works.com \$60,000 2 Partie

Mr. M Professional 14405 Dylan Walker Thursday, June 01, 1950 M dylan53@adventure-works.com \$60,000 2 Partie

Mr. M Professional 14407 Thomas Bryant Sunday, June 04, 1950 M thomas20@adventure-works.com \$60,000 2 Partie

Mr. M Professional 14415 William Davis Friday, November 16, 1951 M william21@adventure-works.com \$60,000 2 Partie

Mr. M Professional 15325 Elijah Hayes Thursday, May 20, 1948 M elijah24@adventure-works.com \$60,000 2 Partie

Mr. M Professional 15331 Jacob Taylor Friday, August 26, 1949 M jacob2@adventure-works.com \$60,000 2 Partie

Mr. M Professional 15322 Jason Sharma Friday, April 04, 1949 M jason27@adventure-works.com \$60,000 2 Partie

Mr. M Professional 15336 Marco Garcia Thursday, September 21, 1950 M marco15@adventure-works.com \$60,000 2 Partie

Mr. M Professional 15329 Noah Zhang Saturday, July 01, 1950 M noeh21@adventure-works.com \$60,000 2 Partie

Mr. M Professional 15343 Carson Barnes Friday, September 07, 1951 M carson2@adventure-works.com \$60,000 2 Partie

Mr. M Professional 15902 Robert Diaz Tuesday, April 25, 1950 M robert33@adventure-works.com \$60,000 2 Partie

Mr. M Professional 15908 David Wilson Thursday, September 20, 1951 M david65@adventure-works.com \$60,000 2 Partie

Mr. M Professional 16806 Luis Zhao Friday, August 04, 1950 M louis4@adventure-works.com \$60,000 2 Partie

Mr. M Professional 16811 Luis Zhang Thursday, April 19, 1951 M luis23@adventure-works.com \$60,000 2 Partie

Mr. M Professional 16813 Carson Diaz Saturday, April 14, 1951 M carson21@adventure-works.com \$60,000 2 Partie

Mr. M Professional 17144 Luis Griffin Friday, February 20, 1948 M luis20@adventure-works.com \$60,000 2 Partie

Mr. M Professional 17749 Steven Richardson Friday, April 22, 1949 M steven19@adventure-works.com \$60,000 2 Partie

Mr. M Professional 17155 Samuel Lewis Thursday, April 21, 1949 M samuel71@adventure-works.com \$60,000 2 Partie

Table: Customer Lookup (18,148 rows)

1 Data is loaded & transformed in the **Power Query Editor**

2 Data models are configured in the **Model View**

3 Table features & calculations are added in the **Data View**

POWER BI WORKFLOW

The screenshot shows the Power BI Desktop interface in the Report View. The main area displays a dashboard with several key performance indicators (KPIs) and visualizations. At the top, there are four large boxes showing \$24.9M Revenue, \$10.5M Profit, 25.2K Orders, and 2.2% Return Rate. Below these are three charts: a line chart for 'Weekly Revenue' over time, a bar chart for 'Orders by Category' (Accessories, Bikes, Clothing), and a table for 'Top 10 Products'. There are also smaller cards for 'Monthly Revenue', 'Monthly Orders', and 'Monthly Returns'. The ribbon at the top has tabs for Power Query Editor, Model View, Data View, and Report View, with Report View selected. The left sidebar shows navigation options like Exec Dashboard, Map, Product Detail, Customer Detail, etc.

1 Data is loaded & transformed in the **Power Query Editor**

2 Data models are configured in the **Model View**

3 Table features & calculations are added in the **Data View**

4 Visuals & reports are designed in the **Report View**

HELPFUL RESOURCES



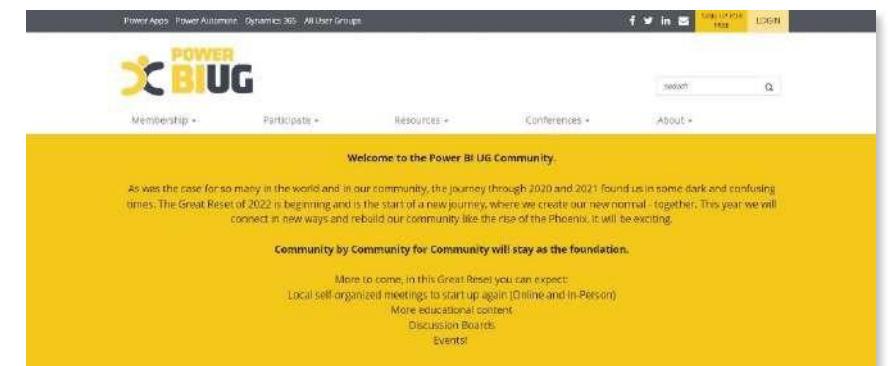
The **Help** tab includes documentation, training videos, sample files, templates, and links to support blogs and communities



The **Microsoft Power BI blog** (powerbi.microsoft.com/blog) publishes monthly summaries to showcase new features



The **Microsoft Power BI YouTube Channel** publishes demos, feature summaries, and advanced tutorials (check out “**Guy in a Cube**” too!)



Power BI User Groups (Power BIUG) are communities of users, which include both local meet-ups and helpful online forums (pbiusergroup.com)

MONTHLY UPDATES

Power BI is updated monthly, so you may notice ongoing changes to settings, options, tools, etc.
Reference the links below to stay up-to-date on product updates and new feature releases:



Power BI Desktop

<https://docs.microsoft.com/en-us/power-bi/fundamentals/desktop-latest-update>



Power BI Service

<https://docs.microsoft.com/en-us/power-bi/fundamentals/service-whats-new>



Power Platform

<https://learn.microsoft.com/en-us/dynamics365/release-plans/>

CONNECTING & SHAPING DATA



CONNECTING & SHAPING DATA



In this section we'll connect to source files and cover some of the most common techniques for **extracting, cleaning, and shaping data** to prepare it for modeling and analysis

TOPICS WE'LL COVER:

Intro to Power Query

Data Connectors

The Query Editor

Connection Modes

Data QA & Profiling

Table Transformations

Calendar Tools

Combining Queries

GOALS FOR THIS SECTION:

- Explore Power BI's query editor and understand the role that Power Query plays in the larger BI workflow
- Introduce different types of connectors and connectivity modes available for getting data into Power BI
- Review tools for checking data quality and key profiling metrics like column distribution, empty values, errors and outliers
- Transform tables using text, numerical and date/time tools, pivot and group records, and create new conditional columns
- Practice combining, modifying and refreshing queries



FRONT-END VS. BACK-END

Power BI Desktop essentially has two distinct environments: a **front-end** and a **back-end**

- The **front-end** includes the **Data, Model & Report** views, where most of the modeling, analysis and visualization takes place
- The **back-end** includes the **Power Query Editor**, where raw data is extracted, transformed, and loaded to the front-end (ETL)

BACK-END

- **Connect & extract** data using pre-built connectors
- **Profile & QA** the data to explore, clean and prepare it for modeling and analysis
- **Transform & shape** tables to add new features, modify values, group records, or sort and filter columns
- **Merge or append** queries to join and combine them prior to loading to the front-end
- Perform **advanced transformations** using custom M code (out of scope for this course)

FRONT-END

- **Build data models** by creating table relationships between primary and foreign keys
- **Add calculated measures & columns** using Data Analysis Expressions (DAX)
- **Design reports** to visualize the data and create interactive, dynamic dashboards
- **Publish & share** your Power BI workbooks using Power BI Service (cloud application)



TYPES OF DATA CONNECTORS

The screenshot shows the Microsoft Power BI ribbon interface. The 'Home' tab is selected, indicated by a green underline. The 'Insert' tab is also visible. A yellow box highlights the 'Get data' button in the 'Insert' tab's dropdown menu. A yellow arrow points from this highlighted button to the 'Get Data' section of the main content area. The 'Get Data' section has a search bar and a list of categories: All, File, Database, Power Platform, Azure, Online Services, and Other. To the right of this list is a large bracket that spans across the entire list of data sources.

All
Excel Workbook
Text/CSV
XML
JSON
Folder
PDF
Parquet
SharePoint folder
SQL Server database
Access database
SQL Server Analysis Services database
Oracle database
IBM Db2 database
IBM Informix database (Beta)
IBM Netezza
MySQL database

Power BI can connect to virtually **any** type of source data, including (*but not limited to*):

- **Flat files & Folders** (csv, text, xlsx, etc.)
- **Databases** (SQL, Access, Oracle, IBM, etc.)
- **Power Platform** (Datasets, Datamarts, Dataflows, Dataverse, etc.)
- **Azure** (Azure SQL, Analysis Services, Databricks, etc.)
- **Online Services** (SharePoint, GitHub, Dynamics 365, Google Analytics, Salesforce, Power BI Service, etc.)
- **Other** (Web feeds, R scripts, Spark, Hadoop, etc.)



POWER QUERY EDITOR

File Home Insert Modeling View Optimize Help External Tools

Get data workbook hub Data Server Data Enter Data Datasource Recent sources Data

Transform data Queries

File Home Transform Add Column View Tools Help

Column From Custom Invoke Custom Examples Column Function General

From Text From Number From Date & Time AI Insights

Conditional Column Index Column Duplicate Column Format Add Extract ABC Parse Statistics Standard Scientific Information Date Time Duration Text Analytics Vision Azure Machine Learning

Queries [13]

Customer Key Prefix First Name Last Name Birth Date Marital Status Gen

	Customer Key	Prefix	First Name	Last Name	Birth Date	Marital Status	Gen
1	11000	Mr.	Jon	Yang	4/8/1966	M	M
2	11001	Mr.	Eugene	Huang	5/14/1965	S	M
3	11002	Mr.	Ruben	Torres	8/12/1965	M	M
4	11003	Ms.	Christy	Zhu	2/15/1968	S	F
5	11004	Mrs.	Elizabeth	Johnson	8/8/1968	S	F
6	11005	Mr.	Julio	Rulz	8/5/1965	S	M
7	11007	Mr.	Marco	Mehta	5/9/1964	M	M
8	11008	Mrs.	Robin	Verhoff	7/7/1964	S	F
9	11009	Mr.	Shannon	Carlson	4/1/1964	S	M
10	11010	Ms.	Jacquelyn	Suarez	2/6/1964	S	F
11	11011	Mr.	Curtis	Lu	11/4/1963	M	M
12	11012	Mrs.	Lauren	Walker	1/18/1968	M	F
13	11013	Mr.	Ian	Jenkins	8/6/1968	M	M
14	11014	Mrs.	Sydney	Bennett	5/9/1968	S	F
15	11015	Ms.	Chloe	Young	2/27/1979	S	F
16	11016	Mr.	Wyatt	Hill	4/28/1979	M	M
17	11017	Mrs.	Shannon	Wang	6/26/1944	S	F
18	11018	Mr.	Clarence	Rai	10/9/1944	S	M
19	11019	Mr.	Luke	Lal	3/7/1978	S	M
20	11020	Mr.	Jordan	King	9/20/1978	S	M
21	11021	Ms.	Destiny	Wilson	9/3/1978	S	F
22	11022	Mr.	Ethan	Zhang	10/12/1978	M	M
23	11023	Mr.	Seth	Edwards	10/11/1978	M	M
24	11024	Mr.	Russell	Xie	9/17/1978	M	M
25	11025		Alejandro	Beck	12/23/1945	M	NA
26	11026	Mr.	Harold	Sai	4/3/1946	S	M
27	11027	Mr.	Jessie	Zhao	12/7/1946	M	M
28	11028	Mrs.	Jill	Jimenez	4/11/1946	M	F
29							

17 COLUMNS, 999+ ROWS Column profiling based on top 1000 rows

PREVIEW DOWNLOADED AT 1:52

Queries [13]

Customer Key Prefix First Name Last Name Birth Date Marital Status Gen

Query Settings

PROPERTIES

- Name: Customer Lookup
- All Properties

APPLIED STEPS

- Source
- Promoted Headers
- Changed Type
- Column Rename
- Capitalized Each Word
- Inserted FullName Column
- Inserted Text Before Delimiter
- Renamed Columns
- Inserted Text Between Delimiter
- Renamed Columns1
- Replaced Value
- Capitalized Each Word1
- Inserted Year
- Renamed Columns2
- Added Conditional Column
- Filtered Rows
- Removed Columns

Queries Pane
(list of all queries)

Query Editing Tools (Table transformations, calculated columns, etc.)

Formula Bar (this is "M" code)

Table Name & Properties

Applied Steps
(like a macro)

Table Preview



QUERY EDITING TOOLS

The **HOME** tab includes **general settings** and **common table transformation tools**

The screenshot shows the Power BI ribbon with the 'Home' tab selected. The 'File' tab is highlighted in blue. The 'Home' tab contains icons for Close & Apply, New Source, Recent Sources, Enter Data, Data source settings, Manage Parameters, Refresh Preview, Advanced Editor, Properties, Choose Columns, Remove Columns, Keep Rows, Remove Rows, Sort, Split Column, Group By, Data Type (Whole Number), Use First Row as Headers, Merge Queries, Append Queries, Combine Files, and Combine. Below the ribbon, there are buttons for Close, New Query, Data Sources, Parameters, Query, Manage Columns, Reduce Rows, Transform, and AI Insights.

The **TRANSFORM** tab includes tools to **modify existing columns** (splitting/grouping, transposing, extracting text, etc.)

The screenshot shows the Power BI ribbon with the 'Transform' tab selected. The 'File' tab is highlighted in blue. The 'Transform' tab contains icons for Group By, Use First Row as Headers, Count Rows, Transpose, Reverse Rows, Detect Data Type, Rename, Pivot Column, Convert to List, Unpivot Columns, Fill, Move, Format, Statistics, Standard, Scientific, Trigonometry, Rounding, Parse, and Information. Below the ribbon, there are buttons for Table, Any Column, Text Column, Number Column, Date & Time Column, and Scripts. On the far right, there are buttons for Run R script and Run Python script.

The **ADD COLUMN** tools **create new columns** (based on conditional rules, text operations, calculations, dates, etc.)

The screenshot shows the Power BI ribbon with the 'Add Column' tab selected. The 'File' tab is highlighted in blue. The 'Add Column' tab contains icons for Column From Examples, Custom Column, Invoke Custom Function, Conditional Column, Index Column, Duplicate Column, Format, Merge Columns, Statistics, Standard, Scientific, Trigonometry, Rounding, Parse, Date, Time, Duration, Text Analytics, Vision, Azure Machine Learning, and AI Insights. Below the ribbon, there are buttons for General, From Text, From Number, From Date & Time, and AI Insights.



BASIC TABLE TRANSFORMATIONS

The screenshot shows the Power BI desktop ribbon with several transformation tools highlighted by yellow arrows:

- Sort values (A-Z, Low-High, etc.)**: Points to the Sort button in the Transform ribbon tab.
- Change data type (date, \$, %, text, etc.)**: Points to the Data Type dropdown in the Transform ribbon tab.
- Promote headers**: Points to the Use First Row as Headers dropdown in the Transform ribbon tab.
- Choose or remove columns**: Points to the Manage Columns group in the ribbon, which includes the Choose Columns and Remove Columns buttons.
- Keep or remove rows**: Points to the Remove Rows group in the ribbon, which includes the Keep Rows, Remove Rows, and Reduce Rows buttons.
- Duplicate, move or rename columns**: Points to the context menu for a column header, which includes options like Copy, Remove, Remove Other Columns, Duplicate Column, Add Column From Examples..., Remove Duplicates, Remove Errors, Change Type, Transform, Replace Values..., Replace Errors..., Group By..., Fill, Unpivot Columns, Unpivot Other Columns, Unpivot Only Selected Columns, Rename..., Move, Drill Down, and Add as New Query.
- Tip: use the "Remove Other Columns" option if you always want a specific set**: A tip for using the Remove Other Columns option in the Choose or remove columns section.
- Tip: use "Remove Duplicates" to create new lookup tables from scratch**: A tip for using the Remove Duplicates option in the Keep or remove rows section.

ASSIGNMENT: TABLE TRANSFORMATIONS





NEW MESSAGE

From: **Ethan T. Langer (Analytics Manager)**
Subject: **Welcome aboard!**

Hello, and welcome to the team!
We're excited that you'll be helping us develop our new internal reports in Power BI. Looks like you've already gotten started, but we have some new data to add to the model.
Could you please create two new queries to connect to the **Product Category Lookup** and **Product Subcategory Lookup** files attached, and help with a few modifications to the product table?
Thanks!
-ETL

[Product Category Lookup](#)
[Product Subcategory Lookup](#)

[Reply](#) [Forward](#)

Key Objectives

1. Create queries to connect to the two new .csv files
2. Name your queries **Product Category Lookup** and **Product Subcategory Lookup**
3. Confirm that column headers have been promoted and that all data types are correct
4. Add a new column to extract all characters before the dash ("–") in the **Product SKU** column, and name it "**SKU Type**"
5. Update the **SKU Type** calculation above to return all characters before *second* dash, instead of the first
6. Replace zeros (**0**) in the **Product Style** column with "**NA**"
7. Close and load to your data model



SOLUTION: TABLE TRANSFORMATIONS

 NEW MESSAGE

From: **Ethan T. Langer (Analytics Manager)**

Subject: **Welcome aboard!**

Hello, and welcome to the team!

We're excited that you'll be helping us develop our new internal reports in Power BI. Looks like you've already gotten started, but we have some new data to add to the model.

Could you please create two new queries to connect to the **Product Category Lookup** and **Product Subcategory Lookup** files attached, and help with a few modifications to the product table?

Thanks!

-ETL

 [Product Category Lookup](#)
[Product Subcategory Lookup](#)

Solution Preview

File Home Transform Add Column View Tools Help

Queries [4] = Table.TransformColumnTypes

1	2	3	4
ProductCategoryKey	CategoryName	Bikes	Components
1	2	2 Components	3 Clothing
2	3	3 Clothing	4 Accessories
3	4	4 Accessories	
4			

Query Settings

PROPERTIES

Name: Product Category Lookup

All Properties

APPLIED STEPS

Source
Promoted Headers
Changed Type

= Table.ReplaceValue#"Renamed Columns","0","NA",Replacer.ReplaceText,

ProductStyle	ProductCost	ProductPrice	SKU Type
	13.09	34.99	HL-U509
	12.03	33.64	SO-B909
	3.40	9.50	SO-B909
	3.40	9.50	SO-B909
	12.03	33.64	HL-U509
	5.71	8.64	CA-1098
	31.72	48.07	LI-0192
	747.97	1,263.46	FR-R92R

Query Settings

PROPERTIES

Name: Product Lookup

All Properties

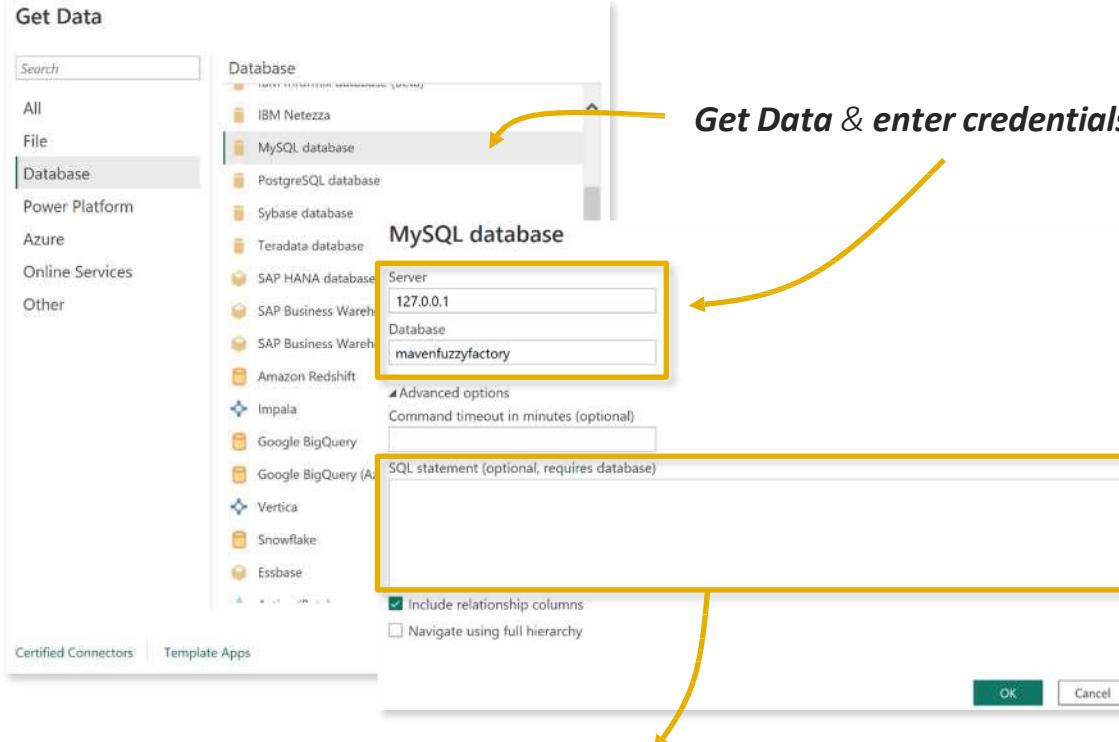
APPLIED STEPS

Source
Promoted Headers
Changed Type
Changed to Currency
Removed Columns
Sorted Rows
Inserted Text Before Delimiter
Renamed Columns
Replaced Value



CONNECTING TO A DATABASE

Power Query can connect to data from various **database sources** including SQL Server, MS Access, MySQL, PostgreSQL, Oracle, SAP, and more



Get Data & enter credentials

The screenshot shows the 'Power Query - Choose data' interface. On the left, a list of tables is shown with several checked. A yellow arrow points to this list with the text 'Select tables & transform'. On the right, a preview of the data is displayed in a table format. The table has columns: website_session_id, created_at, user_id, is_repeat_session, utm_source, and utm_campaign. The data shows 19 rows of session data. At the bottom right are 'Cancel' and 'Transform data' buttons.

website_session_id	created_at	user_id	is_repeat_session	utm_source	utm_campaign
1	3/19/2012, 8:04:16 AM	1	0	gsearch	nonbrand
2	3/19/2012, 8:16:49 AM	2	0	gsearch	nonbrand
3	3/19/2012, 8:26:55 AM	3	0	gsearch	nonbrand
4	3/19/2012, 8:37:33 AM	4	0	gsearch	nonbrand
5	3/19/2012, 9:05:55 AM	5	0	gsearch	nonbrand
6	3/19/2012, 9:05:46 AM	6	0	gsearch	nonbrand
7	3/19/2012, 9:06:27 AM	7	0	gsearch	nonbrand
8	3/19/2012, 9:17:17 AM	8	0	gsearch	nonbrand
9	3/19/2012, 9:27:56 AM	9	0	gsearch	nonbrand
10	3/19/2012, 9:35:37 AM	10	0	gsearch	nonbrand
11	3/19/2012, 9:37:42 AM	11	0	gsearch	nonbrand
12	3/19/2012, 9:39:57 AM	12	0	gsearch	nonbrand
13	3/19/2012, 9:45:29 AM	13	0	gsearch	nonbrand
14	3/19/2012, 9:45:47 AM	14	0	gsearch	nonbrand
15	3/19/2012, 9:57:14 AM	15	0	gsearch	nonbrand
16	3/19/2012, 10:01:35 ...	16	0	gsearch	nonbrand
17	3/19/2012, 10:14:34 ...	17	0	gsearch	nonbrand
18	3/19/2012, 10:16:57 ...	18	0	gsearch	nonbrand
19	3/19/2012, 10:17:38 ...	19	0	gsearch	nonbrand

Write custom or advanced queries with SQL statements (optional)



EXTRACTING DATA FROM THE WEB

Power Query includes a native **Web connector** for importing web-hosted files (csv, xlsx, etc.) or scraping URLs for anything that Power Query can identify as a structured table

List of asset management firms

From Wikipedia, the free encyclopedia

"Asset management company" redirects here. The term may also refer to a [bad bank](#).

An **asset management company** (AMC) is an asset management / investment management company/firm that invests the pooled funds of retail investors in securities in line with the stated investment objectives. For a fee, the company/firm provides more diversification, liquidity, and professional management consulting service than is normally available to individual investors. The diversification of portfolio is done by investing in such securities which are inversely correlated to each other. Money is collected from investors by way of floating various collective investment schemes, e.g. mutual fund schemes. In general, an AMC is a company that is engaged primarily in the business of investing in, and managing, portfolios of securities. A study by consulting firm Casey Quirk, which is owned by Deloitte, found that asset management firms ended 2020 with record highs in both revenue and assets under management.^[1]

Largest companies [edit]

The following is a list of the top 20 asset managers in the world (as of 2022), ranked by total assets under management (AUM)^[2]

Rank	Firm/company	Country	AUM (billion USD)
1	BlackRock	United States	9,570
2	Vanguard Group	United States	8,100
3	Fidelity Investments	United States	4,283
4	UBS	Switzerland	4,380
5	State Street Global Advisors	United States	4,020
6	Morgan Stanley	United States	3,230
7	JPMorgan Chase	United States	2,960
8	Crédit Agricole	France	2,875
9	Allianz	Germany	2,760
10	Capital Group	United States	2,700
11	Goldman Sachs	United States	2,394
12	BNY Mellon	United States	2,266
13	Amundi	France	2,251
14	PIMCO	United States	2,000
15	Legal & General	United Kingdom	1,866
16	Edward Jones Investments	United States	1,700
17	PGIM	United States	1,620
18	Deutsche Bank	Germany	1,615
19	Bank of America	United States	1,571
20	Invesco	United States	1,556

https://en.wikipedia.org/wiki/List_of_asset_management_firms

Navigator

Display Options

- HTML Tables [8]
 - Largest companies[edit]
 - Table 1
 - Table 2
 - Table 3
 - Table 4
 - Table 5
 - Table 6
 - Table 7
- Suggested Tables [4]
 - Table 8
 - Table 9
 - Table 10
 - Table 11
- Text [2]
 - HTML Code
 - Displayed Text

Table View Web View

Largest companies[edit]

Rank	Firm/company	Country	AUM (billion USD)
1	BlackRock	United States	10010
2	Charles Schwab	United States	8140
3	Vanguard Group	United States	8100
4	UBS	Switzerland	4380
5	Fidelity Investments	United States	4283
6	State Street Global Advisors	United States	4020
7	Morgan Stanley	United States	3230
8	JPMorgan Chase	United States	2960
9	Allianz	Germany	2760
10	Capital Group	United States	2700
11	Goldman Sachs	United States	2394
12	BNY Mellon	United States	2266
13	Amundi	France	2251
14	PIMCO	United States	2000
15	Legal & General	United Kingdom	1866
16	Prudential Financial	United States	1620
17	Deutsche Bank	Germany	1615
18	Bank of America	United States	1571
19	Invesco	United States	1556
20	T. Rowe Price	United States	1552



TEXT TOOLS

The screenshot shows the Power BI ribbon with the 'Transform' tab selected. The menu bar includes 'File', 'Home', 'Transform', 'Add Column', 'View', 'Tools', and 'Help'. Under the 'Transform' tab, there are several icons: Group By, Use First Row as Headers, Transpose, Reverse Rows, Detect Data Type, Rename, Pivot Column, Fill, Move, Unpivot Columns, Convert to List, Replace Values, and a dropdown for 'Data Type: Date'. Below the ribbon, the status bar shows 'Table'.

A context menu is open under the 'Text Column' option. It includes 'Split Column', 'Format', 'Merge Columns', 'Extract', and 'Parse'. A yellow arrow points from the 'Text Column' option in the main ribbon to this menu.

The 'Text Column' context menu is expanded, showing options like 'Length', 'First Characters', 'Last Characters', 'Range', 'Text Before Delimiter', 'Text After Delimiter', and 'Text Between Delimiters'. A yellow arrow points from the 'Text Column' option in the main ribbon to this menu.

Split a text column based on a specific delimiter, number of characters, or other attributes

A context menu for 'Split Column' is shown, listing 'By Delimiter', 'By Number of Characters', 'By Positions', 'By Lowercase to Uppercase', 'By Uppercase to Lowercase', 'By Digit to Non-Digit', and 'By Non-Digit to Digit'. A yellow arrow points from the 'Split Column' option in the main ribbon to this menu.

A context menu for 'Format' is shown, listing 'lowercase', 'UPPERCASE', 'Capitalize Each Word', 'Trim', 'Clean', 'Add Prefix', and 'Add Suffix'. A yellow arrow points from the 'Format' option in the main ribbon to this menu.

Extract characters from text based on fixed lengths, first/last characters, ranges or delimiters

HEY THIS IS IMPORTANT!

You can access many tools from both the **Transform** and **Add Column** menus - the difference is whether you want to **ADD** a new column or **OVERWRITE** an existing one

Format a text column to upper, lower or proper case, or add a prefix or suffix

Tip: Use "Trim" to eliminate leading & trailing spaces, or "Clean" to remove non-printable characters



ASSIGNMENT: TEXT TOOLS

  NEW MESSAGE

From: **Ethan T. Langer (Analytics Manager)**
Subject: **Customer domains**

Hi!
We're looking to better understand where our customers may be coming from, based on their email domains.
Could you please create a new column in the customer table that will allow us do this?

Thanks!
-ETL

[Reply](#) [Forward](#)

Key Objectives

1. Duplicate the email address column and name it **"Domain Name"**
2. In the new column, remove all text/characters except for the domain name
3. Use transformation steps to clean up and capitalize the domain names (i.e. **"Adventure Works"**)
4. Save & Apply changes



SOLUTION: TEXT TOOLS

 NEW MESSAGE

From: **Ethan T. Langer (Analytics Manager)**
Subject: **Customer domains**

Hi!
We're looking to better understand where our customers may be coming from, based on their email domains.
Could you please create a new column in the customer table that will allow us do this?

Thanks!
-ETL

Reply **Forward**

Solution Preview

ABC Domain Name
Adventure Works

PROPERTIES

Name
Customer Lookup
[All Properties](#)

APPLIED STEPS

- Source
- Promoted Headers
- Changed Type
- Changed Type1
- Capitalized Each Word
- Customer Full Name
- Duplicated Column
- Renamed Columns
- Extracted Text After Delimiter
- Extracted Text Before Delimiter
- Replaced Value
- Capitalized Each Word1



NUMERICAL TOOLS

The screenshot shows the Power BI ribbon with the 'Transform' tab selected. A yellow box highlights the 'Number Column' section under the 'Text Column' dropdown. Arrows point from three callout boxes below to specific sections: 'Statistics' points to the first column of the main grid; 'Standard' points to the second column; and 'Trigonometry' points to the third column.

Statistics	Standard	Trigonometry	Information
Sum Minimum Maximum Median Average Standard Deviation Count Values Count Distinct Values	Add Multiply Subtract Divide Integer-Divide Modulo Percentage Percent Of	Absolute Value Power Square Root Exponent Logarithm Factorial	Sine Cosine Tangent Arcsine Arccosine Arctangent
			Is Even Is Odd Sign

Statistics functions allow you to evaluate basic stats for a selected column (sum, min/max, average, count, count distinct, etc.)

Note: These tools return a *SINGLE* value, and are commonly used to explore a table rather than prepare it for loading

Standard, Scientific and Trigonometry tools allow you to apply standard operations (addition, multiplication, division, etc.) or more advanced calculations (power, logarithm, sine, tangent, etc.) to each value in a column

Note: Unlike the Statistics tools, these are applied to each row in the table

Information tools allow you to define binary flags (1/0 or TRUE/FALSE) to mark rows as even, odd, positive or negative



ASSIGNMENT: NUMERICAL TOOLS

  NEW MESSAGE

From: **Ethan T. Langer** (Analytics Manager)
Subject: **Need some stats for leadership**

Hi again,

Leadership is asking us to validate some high-level stats about our products and customers. Can you please help me answer the following questions?

We don't really need to store these values anywhere, so make sure to restore the tables back to their original state once you're done pulling the stats.

Thank you!
-ETL

Reply Forward

Key Objectives

1. What is our average product cost?
2. How many colors do we sell our products in?
3. How many distinct customers do we have?
4. What is the maximum annual customer income?
5. Return the tables to their original state



SOLUTION: NUMERICAL TOOLS

  NEW MESSAGE

From: **Ethan T. Langer (Analytics Manager)**
Subject: **Need some stats for leadership**

Hi again,

Leadership is asking us to validate some high-level stats about our products and customers. Can you please help me answer the following questions?

We don't really need to store these values anywhere, so make sure to restore the tables back to their original state once you're done pulling the stats.

Thank you!
-ETL

Reply Forward

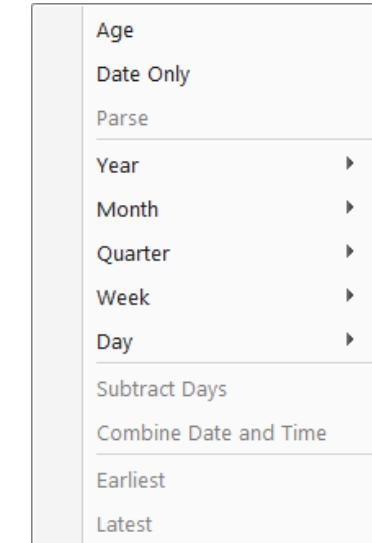
Solution Preview

1. What is our average product cost? (**\$413.66**)
2. How many colors do we sell our products in? (**10**)
3. How many distinct customers do we have? (**18,148**)
4. What is the maximum annual customer income? (**\$170k**)
5. Return the tables to their original state



DATE & TIME TOOLS

The screenshot shows the Power BI ribbon with the "Tools" tab selected. The "Date & Time" tools group is highlighted with a yellow box. It contains icons for "Date", "Time", and "Duration". Below these are dropdown menus for "Date", "Time", and "Duration". A callout arrow points from the "From Date & Time" option in the "Duration" dropdown to a detailed view of the options.



Date & Time tools are relatively straight-forward, and include the following options:

- **Age**: Difference between the current date and the date in each row
- **Date Only**: Removes the time component from a date/time field
- **Year/Month/Quarter/Week/Day**: Extracts individual components from a date field (time-specific options include Hour, Minute, Second, etc.)
- **Earliest/Latest**: Evaluates the earliest or latest date from a column as a single value (can only be accessed from the “Transform” menu)

Note: You will almost always want to perform these operations from the “Add Column” menu to build out new fields, rather than transforming an individual date/time column

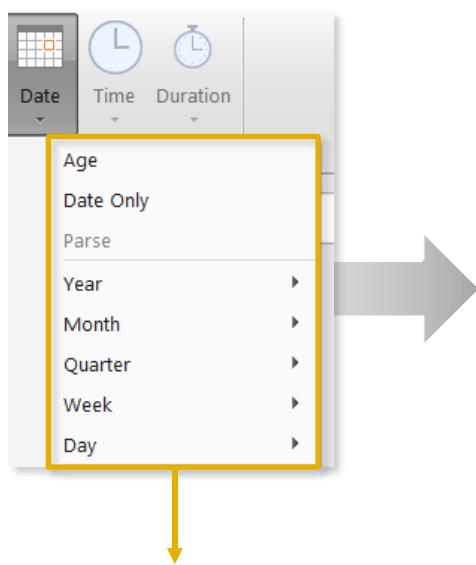


PRO TIP: Load up a table containing a **single date column** and use Date tools to build out an **entire calendar table**



CREATING A CALENDAR TABLE

	Date
1	1/1/2020
2	1/2/2020
3	1/3/2020
4	1/4/2020
5	1/5/2020
6	1/6/2020
7	1/7/2020
8	1/8/2020
9	1/9/2020
10	1/10/2020
11	1/11/2020
12	1/12/2020
13	1/13/2020
14	1/14/2020
15	1/15/2020
16	1/16/2020
17	1/17/2020
18	1/18/2020
19	1/19/2020
20	1/20/2020
21	1/21/2020
22	1/22/2020
23	1/23/2020
24	1/24/2020
25	1/25/2020
26	1/26/2020
27	1/27/2020
28	1/28/2020



Use the **Date** options in the **Add Column** menu to quickly build out an entire calendar table from a list of dates

	Date	Day Name	Start of Week	Start of Month	Month Name
1	1/1/2020	Wednesday	12/29/2019	1/1/2020	January
2	1/2/2020	Thursday	12/29/2019	1/1/2020	January
3	1/3/2020	Friday	12/29/2019	1/1/2020	January
4	1/4/2020	Saturday	12/29/2019	1/1/2020	January
5	1/5/2020	Sunday	1/5/2020	1/1/2020	January
6	1/6/2020	Monday	1/5/2020	1/1/2020	January
7	1/7/2020	Tuesday	1/5/2020	1/1/2020	January
8	1/8/2020	Wednesday	1/5/2020	1/1/2020	January
9	1/9/2020	Thursday	1/5/2020	1/1/2020	January
10	1/10/2020	Friday	1/5/2020	1/1/2020	January
11	1/11/2020	Saturday	1/5/2020	1/1/2020	January
12	1/12/2020	Sunday	1/12/2020	1/1/2020	January
13	1/13/2020	Monday	1/12/2020	1/1/2020	January
14	1/14/2020	Tuesday	1/12/2020	1/1/2020	January
15	1/15/2020	Wednesday	1/12/2020	1/1/2020	January
16	1/16/2020	Thursday	1/12/2020	1/1/2020	January
17	1/17/2020	Friday	1/12/2020	1/1/2020	January
18	1/18/2020	Saturday	1/12/2020	1/1/2020	January
19	1/19/2020	Sunday	1/19/2020	1/1/2020	January
20	1/20/2020	Monday	1/19/2020	1/1/2020	January
21	1/21/2020	Tuesday	1/19/2020	1/1/2020	January
22	1/22/2020	Wednesday	1/19/2020	1/1/2020	January
23	1/23/2020	Thursday	1/19/2020	1/1/2020	January
24	1/24/2020	Friday	1/19/2020	1/1/2020	January
25	1/25/2020	Saturday	1/19/2020	1/1/2020	January
26	1/26/2020	Sunday	1/26/2020	1/1/2020	January
27	1/27/2020	Monday	1/26/2020	1/1/2020	January
28	1/28/2020	Tuesday	1/26/2020	1/1/2020	January

ASSIGNMENT: CALENDAR TABLES



  NEW MESSAGE

From: **Ethan T. Langer (Analytics Manager)**
Subject: **New date fields**

Hi,

We need to add a few fields to our calendar table to help us analyze sales trending over time.

Could you please add the following columns when you get a chance?

Thanks!
-ETL

Reply Forward

Key Objectives

Add the following columns to the calendar table:

1. **Month Name** (e.g. "January")
2. **Month Number** (e.g. "1")
3. **Start of Year** (e.g. "1/1/2020")
4. **Year** (e.g. "2020")



SOLUTION: CALENDAR TABLES

  NEW MESSAGE

From: **Ethan T. Langer (Analytics Manager)**

Subject: **New date fields**

Hi,

We need to add a few fields to our calendar table to help us analyze sales trending over time.

Could you please add the following columns when you get a chance?

Thanks!

-ETL

Reply Forward

Solution Preview

Month Name	Month Number	Start of Year	Year
January	1	1/1/2020	2020
January	1	1/1/2020	2020
January	1	1/1/2020	2020
January	1	1/1/2020	2020
January	1	1/1/2020	2020
January	1	1/1/2020	2020
January	1	1/1/2020	2020
January	1	1/1/2020	2020
January	1	1/1/2020	2020
January	1	1/1/2020	2020
January	1	1/1/2020	2020
January	1	1/1/2020	2020
January	1	1/1/2020	2020
January	1	1/1/2020	2020
January	1	1/1/2020	2020
January	1	1/1/2020	2020
January	1	1/1/2020	2020

PROPERTIES

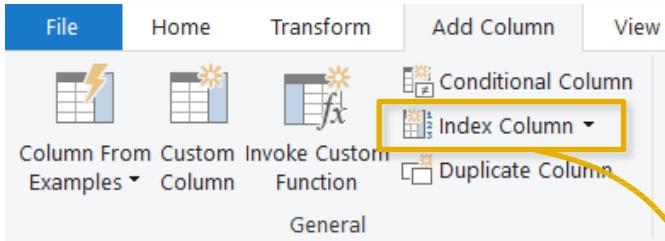
Name
Calendar Lookup
All Properties

APPLIED STEPS

Source
Promoted Headers
Changed Type
Inserted Day Name
Inserted Start of Week
Inserted Start of Month
Inserted Month Name
Inserted Start of Year
Inserted Year
Inserted Month
Renamed Columns



INDEX COLUMNS



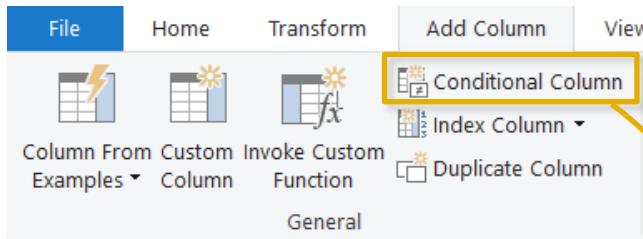
Index Columns contain a list of sequential values that can be used to identify each unique row in a table (*typically starting from 0 or 1*)

These are often used to create **unique IDs** that can be used to form relationships between tables (*more on that later!*)

Index	Order Date	Stock Date	Order Number	Product Key
1	1/1/2020	9/21/2019	SO45080	332
2	1/1/2020	12/5/2019	SO45079	312
3	1/1/2020	10/29/2019	SO45082	350
4	1/1/2020	11/16/2019	SO45081	338
5	1/2/2020	12/15/2019	SO45083	312
6	1/2/2020	10/12/2019	SO45084	310
7	1/2/2020	12/18/2019	SO45086	314
8	1/2/2020	10/9/2019	SO45085	312
9	1/3/2020	10/3/2019	SO45093	312
10	1/3/2020	9/29/2019	SO45090	310
11	1/3/2020	12/11/2019	SO45088	345
12	1/3/2020	10/24/2019	SO45092	313
13	1/3/2020	12/16/2019	SO45089	351
14	1/3/2020	10/26/2019	SO45091	314
15	1/3/2020	9/11/2019	SO45087	350
16	1/3/2020	9/11/2019	SO45094	310
17	1/4/2020	10/30/2019	SO45096	312
18	1/4/2020	10/30/2019	SO45097	313
19	1/4/2020	9/15/2019	SO45098	310
20	1/4/2020	12/7/2019	SO45095	344



CONDITIONAL COLUMNS



Conditional Columns allow you to define new fields based on logical rules and conditions (IF/THEN statements)

Here we're creating a conditional column named **Quantity Type**, which is based on **Order Quantity**:

- If Order Quantity =1, Quantity Type = “**Single Item**”
- Else If Order Quantity >1, Quantity Type = “**Multiple Items**”
- Else; Quantity Type = “**Other**”

New column name: QuantityType

If: Order Quantity equals ABC 123 Then: Single Item

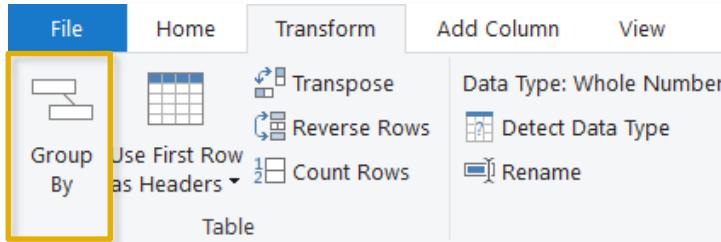
Else If: Order Quantity is greater than ABC 123 Then: Multiple Items

Else: ABC 123 Other

OK Cancel

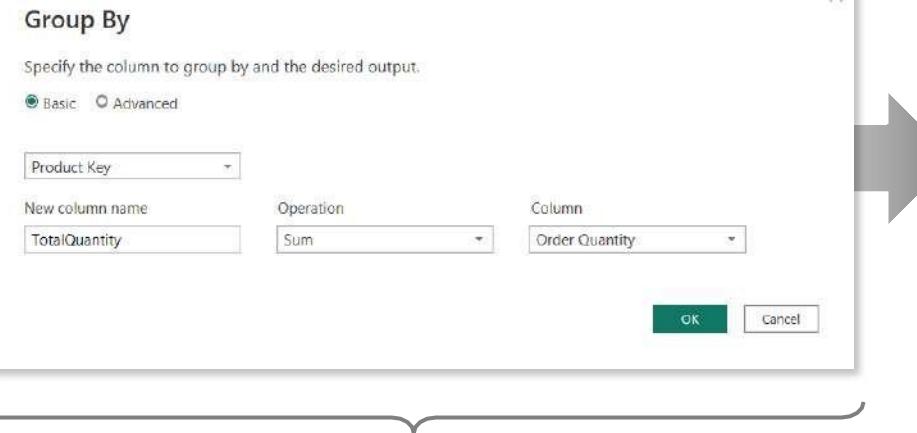


GROUPING & AGGREGATING



Group By allows you to aggregate data at a different level or “grain”
(i.e. group daily records into monthly, aggregate transactions by store, etc.)

	Order Date	Product Key	Customer Key	Order Quantity
1	6/25/2022	214	14719	1
2	10/8/2021	214	21990	1
3	12/30/2021	214	22098	1
4	6/29/2022	214	22748	1
5	8/16/2021	214	27821	1
6	10/9/2021	214	15685	1
7	8/9/2021	214	14951	1
8	1/19/2022	214	23101	1
9	9/23/2021	214	17158	1
10	1/19/2022	214	24196	1
11	6/29/2022	214	12963	1
12	9/13/2021	214	127715	1
13	10/2/2021	214	14846	1
14	7/31/2021	214	11290	1
15	11/24/2021	214	22109	1
16	8/1/2021	214	16982	1
17	10/12/2021	214	20410	1
18	9/10/2021	214	14217	1
19	10/22/2021	214	19642	1
20	8/11/2021	214	11665	1



	Product Key	TotalQuantity
1	214	2099
2	215	1940
3	220	1995
4	223	4151
5	226	392
6	229	408
7	232	424
8	235	381
9	310	169
10	311	139
11	312	179
12	313	168
13	314	157
14	320	65
15	322	39
16	324	72
17	326	65

Here we're transforming a daily, transaction-level table into a summary of **Total Quantity** by **Product Key**

NOTE: Any fields not specified in the Group By settings are lost



GROUPING & AGGREGATING

	Order Date	Product Key	Customer Key	Order Quantity
1	6/25/2022	214	14719	1
2	10/3/2021	214	21990	1
3	12/3/2021	214	22098	1
4	6/29/2022	214	22748	1
5	8/16/2021	214	27821	1
6	10/9/2021	214	15685	1
7	8/3/2021	214	14951	1
8	1/19/2022	214	23101	1
9	9/23/2021	214	17158	1
10	1/19/2022	214	24196	1
11	6/29/2022	214	12963	1
12	9/13/2021	214	12715	1
13	10/2/2021	214	14845	1
14	7/31/2021	214	11290	1
15	11/24/2021	214	22103	1
16	8/1/2021	214	16982	1
17	10/12/2021	214	20410	1
18	9/10/2021	214	14217	1
19	10/22/2021	214	19642	1
20	8/11/2021	214	11666	1

Group By

Specify the columns to group by and one or more outputs.

Basic Advanced

Product Key

Customer Key

Add grouping

New column name

TotalQuantity

Operation

Sum

Column

Order Quantity

Add aggregation

OK

Cancel

	Product Key	Customer Key	TotalQuantity
1	214	19356	1
2	214	15101	1
3	214	12473	1
4	214	12963	1
5	214	26986	1
6	214	13202	1
7	214	14951	1
8	214	11201	1
9	214	19538	1
10	214	22749	1
11	214	15815	1
12	214	19252	1
13	214	14849	1
14	214	11290	1
15	214	27851	1
16	214	16982	1
17	214	22863	2
18	214	19725	1
19	214	15684	1
20	214	11666	1
21	214	26941	1

This time we're transforming the daily, transaction-level table into a summary of **Total Quantity** grouped by both **Product Key** and **Customer Key** (using the "Advanced" option)

NOTE: This is like creating a PivotTable in Excel and pulling in **Sum of Order Quantity** with **Product Key** and **Customer Key** as row labels



PIVOTING & UNPIVOTING

Pivoting describes the process of turning **distinct row values into columns**, and **unpivoting** describes the process of turning **distinct columns into rows**

The diagram illustrates the relationship between two tables. A large yellow curved arrow, oriented from left to right, is labeled "PIVOT" on its upper curve and "UNPIVOT" on its lower curve, symbolizing the reversible nature of the transformation.

Pivoted Table (Top):

	Date	Product Category	North Region	Central Region	South Region
1	7/1/2022	Bikes	10	19	25
2	7/1/2022	Components	14	31	16
3	7/1/2022	Clothing	35	32	46

Unpivoted Table (Bottom):

	Date	Product Category	Region	Quantity Sold
1	7/1/2022	Bikes	North Region	10
2	7/1/2022	Bikes	Central Region	19
3	7/1/2022	Bikes	South Region	25
4	7/1/2022	Components	North Region	14
5	7/1/2022	Components	Central Region	31
6	7/1/2022	Components	South Region	16
7	7/1/2022	Clothing	North Region	35
8	7/1/2022	Clothing	Central Region	32
9	7/1/2022	Clothing	South Region	46

Imagine the table on a hinge; **pivoting** rotates it from **vertical** to **horizontal**, and **unpivoting** rotates it from **horizontal** to **vertical**

NOTE: **Transpose** works very similarly, but doesn't recognize unique values; instead, the entire table is transformed so that each row becomes a column and vice versa



MERGING QUERIES

Merge Queries

Merge

Select a table and matching columns to create a merged table.

Sales Data

Order Date	Product Key	Customer Key	Order Quantity	Index	Stock Date	Order Number	Territory
6/25/2022	214	14719	1	55115	4/20/2022	S053780	
10/6/2021	214	21990	1	14247	7/2/2021	S055746	
12/30/2021	214	22098	1	26322	11/10/2021	S061052	
6/29/2022	214	22748	1	55740	4/9/2022	S074069	
6/29/2022	214	27700	1	55740	7/15/2022	S074069	

Product Lookup

Product Key	Product Subcategory Key	Product SKU	Product Name	Model Name	Category
214	32	HL-U509-R	Sport-100 Helmet, Red	Sport-100	Universal fit, v
215	31	HL-U509	Sport-100 Helmet, Black	Sport-100	Universal fit, v
218	23	SO-B909-M	Mountain Bike Socks, M	Mountain Bike Socks	Combination c
219	23	SO-B909-L	Mountain Bike Socks, L	Mountain Bike Socks	Combination c

Join Kind

Left Outer (all from first, matching from second)

Use fuzzy matching to perform the merge

Fuzzy matching options

The selection matches 56046 of 56046 rows from the first table.

OK Cancel

Merging queries allows you to **join tables** based on a common column (like a lookup in Excel)

In this case we're merging the **Sales Data** table with the **Product Lookup** table, which share a common **Product Key** column

NOTE: Merging **adds columns** to an existing table/query

HEY THIS IS IMPORTANT!

Just because you can merge tables, doesn't mean you should!



In many cases, it's better to keep tables separate and define **relationships** between them in the data model (*more on that soon!*)



APPENDING QUERIES

Merge Queries ▾

Append Queries ▾ **Append**

Combine Files

Combine

Concatenate rows from two tables into a single table.

Two tables Three or more tables

First table
AdventureWorks Sales Data 2020

Second table
AdventureWorks Sales Data 2021

Appending queries allows you to **combine** or **stack** tables sharing the exact same column structure and data types

Here we're appending the **AdventureWorks Sales 2020** table to the **AdventureWorks Sales 2021** table, which is valid since they share identical table structures

NOTE: Appending **adds rows** to an existing table/query



PRO TIP: Use the **Folder** option (Get Data > More > Folder) to **append all files within a specified folder** (assuming they share the same structure); as you add new files, simply refresh the query and they will automatically append!



PRO TIP: APPENDING FILES FROM A FOLDER

The screenshot illustrates a workflow for appending files from a folder in Power BI.

Get Data Interface: The main window shows the "Get Data" interface with various options like "Get data", "Excel", "Data", and "SQL". The "More..." button is highlighted with a yellow box. The "Common data sources" section includes "Excel workbook", "Power BI datasets", "Dataflows", "Dataverse", "SQL Server", "Analysis Services", "Text/CSV", "Web", "OData feed", "Blank query", and "Power BI Template Apps".

Folder Selection: In the "Get Data" interface, the "Folder" option under the "All" category is selected and highlighted with a yellow box. This leads to the "Folder" dialog box.

Folder Dialog: The "Folder" dialog box shows the "Folder path" as "C:\Users\Branislav Poljasevic\Documents\3. PowerBI Desktop\Sales". The "OK" button is visible.

Content View: A preview window displays the contents of the folder, listing three CSV files:

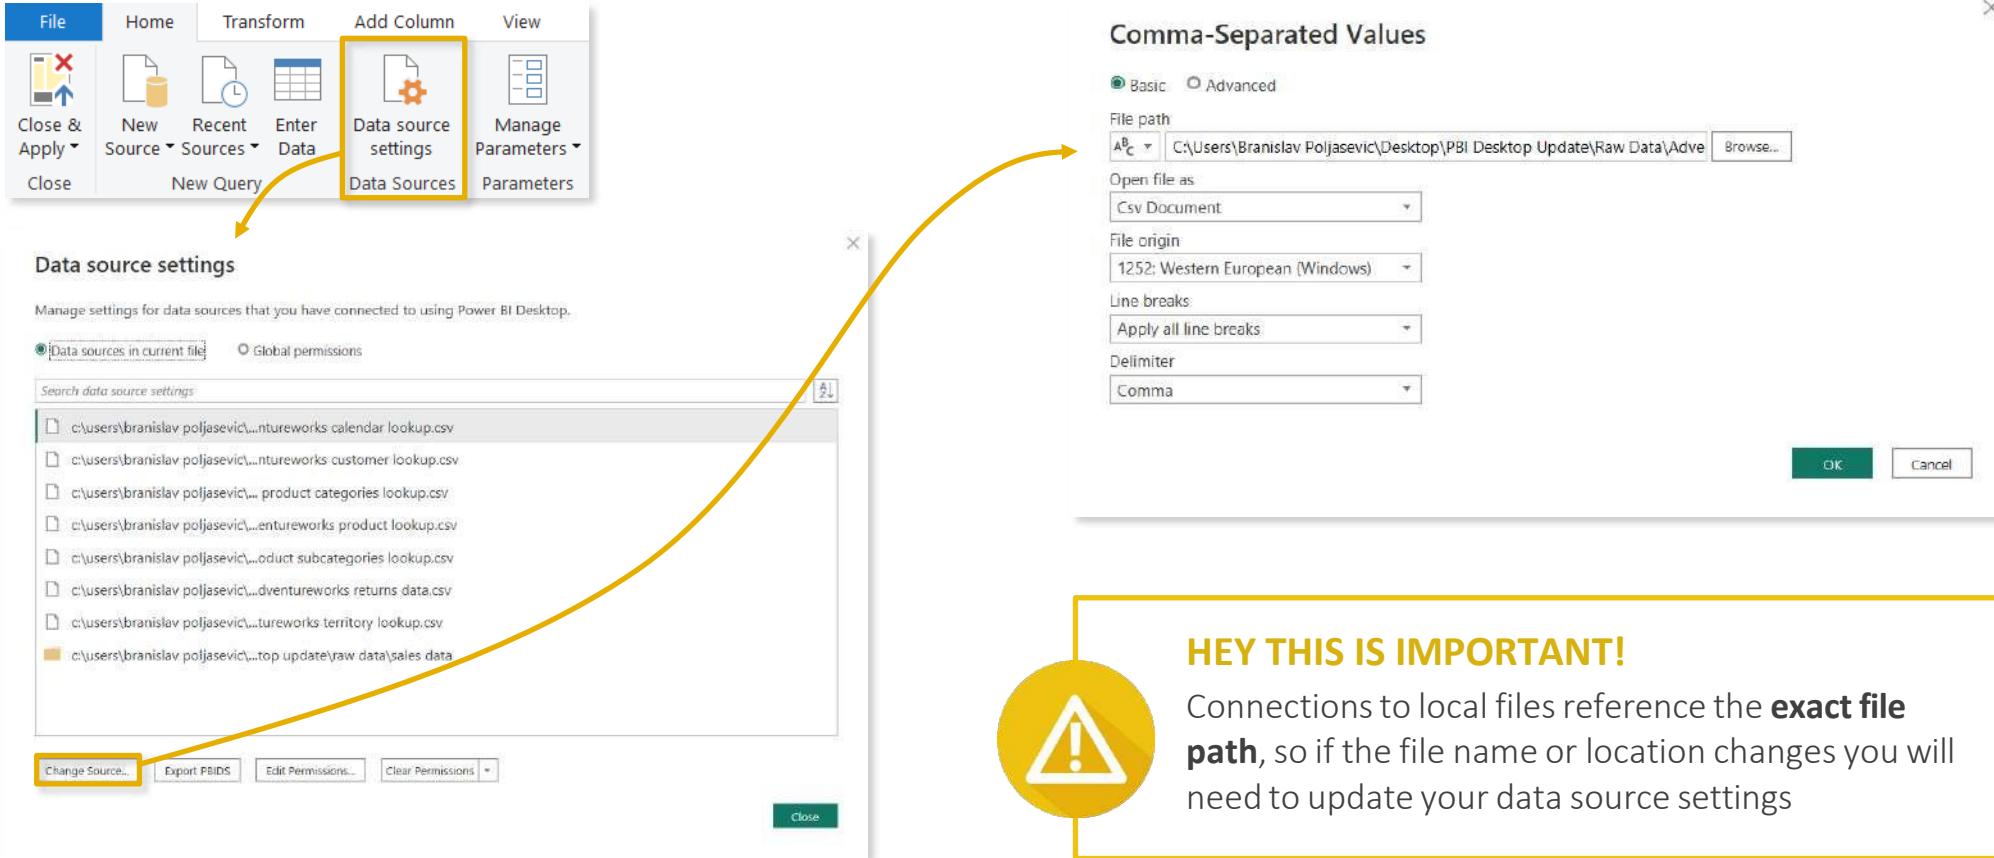
Content	Name	Extension	Date accessed	Date modified	Date created	Attributes
Binary	AdventureWorks Sales Data 2020.csv	.csv	12/11/2022 6:17:52 PM	11/3/2022 4:09:09 PM	12/11/2022 6:17:52 PM	Record
Binary	AdventureWorks Sales Data 2021.csv	.csv	12/11/2022 6:17:52 PM	11/3/2022 4:06:28 PM	12/11/2022 6:17:52 PM	Record
Binary	AdventureWorks Sales Data 2022.csv	.csv	12/11/2022 6:17:52 PM	11/3/2022 7:08:24 PM	12/11/2022 6:17:52 PM	Record

Action Buttons: At the bottom of the preview window, the "Transform Data" button is highlighted with a yellow box.



DATA SOURCE SETTINGS

Data Source Settings allow you to manage existing data connections, file paths and permissions



The screenshot shows two windows from Power BI Desktop. The left window is titled 'Data source settings' and lists several local CSV files connected to the current file. The right window is a 'Comma-Separated Values' configuration dialog, showing the file path as 'C:\Users\Branislav Poljasevic\Desktop\PBI Desktop Update\Raw Data\Adve'. A yellow arrow points from the 'Data source settings' dialog to the 'Data Sources' button in the ribbon, and another yellow arrow points from the 'Data Sources' button to the 'File path' field in the configuration dialog.

Data source settings

Manage settings for data sources that you have connected to using Power BI Desktop.

(Data sources in current file) (Global permissions)

Search data source settings

- c:\users\branislav poljasevic\...ntureworks calendar lookup.csv
- c:\users\branislav poljasevic\...ntureworks customer lookup.csv
- c:\users\branislav poljasevic\... product categories lookup.csv
- c:\users\branislav poljasevic\...ntureworks product lookup.csv
- c:\users\branislav poljasevic\...oduct subcategories lookup.csv
- c:\users\branislav poljasevic\...ntureworks returns data.csv
- c:\users\branislav poljasevic\...tureworks territory lookup.csv
- c:\users\branislav poljasevic\...top update\raw data\sales data**

Change Source... Export PBIDS Edit Permissions... Clear Permissions... Close

Comma-Separated Values

(Basic) (Advanced)

File path: C:\Users\Branislav Poljasevic\Desktop\PBI Desktop Update\Raw Data\Adve

Open file as: Csv Document

File origin: 1252: Western European (Windows)

Line breaks: Apply all line breaks

Delimiter: Comma

OK Cancel

HEY THIS IS IMPORTANT!

Connections to local files reference the **exact file path**, so if the file name or location changes you will need to update your data source settings



PRO TIP: DATA SOURCE PARAMETERS

Use **parameters** to dynamically manage and update connection paths in the Power Query editor

The screenshot shows the Power Query Editor interface. On the left, the ribbon has 'Data source settings' selected. The main area is titled 'Manage Parameters' and shows a list of parameters. A new parameter named 'Database (Fuzzy Factory)' is being edited. The 'Name' field contains 'Database (Fuzzy Factory)'. The 'Type' field is set to 'Text' and has a dropdown menu open, showing 'List of values' which lists 'mavenfuzzyfactory_development' and 'mavenfuzzyfactory_production'. The 'Default Value' field contains 'mavenfuzzyfactory_development' and the 'Current Value' field contains 'mavenfuzzyfactory_production'. The ribbon at the top right has 'Data source settings' highlighted with a yellow box.

Power Query Editor

New Recent Sources Enter Data Data source settings Manage Parameters

Manage Parameters

Name: Database (Fuzzy Factory)

Description:

Required

Type: Text

Suggested Values: List of values

1 mavenfuzzyfactory_development
2 mavenfuzzyfactory_production

Default Value: mavenfuzzyfactory_development

Current Value: mavenfuzzyfactory_production

MySQL database

Server: Server (Fuzzy Factory)

Type: Text

Parameter: Parameter

New Parameter...

Advanced options

Update Server & Database connection text values with **parameters**

Parameter **name**
(Name of the query/table)

Parameter **type**
(Any value, text, date, etc.)

Parameter **value**
(Any value, list, query)

Parameter **type**
(Default & current)



REFRESHING QUERIES

The screenshot shows the Power BI ribbon with the 'Home' tab selected. In the 'Data' section of the ribbon, there is a 'Refresh' button represented by a document with a circular arrow icon. A yellow box highlights this button, and a yellow arrow points from it down towards the explanatory text below.

By default, **all queries** will refresh when you use the **Refresh** command from the **Home** tab

The screenshot shows the Power BI ribbon with the 'Transform' tab selected. In the 'Data Sources' section, there is a 'Refresh' button represented by a document with a circular arrow icon. A yellow box highlights this button.

Queries [13]

Transform File from Sales Data [2]

Other Queries [9]

Calendar Lookup

Customer Lookup

Product Categories

Product Lookup

Product Subcategory

Returns Data

Sales Data

Territory Lookup

Rolling Calendar

- Copy
- Paste
- Delete
- Rename
- Enable load
- Include in report refresh** (highlighted)
- Duplicate
- Reference
- Move To Group
- Move Up
- Move Down
- Create Function...
- Convert To Parameter
- Advanced Editor
- Properties...

From the Query Editor, uncheck **Include in report refresh** to exclude individual queries from the refresh

PRO TIP: Exclude queries from refresh that don't change often (like lookups or static data tables)





POWER QUERY BEST PRACTICES



Get organized before connecting and loading data

- *Define clear and intuitive table/query names from the start, and establish an organized file/folder structure if you are working with local flat files to avoid changes to file names or paths*



Disable report refresh for any static data sources

- *There's no need to constantly refresh data sources that don't change, like lookups or static data tables*



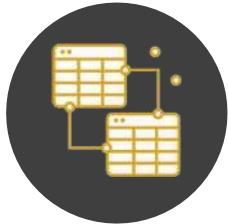
When working with large tables, only load the data you need

- *Don't include hourly data when you only need daily, or transaction-level data when only need a product-level summary (extra data will only slow your report down!)*

CREATING A DATA MODEL



CREATING A DATA MODEL



In this section we'll cover **foundational data modeling topics** like normalization, fact and dimension tables, primary and foreign keys, relationship cardinality and filter flow

TOPICS WE'LL COVER:

Data Modeling 101

Normalization

Facts & Dimensions

Primary & Foreign Keys

Cardinality

Filter Flow

Common Schemas

Hierarchies

GOALS FOR THIS SECTION:

- Understand the basic principles of data modeling, including normalization, fact & dimension tables and common schemas
- Create table relationships using primary and foreign keys, and discuss different types of relationship cardinality
- Configure report filters and trace filter context as it flows between related tables in the model
- Explore data modeling options like hierarchies, data categories and hidden fields



WHAT IS A DATA MODEL?

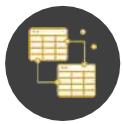
The screenshot shows the Power BI Data Model view. On the left, there's a navigation bar with icons for Home, Reports, and Data. In the center, three tables are listed:

- Product Lookup**: Contains columns for Model Name, Product Color, Product Cost, Product Description, Product Key, and Product Name. A 'Collapse' button is at the bottom.
- Sales Data**: Contains columns for Customer Key, Index, Order Date, Order Line Item, Order Number, Order Quantity, Product Key, Stock Date, Territory Key, and a 'Collapse' button.
- Returns Data**: Contains columns for Product Key, Return Date, Return Quantity, and Territory Key. A 'Collapse' button is at the bottom.

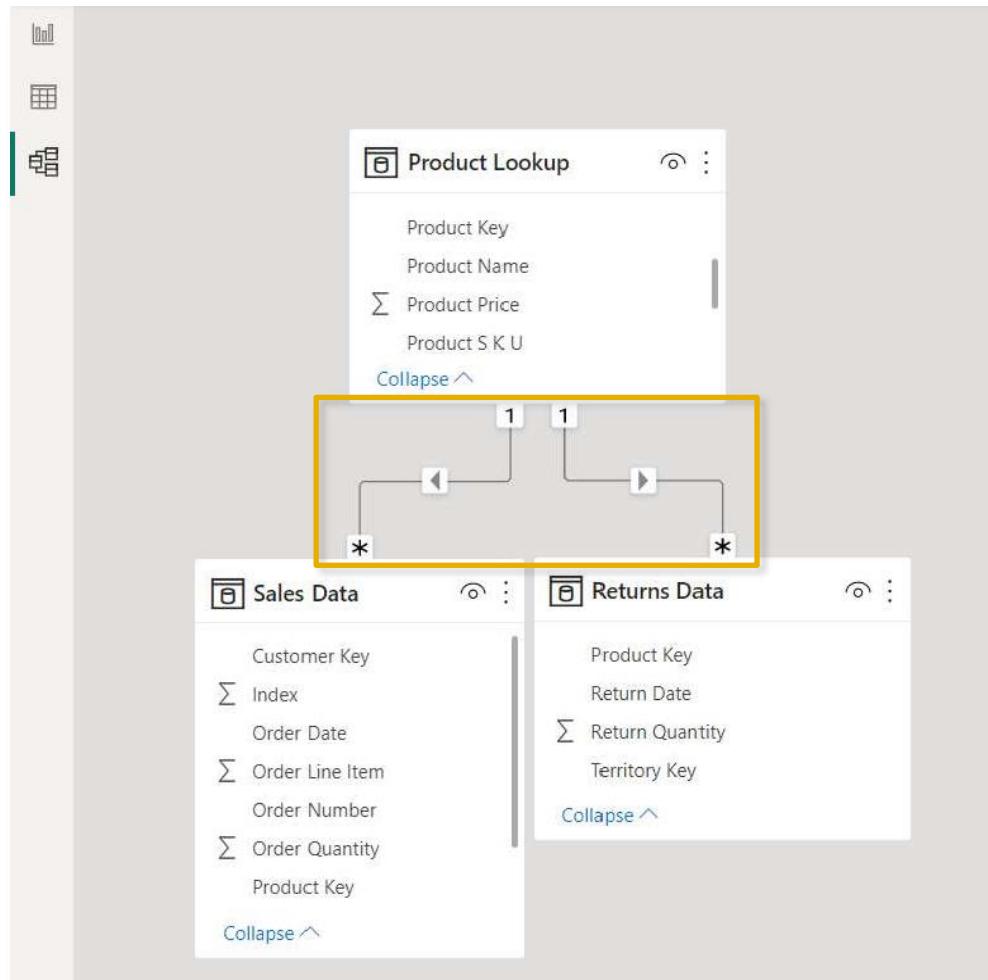
This **IS NOT** a data model 😞

- This is a collection of independent tables, which share no connections or relationships
- If you tried to visualize **Orders** and **Returns** by **Product**, this is what you'd get

ProductName	OrderQuantity	ReturnQuantity
All-Purpose Bike Stand	84,174	1,828
AWC Logo Cap	84,174	1,828
Bike Wash - Dissolver	84,174	1,828
Cable Lock	84,174	1,828
Chain	84,174	1,828
Classic Vest, L	84,174	1,828
Classic Vest, M	84,174	1,828
Classic Vest, S	84,174	1,828
Fender Set - Mountain	84,174	1,828
Total	84,174	1,828



WHAT IS A DATA MODEL?



This **IS** a data model! 😊

- The tables are connected via relationships, based on a common field (Product Key)
- Now **Sales** and **Returns** data can be filtered using fields from the **Product Lookup** table!

ProductName	OrderQuantity	ReturnQuantity
All-Purpose Bike Stand	234	8
AWC Logo Cap	4,151	46
Bike Wash - Dissolver	1,706	25
Classic Vest, L	182	4
Classic Vest, M	182	7
Classic Vest, S	157	8
Fender Set - Mountain	3,960	54
Half-Finger Gloves, L	840	18
Half-Finger Gloves, M	918	16
Total	84,174	1,828



DATABASE NORMALIZATION

Normalization is the process of organizing the tables and columns in a relational database to reduce redundancy and preserve data integrity. It's commonly used to:

- **Eliminate redundant data** to decrease table sizes and improve processing speed & efficiency
- **Minimize errors and anomalies** from data modifications (inserting, updating or deleting records)
- **Simplify queries** and structure the database for meaningful analysis

 In a normalized database, each table should serve a **distinct** and **specific** purpose
(i.e. product information, transaction records, customer attributes, store details, etc.)

date	product_id	quantity	product_brand	product_name	product_sku	product_weight
1/1/1997	869	5	Nationeel	Nationeel Grape Fruit Roll	52382137179	17
1/7/1997	869	2	Nationeel	Nationeel Grape Fruit Roll	52382137179	17
1/3/1997	1	4	Washington	Washington Berry Juice	90748583674	8.39
1/1/1997	1472	3	Fort West	Fort West Fudge Cookies	37276054024	8.28
1/6/1997	1472	2	Fort West	Fort West Fudge Cookies	37276054024	8.28
1/5/1997	2	4	Washington	Washington Mango Drink	96516502499	7.42
1/1/1997	76	4	Red Spade	Red Spade Sliced Chicken	62054644227	18.1
1/1/1997	76	2	Red Spade	Red Spade Sliced Chicken	62054644227	18.1
1/5/1997	3	2	Washington	Washington Strawberry Drink	58427771925	13.1
1/7/1997	3	2	Washington	Washington Strawberry Drink	58427771925	13.1
1/1/1997	320	3	Excellent	Excellent Cranberry Juice	36570182442	16.4

Models that aren't normalized contain **redundant**, **duplicate data**. In this case, all of the product-specific fields could be stored in a separate table containing a unique record for each **product id**

This may not seem critical now, but minor inefficiencies can become major problems at scale!



FACT & DIMENSION TABLES

Data models generally contain two types of tables: **fact** (“data”) tables, and **dimension** (“lookup”) tables:

- **Fact tables** contain **numerical values** or metrics used for summarization (*sales, orders, transactions, pageviews, etc.*)
- **Dimension tables** contain **descriptive attributes** used for filtering or grouping (*products, customers, dates, stores, etc.*)

date	product_id	quantity
1/1/1997	869	5
1/1/1997	1472	3
1/1/1997	76	4
1/1/1997	320	3
1/1/1997	4	4
1/1/1997	952	4
1/1/1997	1222	4
1/1/1997	517	4
1/1/1997	1359	4
1/1/1997	357	4
1/1/1997	1426	5
1/1/1997	190	4
1/1/1997	367	4
1/1/1997	250	5
1/1/1997	600	4
1/1/1997	702	5

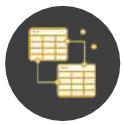
This **Fact** table contains **quantity** values, along with **date** and **product_id** fields

date	day_of_month	month	year	weekday	week_of_year	week_ending	month_name	quarter
1/1/1997	1	1	1997	Wednesday	1	1/5/1997	January	Q1
1/2/1997	2	1	1997	Thursday	1	1/5/1997	January	Q1
1/3/1997	3	1	1997	Friday	1	1/5/1997	January	Q1
1/4/1997	4	1	1997	Saturday	1	1/5/1997	January	Q1
1/5/1997	5	1	1997	Sunday	2	1/5/1997	January	Q1
1/6/1997	6	1	1997	Monday	2	1/12/1997	January	Q1

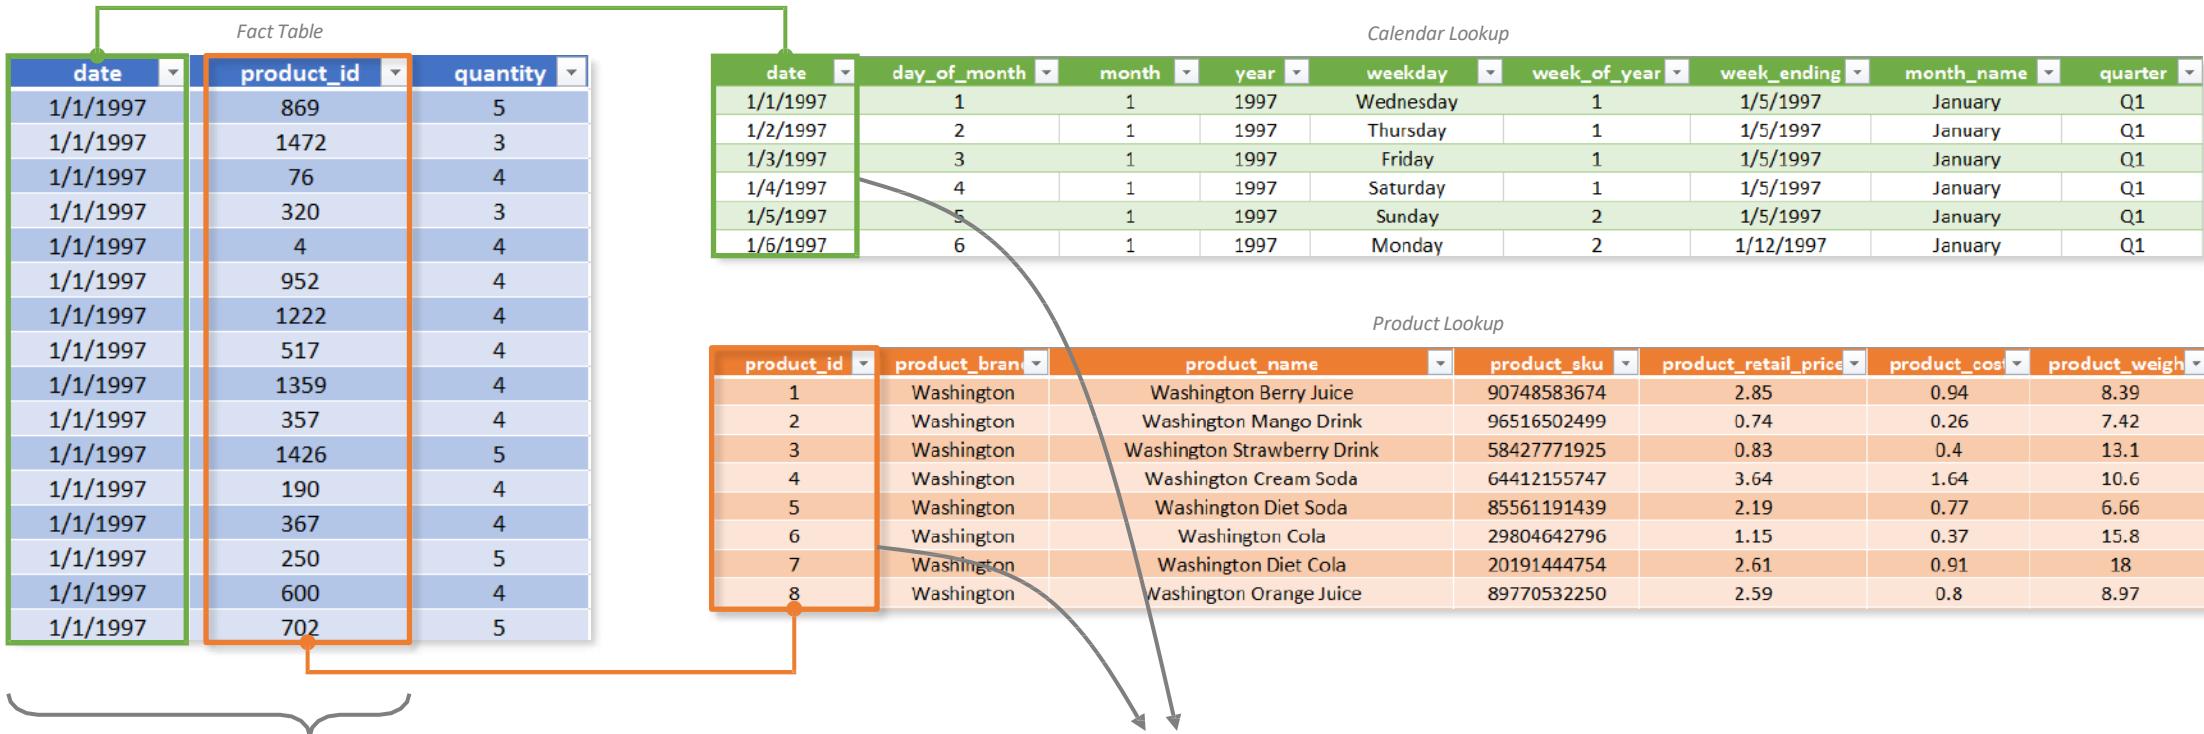
This **Calendar Lookup** table contains attributes about each **date** (month, year, quarter, etc.)

product_id	product_brand	product_name	product_sku	product_retail_price	product_cost	product_weight
1	Washington	Washington Berry Juice	90748583674	2.85	0.94	8.39
2	Washington	Washington Mango Drink	96516502499	0.74	0.26	7.42
3	Washington	Washington Strawberry Drink	58427771925	0.83	0.4	13.1
4	Washington	Washington Cream Soda	64412155747	3.64	1.64	10.6
5	Washington	Washington Diet Soda	85561191439	2.19	0.77	6.66
6	Washington	Washington Cola	29804642796	1.15	0.37	15.8
7	Washington	Washington Diet Cola	20191444754	2.61	0.91	18
8	Washington	Washington Orange Juice	89770532250	2.59	0.8	8.97

This **Product Lookup** table contains attributes about each **product_id** (brand, SKU, price, etc.)



PRIMARY & FOREIGN KEYS



These are **foreign keys (FK)**

*They contain multiple instances of each value, and relate to **primary keys** in dimension tables*

These are **primary keys (PK)**

*They uniquely identify each row of the table, and relate to **foreign keys** in fact tables*



RELATIONSHIPS VS. MERGED TABLES



*Can't I just merge queries or use lookup functions to **pull everything into one single table**?*

- Anonymous confused man

Original Fact Table fields			Attributes from Calendar Lookup table								Attributes from Product Lookup table		
date	product_id	quantity	day_of_month	month	year	weekday	month_name	quarter	product_brand	product_name	product_sku	product_weight	
1/1/1997	869	5	1	1	1997	Wednesday	January	Q1	Nationaleel	Nationaleel Grape Fruit Roll	52382137179	17	
1/7/1997	869	2	7	1	1997	Tuesday	January	Q1	Nationaleel	Nationaleel Grape Fruit Roll	52382137179	17	
1/3/1997	1	4	3	1	1997	Friday	January	Q1	Washington	Washington Berry Juice	90748583674	8.39	
1/1/1997	1472	3	1	1	1997	Wednesday	January	Q1	Fort West	Fort West Fudge Cookies	37276054024	8.28	
1/6/1997	1472	2	6	1	1997	Monday	January	Q1	Fort West	Fort West Fudge Cookies	37276054024	8.28	
1/5/1997	2	4	5	1	1997	Sunday	January	Q1	Washington	Washington Mango Drink	96516502499	7.42	
1/1/1997	76	4	1	1	1997	Wednesday	January	Q1	Red Spade	Red Spade Sliced Chicken	62054644227	18.1	
1/1/1997	76	2	1	1	1997	Wednesday	January	Q1	Red Spade	Red Spade Sliced Chicken	62054644227	18.1	
1/5/1997	3	2	5	1	1997	Sunday	January	Q1	Washington	Washington Strawberry Drink	58427771925	13.1	
1/7/1997	3	2	7	1	1997	Tuesday	January	Q1	Washington	Washington Strawberry Drink	58427771925	13.1	
1/1/1997	320	3	1	1	1997	Wednesday	January	Q1	Excellent	Excellent Cranberry Juice	36570182442	16.4	

You can, **but it's extremely inefficient!**

- Merging tables creates **redundancy** and often requires **significantly more memory and processing power** to analyze compared to a relational model with multiple small tables



THE MODEL VIEW

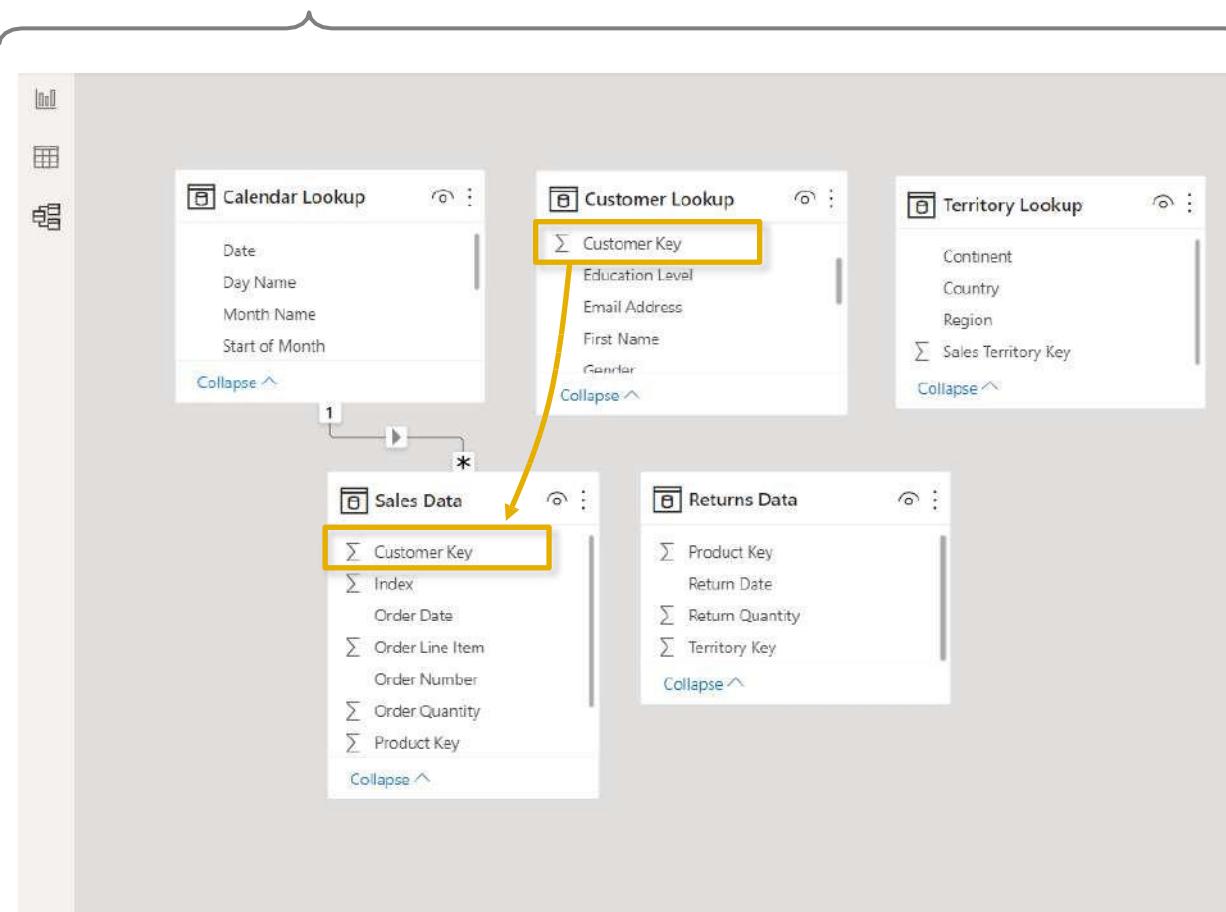
The screenshot illustrates the Power BI Model View interface, which is used for managing data models. The interface includes:

- Menu Ribbon (Home, Help)**: Located at the top left, featuring standard Office-style menu items like File, Home, and Help.
- Model canvas**: The central workspace where data tables and their relationships are visualized. It shows entities like Customer Lookup, Product Lookup, Sales Data, and Returns Data, connected by various relationships.
- Data / Field List**: A pane on the right side listing all fields and entities in the model, such as Day of Week, Month Name, and Product Lookup.
- Properties pane**: A pane on the right side of the model canvas that provides detailed settings for selected objects, including General, Synonyms, Row label, Key column, and Is hidden options.
- View Options**: A set of controls at the bottom right for adjusting the view, including Zoom, Reset Layout, and Fit to Page.
- Model layout tabs**: Buttons at the bottom left for navigating between All tables and a specific table selection.

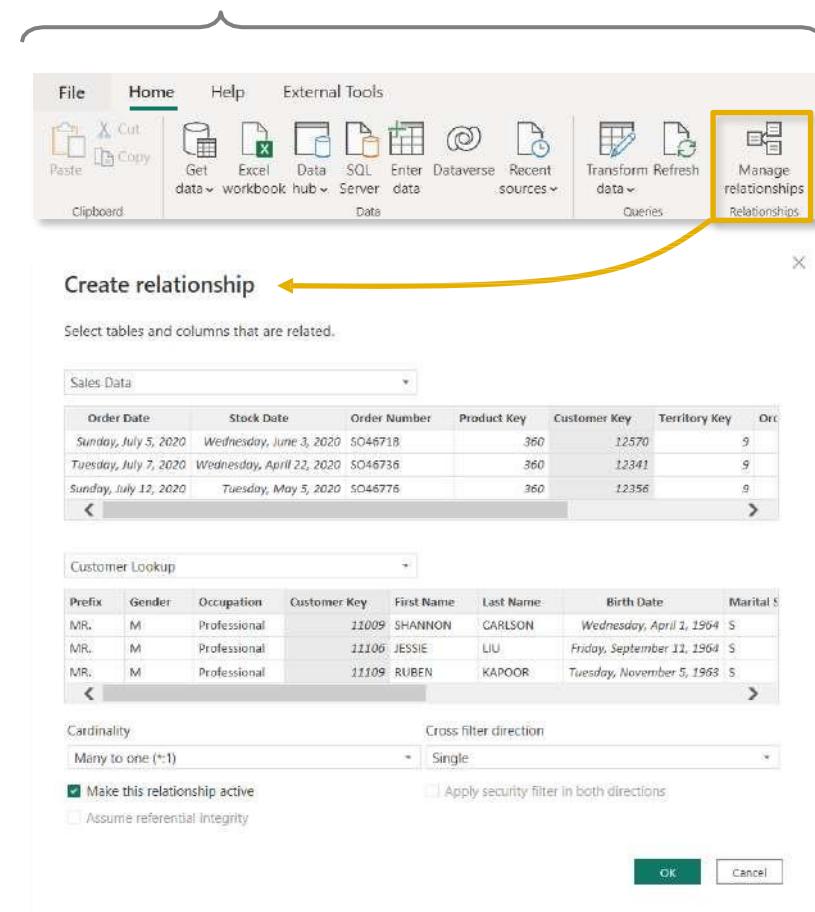


CREATING TABLE RELATIONSHIPS

OPTION 1: Click and drag to connect primary and foreign keys within the **Model** view



OPTION 2: Add or detect relationships using the **Manage Relationships** dialog box





MANAGING & EDITING RELATIONSHIPS

The screenshot shows the Power BI ribbon with the 'Home' tab selected. In the 'Column tools' section, the 'Relationships' icon (a cube with three lines) is highlighted with a yellow box and a callout arrow pointing to the 'Manage relationships' dialog box.

Manage relationships

This dialog box lists existing relationships between tables. Each row shows a checked 'Active' column, a 'From: Table (Column)' column, and a 'To: Table (Column)' column. Relationships listed include Product Lookup, Product Subcategories Lookup, Sales Data (Customer Key), Sales Data (Order Date), and Sales Data (Product Key).

Action buttons: New..., Autodetect..., Edit... (highlighted with a yellow box), Delete, Close.

Launch the **Manage Relationships** dialog box or double-click a relationship to modify it

Edit relationship

Select tables and columns that are related.

The 'Edit relationship' dialog box displays two tables: 'Sales Data' and 'Customer Lookup'. The 'Sales Data' table has columns: Order Date, Stock Date, Order Number, Product Key, Customer Key, Territory Key, and Orc. The 'Customer Lookup' table has columns: Prefix, Gender, Occupation, Customer Key, First Name, Last Name, Birth Date, and Marital Status.

Cardinality: Many to one (*:1)

Cross filter direction: Single

Make this relationship active

Assume referential integrity

Buttons: OK, Cancel.

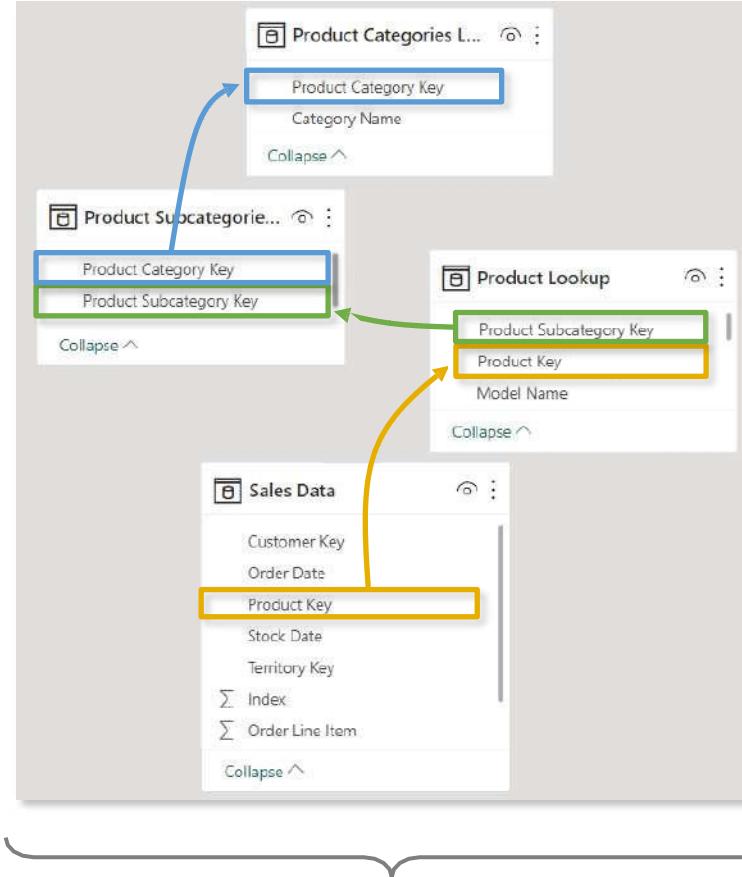
Editing tools allow you to **activate or deactivate** relationships and manage **cardinality** and **filter direction** – more on that soon!



STAR & SNOWFLAKE SCHEMAS



A **star schema** is the simplest and most common type of data model, characterized by a single fact table surrounded by related dimension tables



A **snowflake schema** is an extension of a star, and includes relationships between dimension tables and related sub-dimension tables

ASSIGNMENT: TABLE RELATIONSHIPS





NEW MESSAGE

From: **Dana Modelle (Analyst)**

Subject: **Need a favor...**

Hey there,
Ethan shared the data model you've been working on, and we
might have an issue...

Last night I left my laptop open, and my cat Dennis somehow got
his paws on our model. Now all the relationships are gone!

Could you please rebuild the model, including all three product
tables? I owe you one!

-Dana

Reply **Forward**

Key Objectives

1. Delete all existing table relationships
2. Create a star schema by creating relationships between the Sales, Calendar, Customer, Product and Territories tables
3. Connect all three product tables (Product, Subcategory, Category) in a snowflake schema
4. Use the matrix visual to confirm that you can filter Order Quantity values using fields from each dimension table



SOLUTION: TABLE RELATIONSHIPS

  NEW MESSAGE

From: **Dana Modelle (Analyst)**
Subject: **Need a favor...**

Hey there,
Ethan shared the data model you've been working on, and we might have an issue...

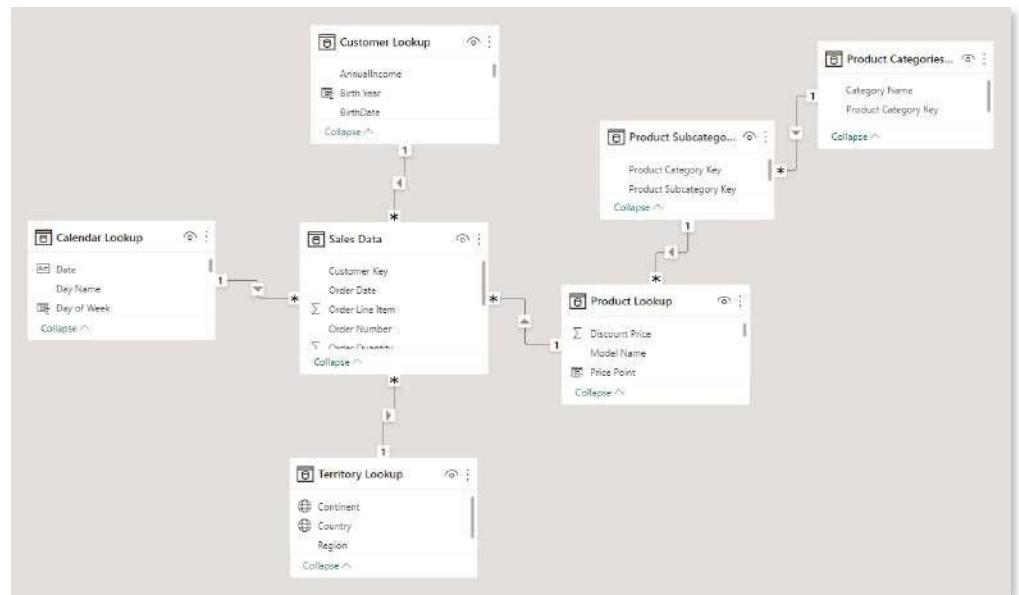
Last night I left my laptop open, and my cat Dennis somehow got his paws on our model. Now all the relationships are gone!

Could you please rebuild the model, including all three product tables? I owe you one!

-Dana

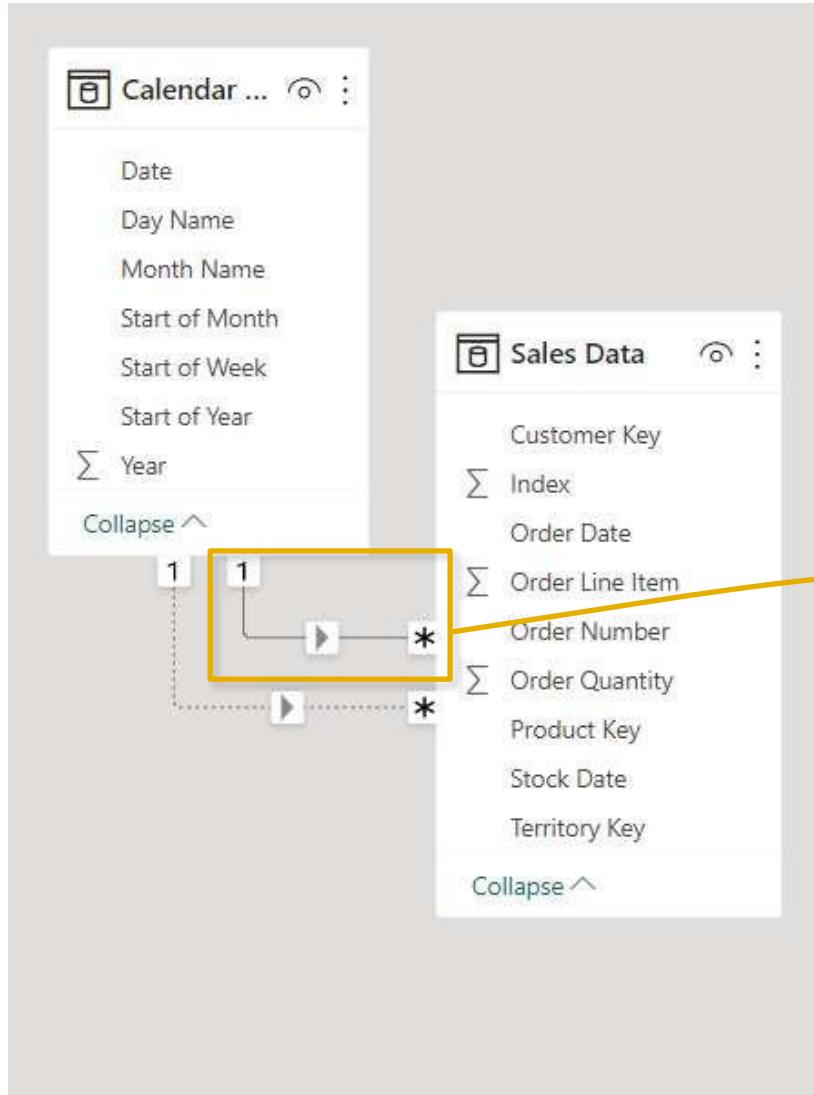
[Reply](#) [Forward](#)

Solution Preview





PRO TIP: ACTIVE & INACTIVE RELATIONSHIPS



Edit relationship

Select tables and columns that are related.

Sales Data

Order Date	Stock Date	Order Number	Product Key	Customer Key	Territory Key	Order Line Item	C
7/5/2020	6/3/2020	SD46718	360	12570	9	1	1
7/7/2020	4/22/2020	SD46736	360	12341	9	1	1
7/12/2020	5/5/2020	SD46776	360	12356	9	1	1

Calendar Lookup

Date	Day Name	Start of Week	Start of Month	Month Name	Start of Year	Year
1/1/2020	Wednesday	12/29/2019	1/1/2020	January	1/1/2020	2020
1/2/2020	Thursday	12/29/2019	1/1/2020	January	1/1/2020	2020
1/3/2020	Friday	12/29/2019	1/1/2020	January	1/1/2020	2020

Cardinality: Many to one (*:1) **Cross filter direction**: Single

Make this relationship active Apply security filter in both directions Assume referential integrity

OK Cancel

Properties

Relationship

Table	Column
Sales Data	Order Date

Cardinality: Many to one (*:1)

Table	Column
Calendar Lookup	Date

Make this relationship active: Yes (switch is on)

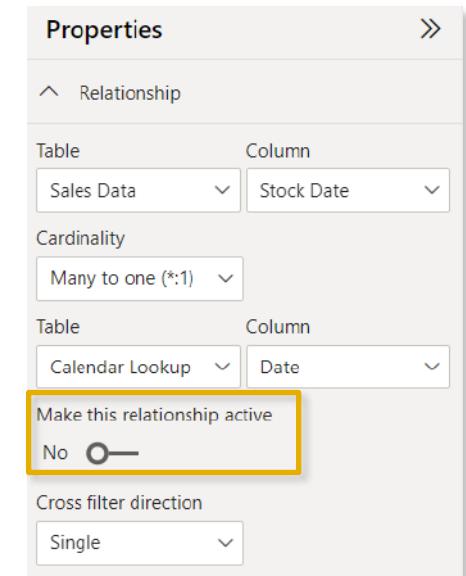
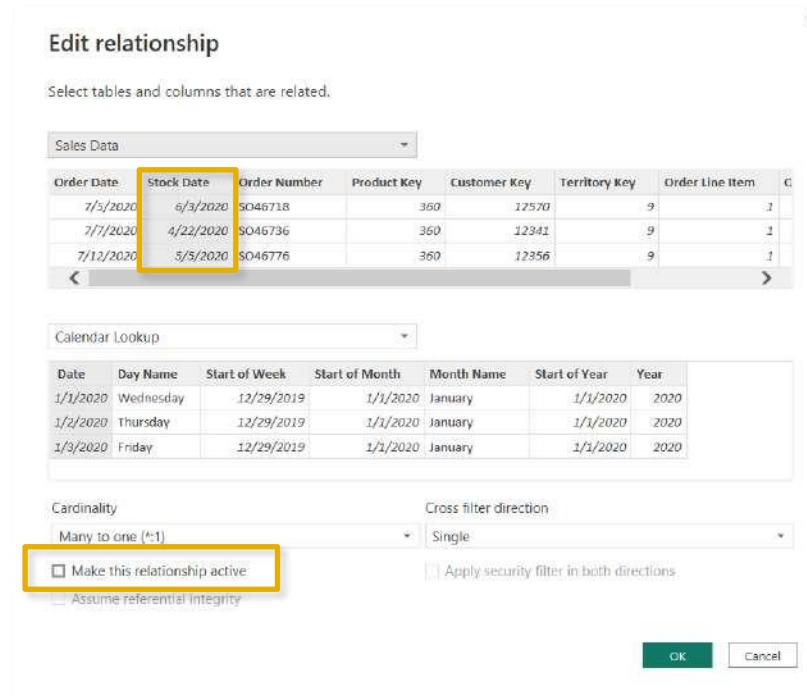
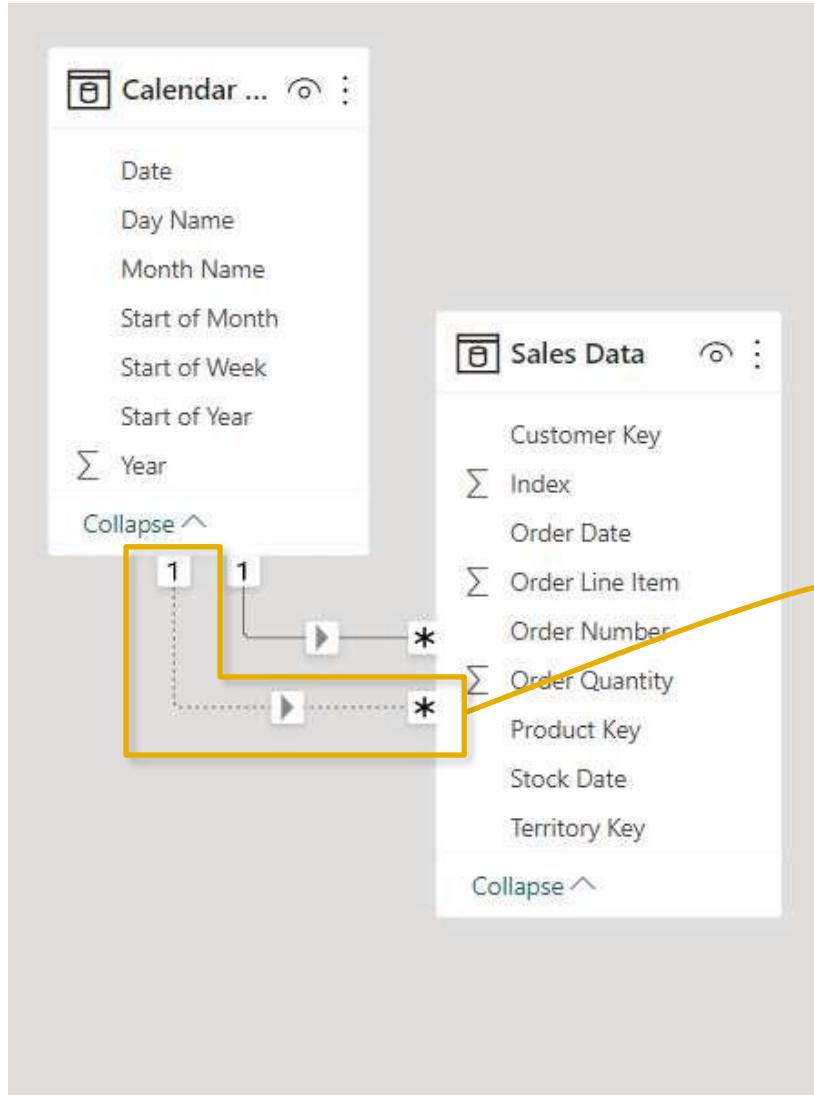
Cross filter direction: Single

The **Sales Data** table contains two date fields (**Order Date** & **Stock Date**), but there can only be **one active relationship** to the Date key in the Calendar table

You can set relationships to active or inactive from either the **Edit Relationships** dialog box or the **Properties** (you must deactivate one before activating another)



PRO TIP: ACTIVE & INACTIVE RELATIONSHIPS

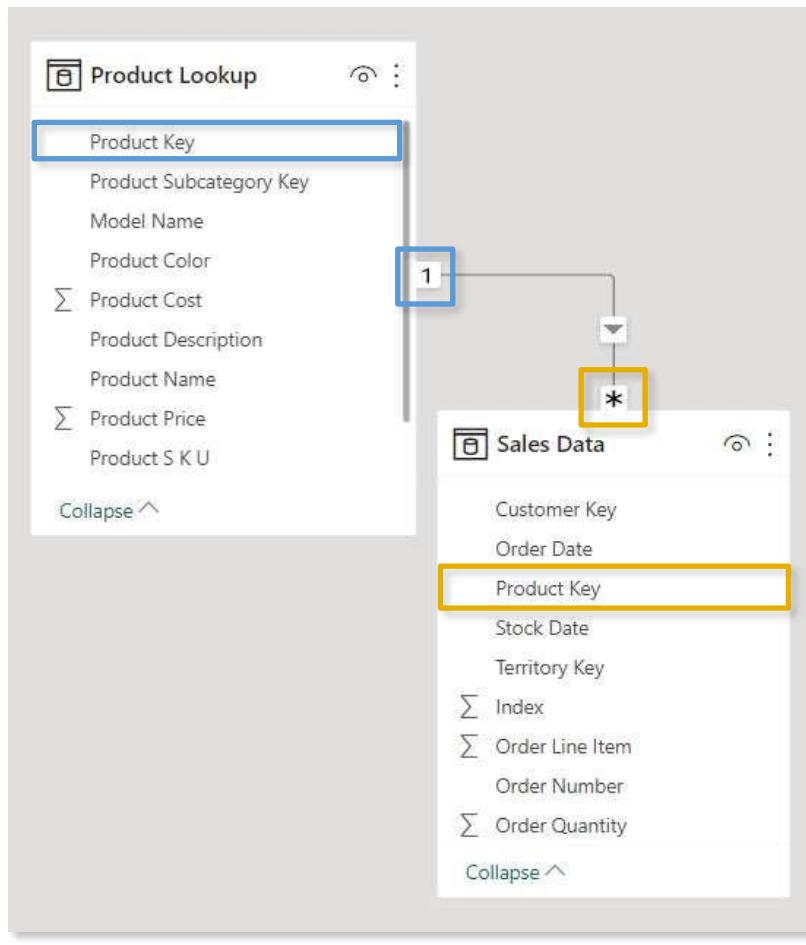


The **Sales Data** table contains two date fields (**Order Date** & **Stock Date**), but there can only be **one active relationship** to the Date key in the Calendar table

You can set relationships to active or inactive from either the **Edit Relationships** dialog box or the **Properties** (you must deactivate one before activating another)



RELATIONSHIP CARDINALITY



Cardinality refers to the uniqueness of values in a column

- Ideally, all relationships in the data model should follow a **one-to-many** cardinality: **one** instance of each primary key, and **many** instances of each foreign key

*In this example there is only **ONE instance of each Product Key** in the Product table (noted by a “1”), since each row contains **attributes of a single product** (name, SKU, description, price, etc.)*

*There are **MANY instances of each Product Key** in the Sales table (noted by an asterisk *), since there are **multiple sales for each product***



EXAMPLE: ONE-TO-ONE CARDINALITY

Product Lookup

product_id	product_name	product_sku
4	Washington Cream Soda	64412155747
5	Washington Diet Soda	85561191439
7	Washington Diet Cola	20191444754
8	Washington Orange Juice	89770532250

Price Lookup

product_id	product_price
4	\$3.64
5	\$2.19
7	\$2.61
8	\$2.59

- Connecting the two tables above using **product_id** creates a **one-to-one relationship**, since each product ID only appears once in each table
- This isn't necessarily a "bad" relationship, but you can simplify the model by merging the tables into a single, valid dimension table

product_id	product_name	product_sku	product_price
4	Washington Cream Soda	64412155747	\$3.64
5	Washington Diet Soda	85561191439	\$2.19
7	Washington Diet Cola	20191444754	\$2.61
8	Washington Orange Juice	89770532250	\$2.59

NOTE: this still respects the rules of normalization, since all rows are unique and capture product-specific attributes



EXAMPLE: MANY-TO-MANY CARDINALITY

Product Lookup

product_id	product_name	product_sku
4	Washington Cream Soda	64412155747
4	Washington Diet Cream Soda	81727382373
5	Washington Diet Soda	85561191439
7	Washington Diet Cola	20191444754
8	Washington Orange Juice	89770532250

Sales

date	product_id	transactions
1/1/2017	4	12
1/2/2017	4	9
1/3/2017	4	11
1/1/2017	5	16
1/2/2017	5	19
1/1/2017	7	11

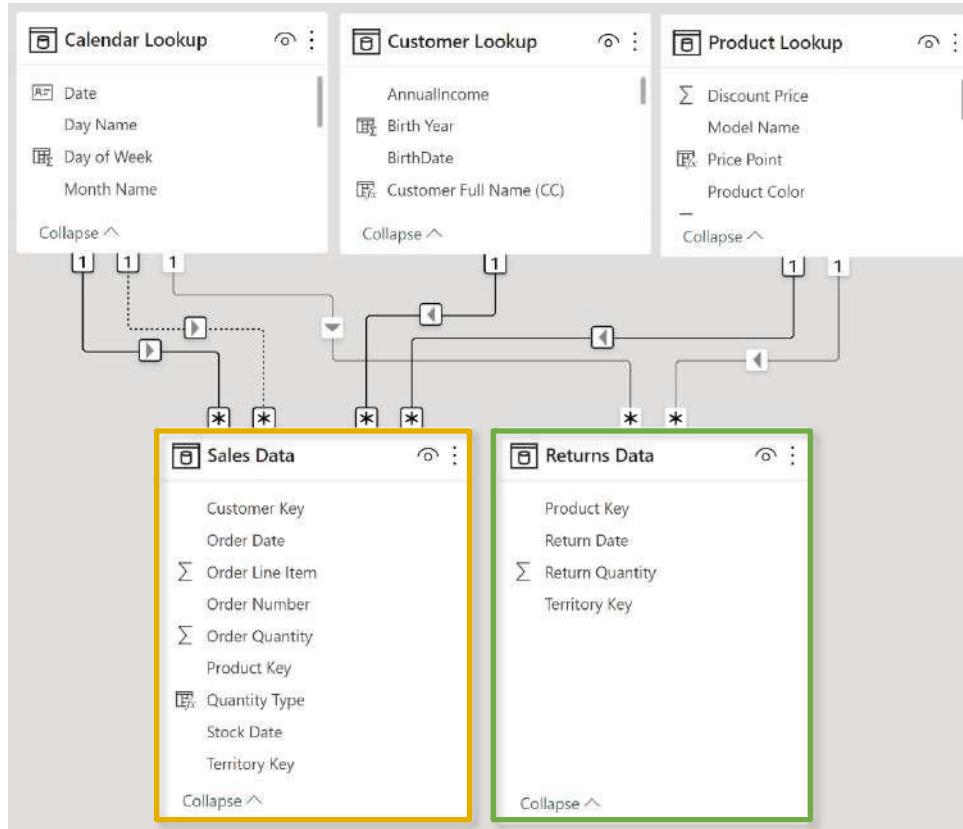


This relationship has cardinality Many-Many. This should only be used if it is expected that neither column (product_id and product_id) contains unique values, and that the significantly different behavior of Many-many relationships is understood. [Learn more](#)

- If we try to connect the tables above using **product_id**, we'll get a **many-to-many relationship** warning since there are multiple instances of product_id in both tables
- Even if we force this relationship, how would we know which product was actually sold on each date – **Cream Soda** or **Diet Cream Soda**?



CONNECTING MULTIPLE FACT TABLES



This model contains two fact tables: **Sales Data** and **Returns Data**

- Since there is no primary/foreign key relationship, we can't connect them directly to each other
- But we *can* connect each fact table to related lookups, which allows us to filter both sales and returns data **using fields from any shared lookup tables**
- We can view orders and returns by product since both tables relate to Product Lookup, but we can't view returns by customer since no relationship exists

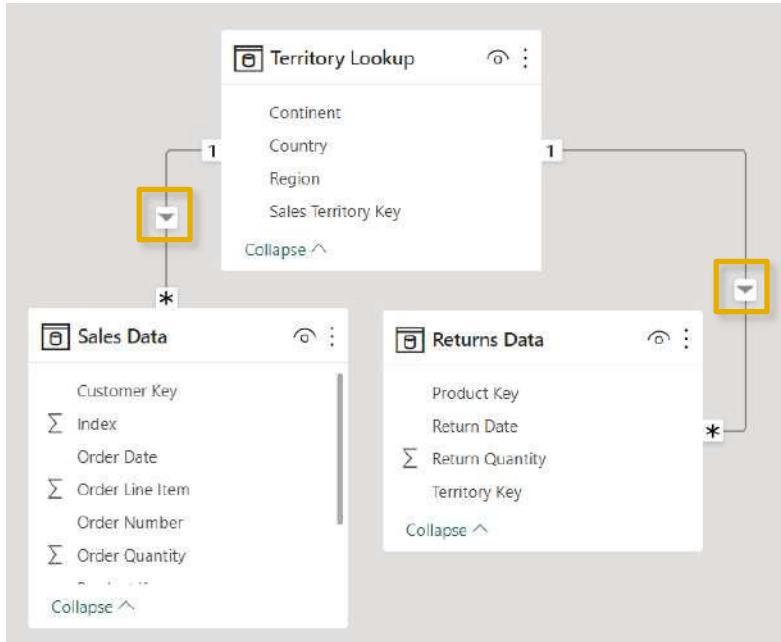


HEY THIS IS IMPORTANT!

Generally speaking, fact tables should **connect through shared dimension tables, not directly to each other**



FILTER CONTEXT & FLOW



Here we have two data tables (**Sales Data** and **Returns Data**), connected to **Territory Lookup**

The arrows show the **filter direction**, and point from the one (**1**) side of the relationship to the many (*) side

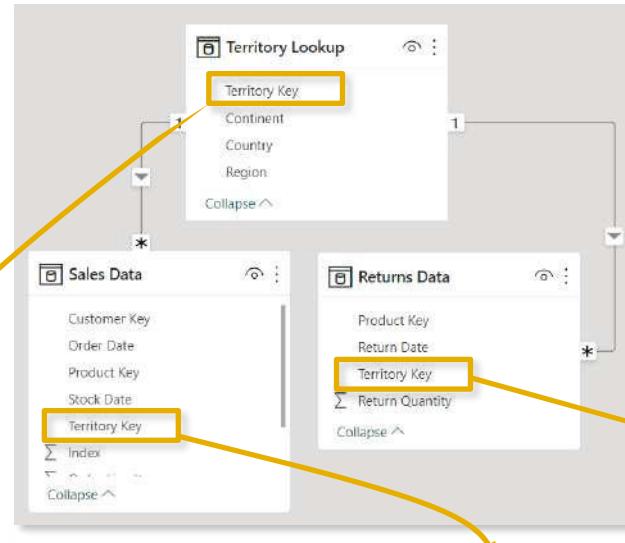
- When you filter a table, that **filter context** is passed to any related “downstream” tables, following the arrow’s direction
- Filter context CANNOT flow “upstream”



PRO TIP: Arrange lookup tables above fact tables in your model as a visual reminder that **filters always flow downstream**



EXAMPLE: FILTER FLOW



TerritoryKey	OrderQuantity	ReturnQuantity
1	12,513	270
2	40	
3	30	
4	17,191	362
5	49	1
6	18,494	38
7	7,862	186
8	7,950	163
9	17,951	404
10	9,694	204
Total	84,174	1,828

Filtering by **Territory Lookup**[Territory Key]

TerritoryKey	OrderQuantity	ReturnQuantity
1	12,513	1,828
2	40	1,828
3	30	1,828
4	17,191	1,828
5	49	1,828
6	18,494	1,828
7	7,862	1,828
8	7,950	1,828
9	17,951	1,828
10	9,694	1,828
Total	84,174	1,828

Filtering by **Sales Data**[Territory Key]

TerritoryKey	OrderQuantity	ReturnQuantity
1	84,174	270
4	84,174	362
5	84,174	1
6	174	238
7	4,174	186
8	84,174	163
9	84,174	404
10	84,174	204
Total	84,174	1,828

Filtering by **Returns Data**[Territory Key]

In this model, the only way to filter both **Sales** and **Returns** data by **Territory** is to use the **Territory Key** from the lookup table, which is upstream and related to both fact tables

- Filtering using Territory Key from the **Sales** table yields **incorrect Returns values**, since the filter context can't flow to any other table
- Filtering using Territory Key from the **Returns** table yields **incorrect Sales values**, and is limited to territories that exist in the returns table



BI-DIRECTIONAL FILTERS

Edit relationship

Select tables and columns that are related.

Sales Data

Order Date	Stock Date	Order Number	Product Key	Customer Key	Territory Key	Or
Sunday, July 5, 2020	Wednesday, June 3, 2020	SO46718	360	12570	9	
Tuesday, July 7, 2020	Wednesday, April 22, 2020	SO46736	360	12341	9	
Sunday, July 12, 2020	Tuesday, May 5, 2020	SO46776	360	12356	9	

Territory Lookup

Region	Country	Continent	Sales Territory Key
Northwest	United States	North America	1
Northeast	United States	North America	2
Central	United States	North America	3

Cardinality

Many to one (*:1)

Make this relationship active

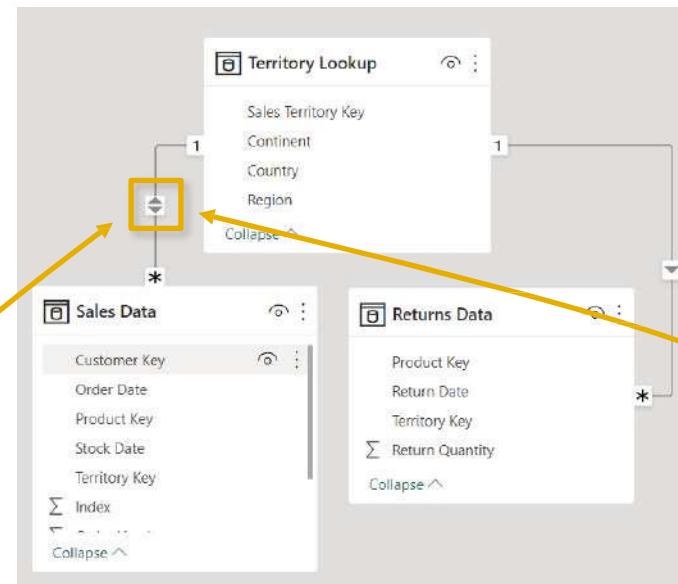
Assume referential integrity

Cross filter direction

Both

Apply security filter in both directions

OK Cancel



Properties

Relationship

Table Column

Sales Data Territory Key

Cardinality

Many to one (*:1)

Table Column

Territory Lookup Sales Territory Key

Make this relationship active

Yes

Cross filter direction

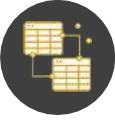
Both

Apply security filter in both directions

No

Updating the **cross-filter direction** from **Single** to **Both** allows filter context to flow in either direction

- In this example, filters applied to the **Sales** table can pass up to the **Territory Lookup** table, then down to **Returns**



EXAMPLE: BI-DIRECTIONAL FILTERS

With two-way cross-filtering enabled between **Sales** and **Territory**, we now see correct values using **Territory Key** from *either* table

- Filter context can now pass up to the **Territory Lookup** table, then downstream to **Returns**
- However, we still see incorrect values when filtering using Territory Key from the **Returns** table, since the filter context is isolated to that single table

Filtering by Territory Lookup[Territory Key]

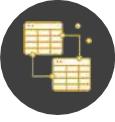
TerritoryKey	OrderQuantity	ReturnQuantity
1	12,513	270
2	40	
3	30	
4	17,191	362
5	49	1
6	894	238
7	7,862	186
8	7,950	163
9	17,951	404
10	9,694	204
Total	84,174	1,828

Filtering by Sales Data[Territory Key]

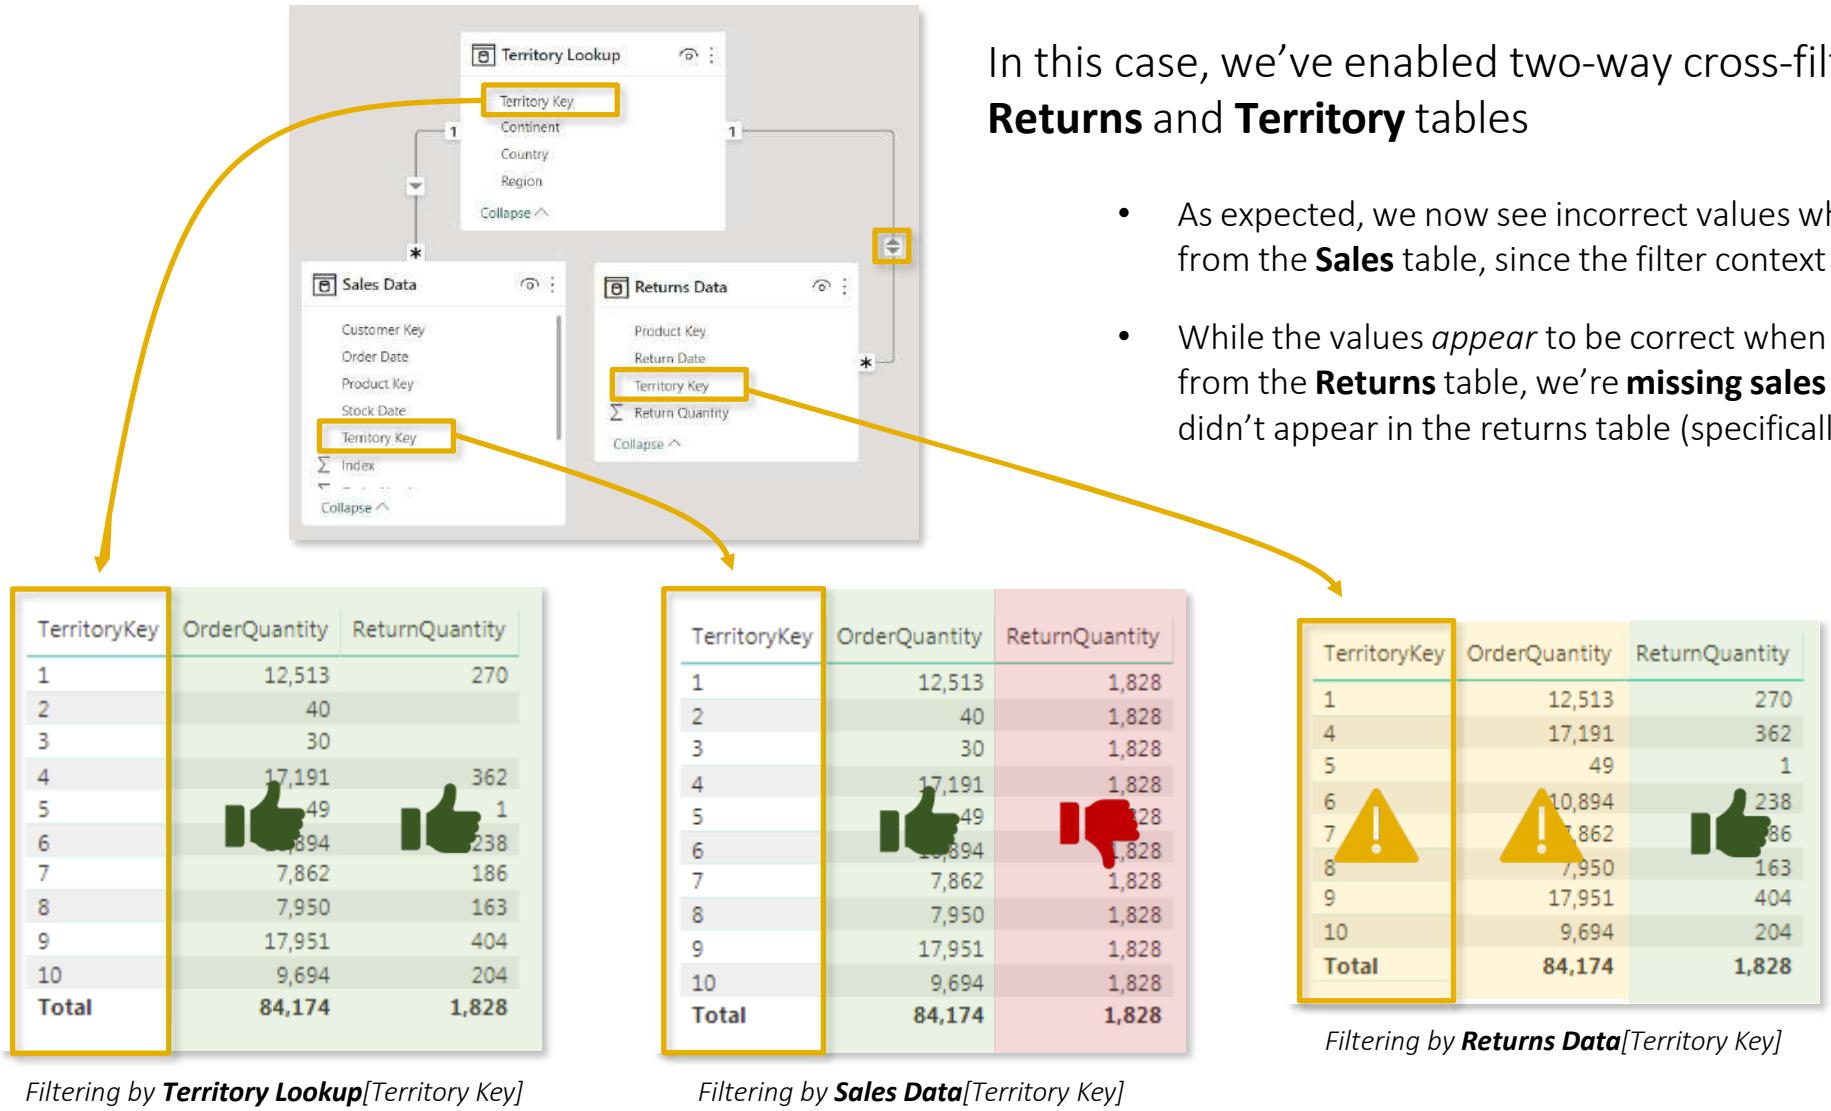
TerritoryKey	OrderQuantity	ReturnQuantity
1	12,513	270
2	40	
3	30	
4	17,191	362
5	49	1
6	894	238
7	7,862	186
8	7,950	163
9	17,951	404
10	9,694	204
Total	84,174	1,828

Filtering by Returns Data[Territory Key]

TerritoryKey	OrderQuantity	ReturnQuantity
1	84,174	270
4	84,174	362
5	84,174	1
6	174	238
7	4,174	186
8	84,174	163
9	84,174	404
10	84,174	204
Total	84,174	1,828



EXAMPLE: BI-DIRECTIONAL FILTERS



In this case, we've enabled two-way cross-filtering between the **Returns** and **Territory** tables

- As expected, we now see incorrect values when filtering using Territory Key from the **Sales** table, since the filter context is isolated to that single table
- While the values *appear* to be correct when filtering using Territory Key from the **Returns** table, we're **missing sales data** from any territories that didn't appear in the returns table (specifically Territories **2 & 3**)

Territories 2 & 3 don't exist in the **Returns** table, so they aren't included in the filter context that passes to **Territory Lookup** and **Sales**

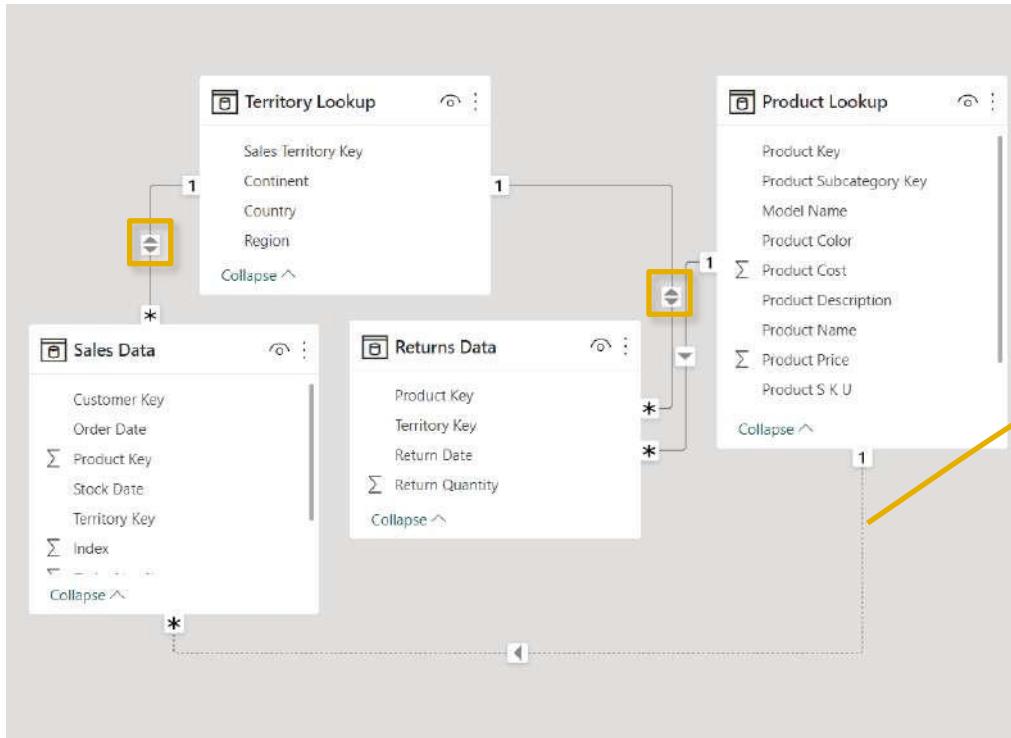
Filtering by **Returns Data**[**Territory Key**]

Filtering by **Territory Lookup**[**Territory Key**]

Filtering by **Sales Data**[**Territory Key**]



AMBIGUITY



Use two-way filters carefully, and **only when necessary**

- Using multiple two-way filters can cause **ambiguity** by introducing multiple filter paths between tables

! You can't create a direct active relationship between Sales_Data and Product_Lookup because that would introduce ambiguity between tables Product_Lookup and Territory_Lookup. To make this relationship active, deactivate or delete one of the relationships between Product_Lookup and Territory_Lookup first.

In this example, filter context from the **Product** table can pass down to **Returns** and up to **Territory Lookup**, which would be filtered based on the Territory Keys passed from the Returns table

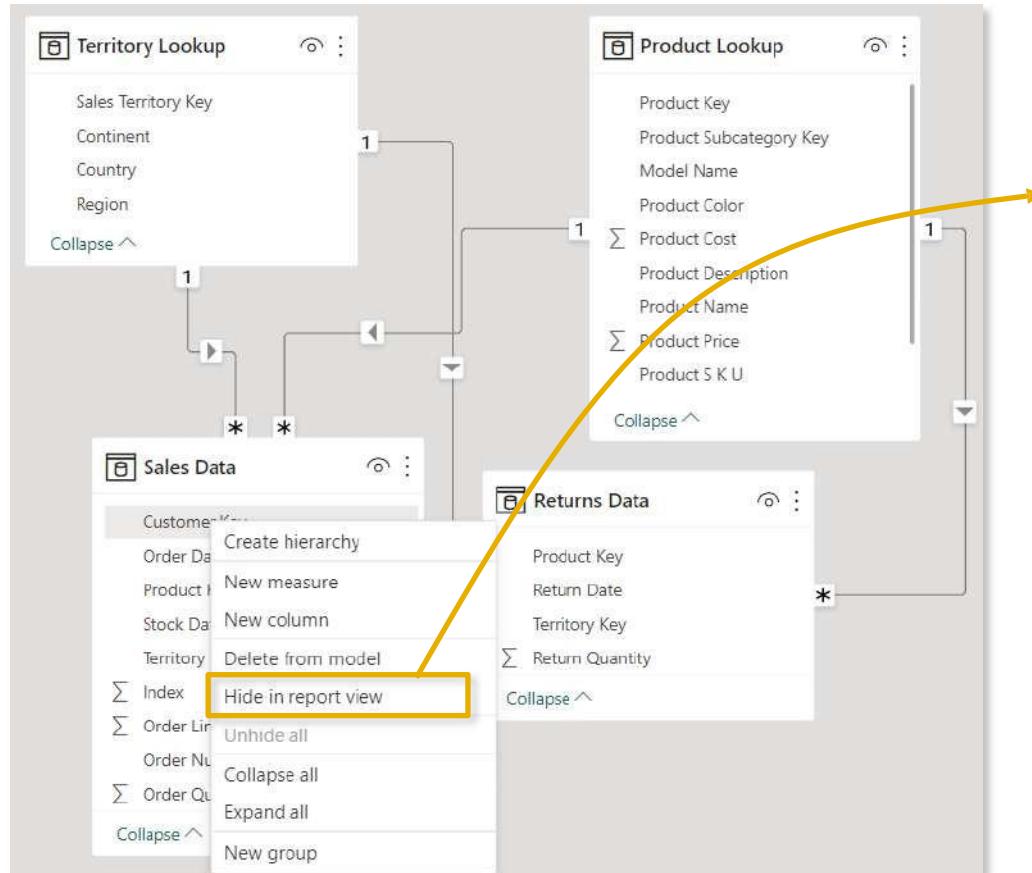
With an active relationship between **Product** and **Sales** as well, filter context could pass through **either the Sales or Returns table to reach the Territory Lookup table**, which could yield conflicting filter context



PRO TIP: Design your models with **one-way filters** and **1:many cardinality** unless more complex relationships are absolutely necessary



HIDING FIELDS

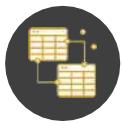


Hide in Report View makes fields inaccessible from the Report tab, but still available in **Data** and **Model** views

- This can be controlled by right-clicking a field in the Data or Model view, or by selecting “**Is hidden**” in the Properties pane
- This is commonly used to prevent users from filtering using invalid fields, reduce clutter, or to hide irrelevant metrics from view



PRO TIP: Hide the **foreign keys** in fact tables to force users to filter using **primary keys** in dimension tables



ASSIGNMENT: FILTER FLOW

 NEW MESSAGE

From: **Dana Modelle (Analyst)**

Subject: **Larry's gone rogue!**

Hey there, we've got another problem.

Larry from Sales just sent me this screenshot. I think he must have downloaded our Power BI model and messed with some relationships, because I KNOW we had sales for product 338.

Can you help diagnose what's going on, and prevent him from doing this again?

-Dana

P.S. Kevin says hi 

[Reply](#) [Forward](#)

Key Objectives

1. Replicate Larry's matrix below to diagnose what he must have done to the model*

Product Key	Sum of Order Quantity	Sum of Return Quantity
322	55	2
324	72	3
326	65	3
328	75	4
330	51	6
332	64	2
334	63	2
336	50	1
340	56	1
342	72	1
346	24	2

No sales for 338!?!?

- Which product is #338?
 - Why didn't Larry's matrix show any orders?
2. Hide any remaining foreign keys to prevent other users from making the same mistake

*Hint: you may need to temporarily change a relationship to bi-directional



SOLUTION: FILTER FLOW

 NEW MESSAGE

From: **Dana Modelle (Analyst)**

Subject: **Larry's gone rogue!**

Hey there, we've got another problem.

Larry from Sales just sent me this screenshot. I think he must have downloaded our Power BI model and messed with some relationships, because I KNOW we had sales for product 338.

Can you help diagnose what's going on, and prevent him from doing this again?

-Dana

P.S. Kevin says hi 

Reply Forward

Solution Preview

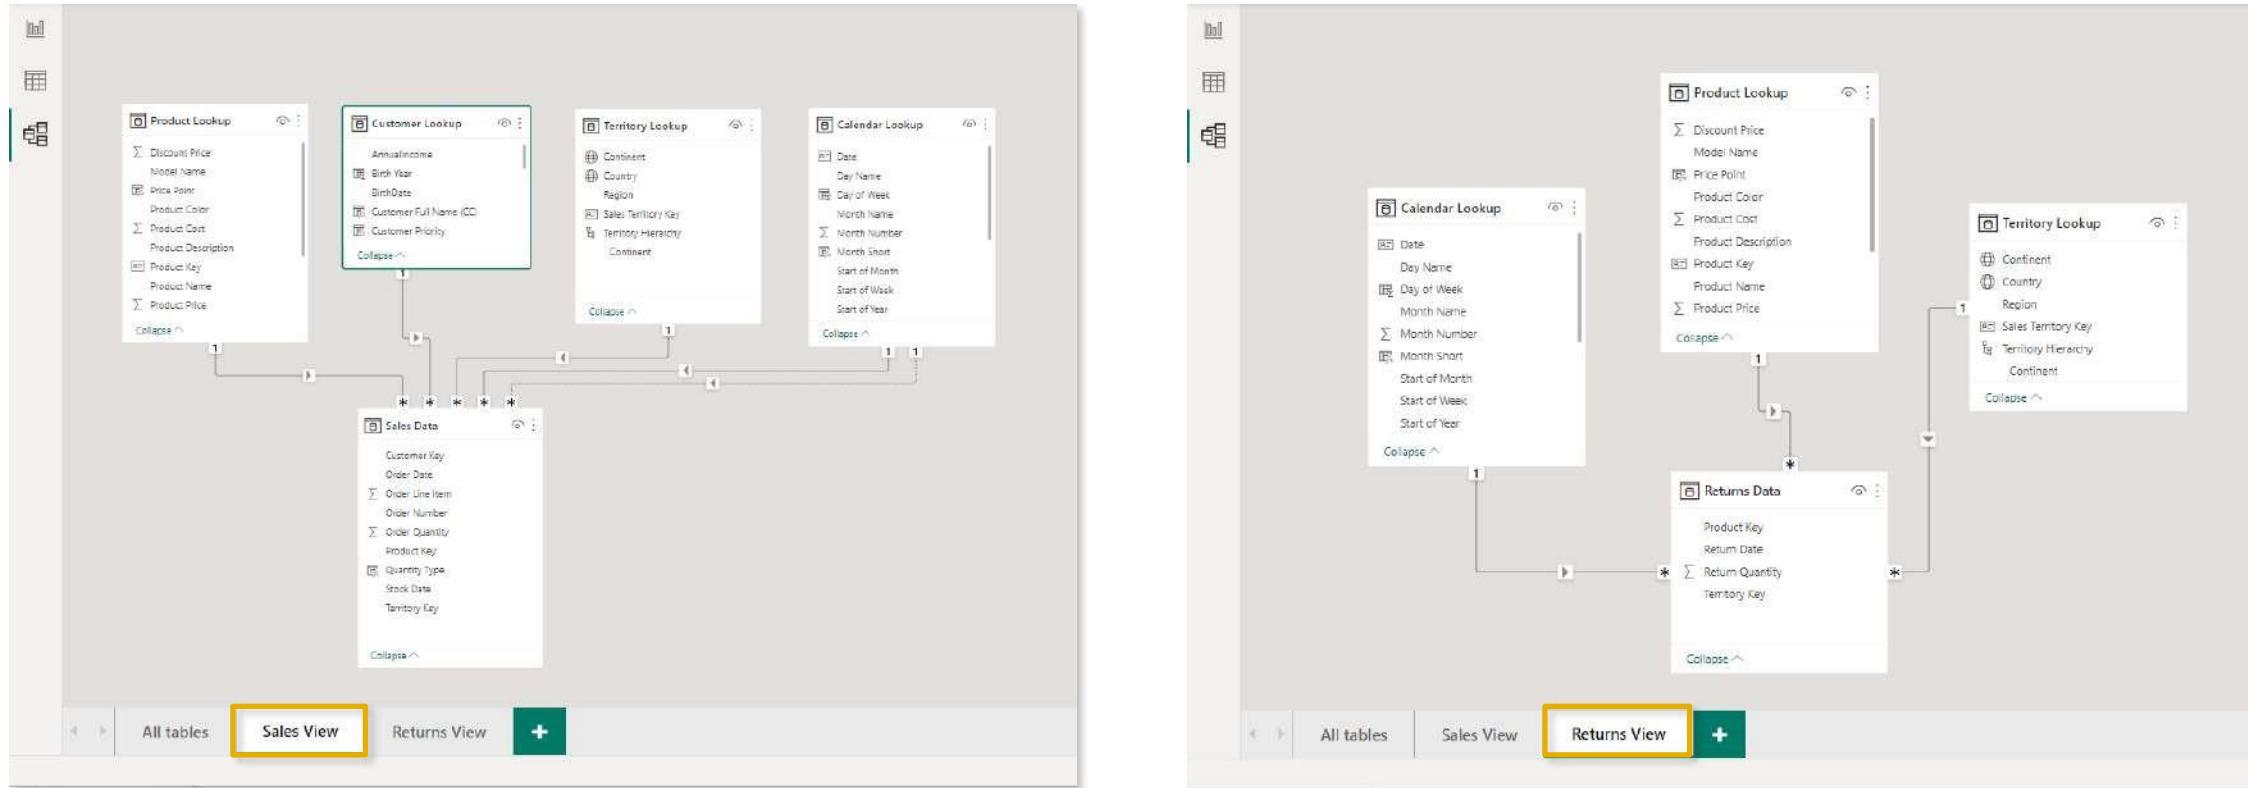
1. Larry must have changed the relationship between **Returns Data** and **Product Lookup** to **bi-directional**, and filtered his matrix using `product_id` from the Returns table
 - Road bike (Road-650 Black, 44)
 - Product 338 doesn't exist in the Returns table, so it was excluded when that filter context passed to the Sales table

2.

>Returns Data	:
Product Key	⟲
Return Date	⟲
Σ Return Quantity	⟲
Σ Territory Key	⟲



PRO TIP: MODEL LAYOUTS



Model layouts allow you to create custom views to show specific portions of large, complex models

- Here we've created a **Sales View** displaying only tables related to sales, and a **Returns View** displaying only tables related to returns (**Note:** this doesn't actually create duplicate tables)



DATA FORMATS & CATEGORIES

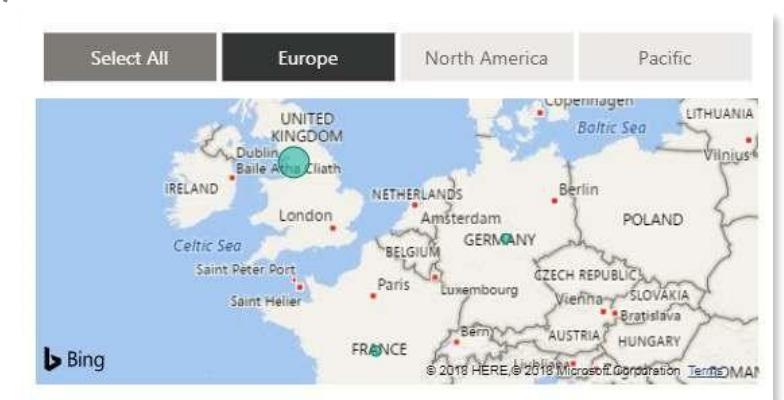
Customize **data formats** from the Column tools menu in the **Data** view or the Properties pane in the **Model** view

The screenshot shows the Power BI desktop interface. A yellow arrow points from the text above to the 'Column tools' tab in the ribbon. Another yellow arrow points from the text below to the 'Data category' dropdown menu, which is open and displays a list of categories including 'Country', 'Uncategorized', 'Address', 'Place', 'City', 'County', 'State or Province', 'Postal code', 'Latitude', 'Longitude', 'Web URL', 'Image URL', and 'Barcode'. The 'Country' option is selected.

Region	Country	Continent	Sales Territory Key	
Northwest	United States	North America		1
Northeast	United States	North America		2
Central	United States	North America		3
Southwest	United States	North America		4
Southeast	United States	North America		5
Canada	Canada	North America		6
France	France	Europe		7
Germany	Germany	Europe		8
Australia	Australia	Pacific		9
United Kingdom	United Kingdom	Europe		10

Assign **data categories** for geospatial fields, URLs or barcodes

- This is commonly used to help Power BI map location-based fields like addresses, countries, cities, coordinates, zip codes, etc.





HIERARCHIES

Hierarchies are groups of columns that reflect multiple levels of granularity

- For example, a **Geography hierarchy** might include **Country**, **State** and **City** fields
- Hierarchies are treated as a **single item** in tables and reports, allowing users to “drill up” and “drill down” through each level

The figure consists of three side-by-side screenshots of the Microsoft Power BI Data pane:

- Screenshot 1:** Shows the context menu for the "Product Price" column. The "Create hierarchy" option is highlighted with a yellow box and connected by a yellow arrow to the same option in Screenshot 2.
- Screenshot 2:** Shows the "Territory Lookup" section expanded. The "Territory Hierarchy" node is selected and highlighted with a yellow box. A yellow arrow points from Screenshot 1 to this node.
- Screenshot 3:** Shows the context menu for the "Country" column. The "Add to hierarchy" option is highlighted with a yellow box and connected by a yellow arrow to the same option in Screenshot 1.

In the **Data** pane, right-click a field and select **Create hierarchy**

This hierarchy contains “Continent”, and is named “**Territory Hierarchy**”

Right-click another field (like “Country”) and select **Add to Hierarchy** (or drag it in!)



ASSIGNMENT: HIERARCHIES

  NEW MESSAGE

From: **Dana Modelle (Analyst)**

Subject: **Adding a date hierarchy**

Good morning!

Hoping you can help with a quick request.

Since we'll be doing a lot of time-series analysis, Ethan asked us to add a date hierarchy to the model so that users can quickly view trends at any level of granularity (year, month, day, etc.)

Please get that added before our afternoon call. Thanks!

-Dana

Reply Forward

Key Objectives

1. Create a new hierarchy based on the **Start of Year** field, and name it "**Date Hierarchy**"
2. Right-click or drag to add fields until your hierarchy contains the following (in this order):
 - **Start of Year**
 - **Start of Month**
 - **Start of Week**
 - **Date**
3. Add your new hierarchy to the matrix visual (on rows) and practice drilling up and down between each level of granularity



SOLUTION: HIERARCHIES

  NEW MESSAGE

From: **Dana Modelle (Analyst)**
Subject: **Adding a date hierarchy**

Good morning!

Hoping you can help with a quick request.

Since we'll be doing a lot of time-series analysis, Ethan asked us to add a date hierarchy to the model so that users can quickly view trends at any level of granularity (year, month, day, etc.)

Please get that added before our afternoon call. Thanks!

-Dana

Reply Forward

Solution Preview

Data >

Search

> Measure Table

✓ Calendar Lookup

- Average Revenue per Customer
- Date
- ✓ Date Hierarchy
- Start of Year
- Start of Month
- Start of Week

A Date



DATA MODEL BEST PRACTICES



Focus on building a normalized model from the start

- *Leverage relationships and make sure that each table serves a clear, distinct purpose*



Organize dimension tables above data tables in your model

- *This serves as a visual reminder that filters always flow “downstream”*



Avoid complex relationships unless absolutely necessary

- *Aim to use 1-to-many table relationships and one-way filters whenever possible*



Hide fields from report view to prevent invalid filter context

- *This forces report users to filter using primary keys from dimension tables*

CALCULATED FIELDS WITH DAX

CALCULATED FIELDS WITH DAX



In this section we'll use **Data Analysis Expressions (DAX)** to add calculated columns & measures to our model, and introduce topics like row & filter context, iterators and more

TOPICS WE'LL COVER:

DAX 101

Columns & Measures

Row & Filter Context

DAX Syntax

Common Functions

Calculate

Iterators

Time Intelligence

GOALS FOR THIS SECTION:

- Introduce DAX fundamentals and learn when to use calculated columns and measures
- Understand the difference between row context and filter context, and how they impact DAX calculations
- Learn DAX formula syntax, basic operators and common function categories (*math, logical, text, date/time, filter, etc.*)
- Explore nested functions, and more complex topics like iterators and time intelligence patterns

MEET DAX



Data Analysis Expressions (commonly known as **DAX**) is the formula language that drives the Power BI front-end. With DAX, you can:

- Go beyond the capabilities of traditional spreadsheet formulas, with powerful and flexible functions built specifically to work with relational data models
- Add **calculated columns** (*for filtering*) and **measures** (*for aggregation*) to enhance data models

Two ways to use DAX

Calculated Columns

Customer Data					
Marital Status	Email Address	Annual Income	Total Children	Education Level	Parent
M	emma32@adventure-works.com	70000	5	Bachelors	Yes
M	barry20@adventure-works.com	40000	5	High School	Yes
M	martha13@adventure-works.com	70000	5	High School	Yes
S	tamara16@adventure-works.com	40000	5	High School	Yes
S	gerald21@adventure-works.com	130000	5	Bachelors	Yes
M	alexa8@adventure-works.com	40000	5	High School	Yes
M	jack53@adventure-works.com	70000	5	Graduate Degree	Yes
S	ricky1@adventure-works.com	100000	5	Bachelors	Yes
M	keith4@adventure-works.com	70000	5	Partial College	Yes
M	latoya19@adventure-works.com	70000	5	Bachelors	Yes

Measures

The screenshot shows the Power BI Query Editor with a context menu open. The 'New measure' option is highlighted. Below it, two measures are defined:

```
Total Orders = DISTINCTCOUNT(Sales_Data[OrderNumber])
Total Revenue = SUMX(Sales_Data, Sales_Data[OrderQuantity] * RELATED(Product_Lookup[ProductPrice]))
Quantity Ordered = SUM(Sales_Data[OrderQuantity])
```



M VS. DAX

M and **DAX** are two distinct functional languages used within Power BI Desktop:

- **M** is used in the Power Query editor, and is designed specifically for extracting, transforming and loading data
- **DAX** is used in the Power BI front-end, and is designed specifically for analyzing relational data models

M

Query Editor:

Properties pane shows 'Territory Lookup' selected.

Applied Steps pane shows 'Source', 'Promoted Headers', and 'Changed Type' (highlighted with a yellow box). A yellow arrow points from this step to the M code below.

```
#"Changed Type" = Table.TransformColumnTypes(          // Adding a new step
    #"Promoted Headers",                                // after we promoted headers
    {
        {"SalesTerritoryKey", Int64.Type},               // that changes column datatypes
        {"Region", type text},
        {"Country", type text},
        {"Continent", type text}
    }
)
```

DAX

Report View:

Category Name	Total Returns	Bike Returns
Accessories	1,115	
Bikes	427	427
Clothing	267	
Total	1,809	427

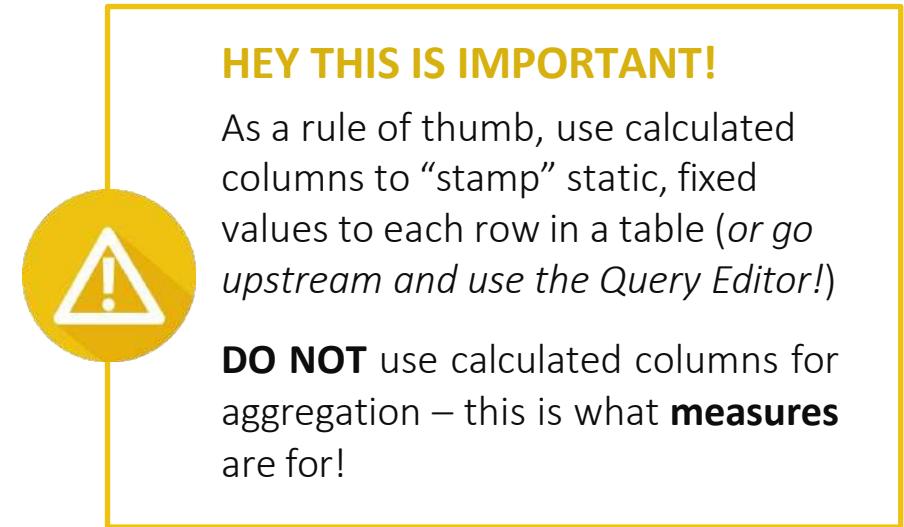
```
1 Bike Returns =
2 CALCULATE(
3     [Total Returns],                                     // Counting total returns
4     'Product Categories Lookup'[Category Name] = "Bikes" // filtered for bikes only
5 )
```



CALCULATED COLUMNS

Calculated columns allow you to add new, formula-based columns to tables in a model

- Calculated columns refer to **entire tables or columns** (*no A1-style cell references*)
- Calculated columns **generate values for each row**, which are visible within tables in the Data view
- Calculated columns understand **row context**; they're great for defining properties based on information in each row, but generally useless for aggregation (*sum, count, etc.*)



PRO TIP:

Calculated columns are typically used for **filtering & grouping** data, rather than creating aggregate numerical values



EXAMPLE: CALCULATED COLUMNS

Email Address	Annual Income	Total Children	Education Level	Parent
emma32@adventure-works.com	70000	5	Bachelors	Yes
barry20@adventure-works.com	40000	5	High School	Yes
martha13@adventure-works.com	70000	5	High School	Yes
tamara16@adventure-works.com	40000	5	High School	Yes
gerald21@adventure-works.com	130000	5	Bachelors	Yes
alexa8@adventure-works.com	40000	5	High School	Yes
jack53@adventure-works.com	70000	5	Graduate Degree	Yes
ricky1@adventure-works.com	100000	5	Bachelors	Yes
keith4@adventure-works.com	70000	5	Partial College	Yes
latoya19@adventure-works.com	70000	5	Bachelors	Yes

Here we're using an aggregation function (SUM) to calculate a new column named **TotalQuantity**

- Since this is an aggregation function, **the same grand total** is returned in *every row* of the table
- This is **not a valid use** of calculated columns; these values are statically “stamped” onto the table and can’t be filtered, sliced, etc.

Stock Date	Order Number	Product Key	Customer Key	Territory Key	Order Line Item	Order Quantity	Index	TotalQuantity
6/3/2020	SO46718	360	12570	9	1	1	1205	84174
4/22/2020	SO46736	360	12341	9	1	1	1228	84174
5/5/2020	SO46776	360	12356	9	1	1	1267	84174
6/22/2020	SO46808	360	12347	9	1	1	1299	84174
5/11/2020	SO46826	360	12575	9	1	1	1314	84174
4/21/2020	SO47075	360	12685	9	1	1	1421	84174
5/1/2020	SO47098	360	12657	9	1	1	1445	84174
4/21/2020	SO47149	360	12669	9	1	1	1495	84174
6/4/2020	SO47212	360	12580	9	1	1	1550	84174
6/25/2020	SO47302	360	12670	9	1	1	1649	84174
8/12/2020	SO47328	360	12681	9	1	1	1669	84174
8/19/2020	SO47346	360	12585	9	1	1	1690	84174
6/12/2020	SO47744	360	12989	9	1	1	1900	84174
7/28/2020	SO47745	360	12998	9	1	1	1904	84174
8/22/2020	SO47753	360	13020	9	1	1	1912	84174

In this case we've added a **calculated column** named **Parent**, which equals “**Yes**” if the [Total Children] field is greater than 0, and “**No**” otherwise

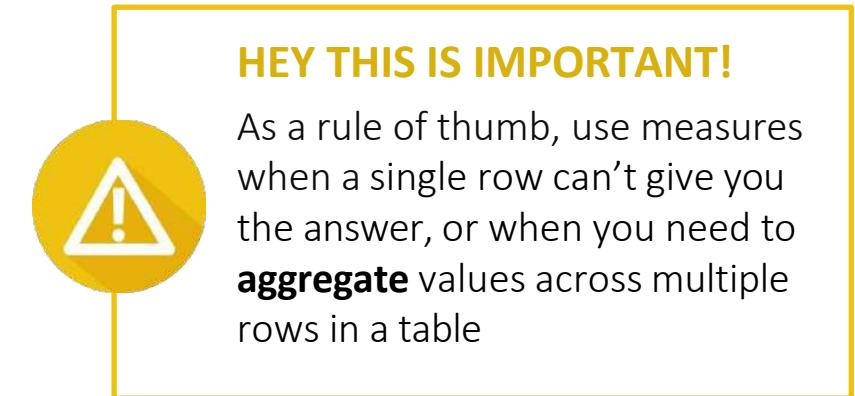
- Since calculated columns understand **row context**, a new value is calculated in each row based on the value in the [Total Children] column
- This is a **valid use** of calculated columns; it creates a new row “property” that we can use to filter or segment any related data within the model



DAX MEASURES

Measures are DAX formulas used to generate new calculated values

- Like calculated columns, measures reference **entire tables** or **columns** (*no A1-style cell references*)
- Unlike calculated columns, **measures** aren't visible within tables; they can only be "seen" within a visualization like a chart or matrix (*similar to a calculated field in a PivotTable*)
- Measures evaluate based on **filter context**, which means they recalculate when the fields or filters around them change



PRO TIP:

Use measures to create **numerical, calculated values** that can be analyzed in the "**values**" field of a report visual



IMPLICIT VS. EXPLICIT MEASURES

The screenshot shows the 'Build a visual' interface in Power BI. On the left, under 'Y-axis', there is a box containing 'Sum of Order ...'. This box is highlighted with a yellow border. Below it is a button labeled '+Add data'. A yellow arrow points from this button to the 'Order Quantity' item in the 'Select data' dialog box on the right. The 'Select data' dialog lists various data items, with 'Order Quantity' being the one selected.

Example of an **implicit measure**

Implicit measures are created when you drag raw numerical fields into a report visual and manually select an aggregation mode (*Sum, Average, Min, Max, Count, etc.*)

Explicit measures are created when you actually write a DAX formula and define a new measure that can be used within the model

HEY THIS IS IMPORTANT!

Implicit measures are only accessible within the **specific visualization** in which they were created, and cannot be referenced elsewhere

Explicit measures can be used **anywhere in the report**, and referenced by other DAX calculations to create “measure trees”



QUICK MEASURES

Quick measures automatically create formulas based on pre-built templates or natural language prompts

Quick measure

Select a calculation to create a measure or describe the measure you need and we'll generate suggestions in DAX, which you can customize later.

Calculations Suggestions

Weighted average per category

Calculate a weighted average of the base value for each category. Multiply the value by weight for each category, sum total, and then divide by the sum total of the weight. [Learn more](#)

Base value ○

Sum of Order Quantity

Weight ○

Add data

Category ○

Add data

Quick measure **calculations** can be used to build measures using **predefined templates** (*weighted averages, percent difference, time intelligence, etc.*)

Quick measure

Select a calculation to create a measure or describe the measure you need and we'll generate suggestions in DAX, which you can customize later.

Calculations **Suggestions**

Sum of quantity sold by calendar lookup year

Generate

Suggested measures

Total quantity sold per year

Quick measure **suggestions** can be used to find suggested measures based on **natural language queries** (i.e. “sum of quantity sold by calendar year”)

PRO TIP:



Quick measures can be a great learning tool for beginners or for building more complex formulas but use them with caution; **mastering DAX requires a deep understanding of the underlying theory!**



RECAP: CALCULATED COLUMNS VS. MEASURES

CALCULATED COLUMNS

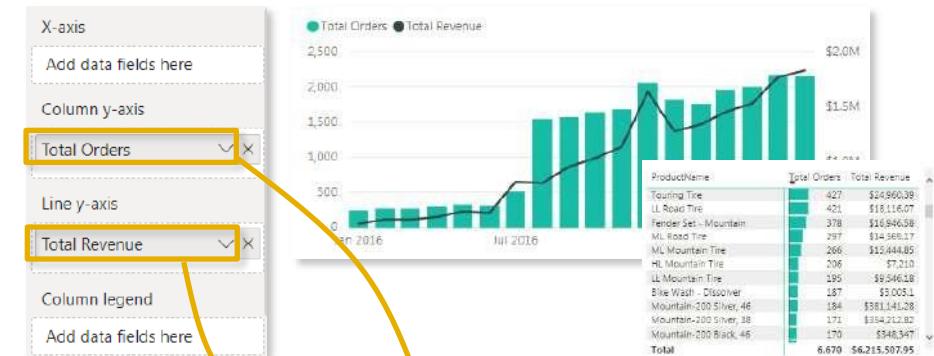
- Values are calculated based on information from each row of a table (**row context**)
- Appends static values to each row in a table and stores them in the model (*which increases file size*)
- Recalculate on data source refresh or when changes are made to component columns
- Primarily used for **filtering** data in reports

Birth Date	Marital Status	Email Address	Annual Income	Total Children	Education Level	Parent
9/3/1943	M	emma32@adventure-works.com	70000	5	Bachelors	Yes
9/14/1967	M	barry20@adventure-works.com	40000	5	High School	Yes
8/5/1945	M	martha13@adventure-works.com	70000	5	High School	Yes
6/4/1946	S	tamara16@adventure-works.com	40000	5	High School	Yes
10/16/1970	S	gerald21@adventure-works.com	130000	5	Bachelors	Yes
5/10/1945	M	alexas8@adventure-works.com	40000	5	High School	Yes
9/24/1938	M	jack53@adventure-works.com	70000	5	Graduate Degree	Yes
7/21/1959	S	ricky1@adventure-works.com	100000	5	Bachelors	Yes
1/6/1962	M	keith4@adventure-works.com	70000	5	Partial College	Yes
8/15/1962	M	latoya19@adventure-works.com	70000	5	Bachelors	Yes
1/26/1967	S	mical11@adventure-works.com	70000	5	Bachelor	Yes
3/8/1946	M	mindy22@adventure-works.com	80000	5	Partial College	Yes
6/11/1960	M	teresa8@adventure-works.com	70000	5	Partial College	Yes

Calculated columns “live” in **tables**

MEASURES

- Values are calculated based on information from any filters in the report (**filter context**)
- Does not create new data in the tables themselves (*doesn’t increase file size*)
- Recalculate in response to any change to filters within the report
- Primarily used for **aggregating values** in report visuals



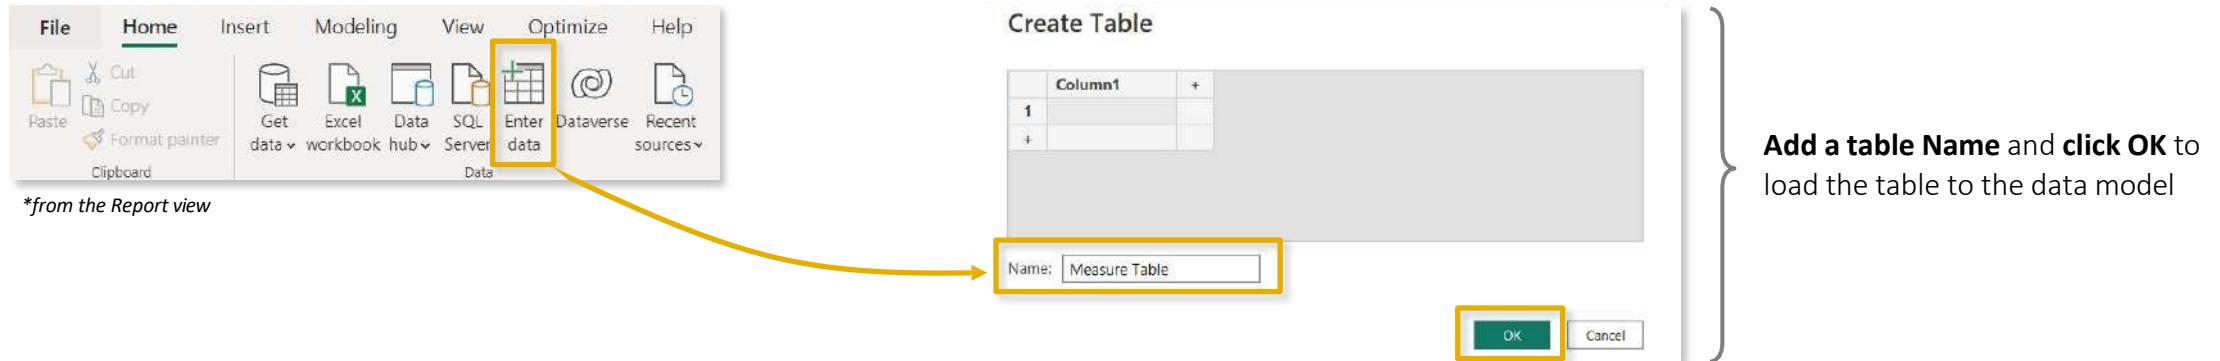
Measures “live” in **visuals**



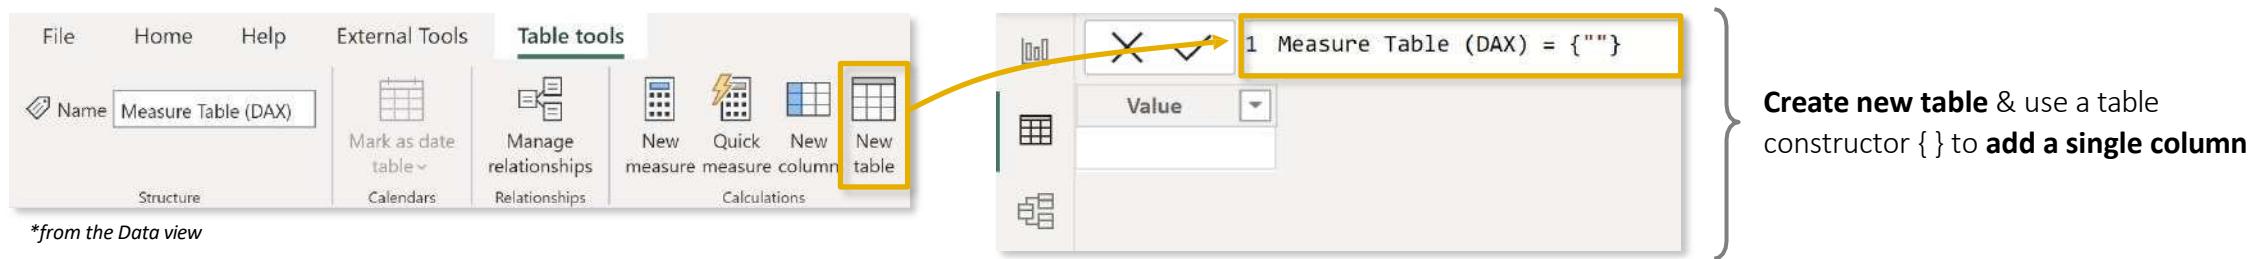
PRO TIP: MEASURE TABLES

It's a common best practice to **create a dedicated table to store your measures**; this will help you stay organized, find measures quickly, and allow you to group related measures into folders

Option 1: Enter Data into Power Query (loads the table to the data model – table is visible in Power Query)



Option 2: Create a calculated table using DAX directly in the model (table is not visible in Power Query)





FILTER CONTEXT

Measures are evaluated based on **filter context**, which means that they recalculate whenever the fields or filters around them change

Top 10 Products	Orders	Revenue	Return %
Water Bottle - 30 oz.	3,983	\$39,755	1.95%
Patch Kit/8 Patches	2,952	\$13,506	1.61%
Mountain Tire Tube	2,846	\$28,333	1.64%
Road Tire Tube	2,173	\$17,265	1.55%
Sport-100 Helmet, Red	2,099	\$73,444	3.33%
AWC Logo Cap	2,062	\$35,865	1.11%
Sport-100 Helmet, Blue	1,995	\$67,112	3.31%
Fender Set - Mountain	1,975	\$87,041	1.36%
Sport-100 Helmet, Black	1,940	\$65,262	2.68%
Mountain Bottle Cage	1,896	\$38,062	2.02%
Total	15,587	\$465,644	1.85%

For this value in the matrix (2,846), the **Orders** measure is calculated based on the following filter context: *Products[Product Name] = “Mountain Tire Tube”*

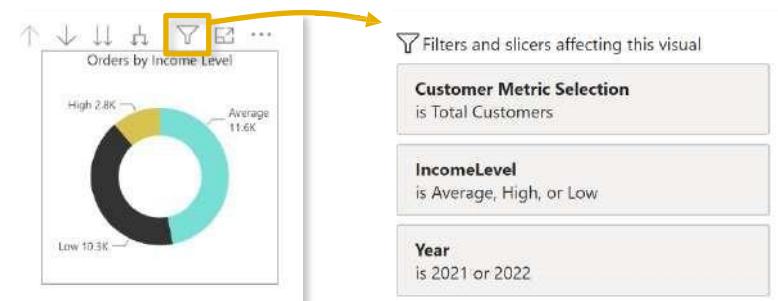
- This allows the measure to return the total order quantity for each product specifically (or whatever context the row and column labels dictate – years, countries, categories, customer names, etc.)

This total (15,587) does **NOT** calculate by summing the values above; it evaluates as an independent measure with **no filter context** applied

- IMPORTANT:** Every measure value in a report evaluates **independently** (like an island) and calculates based on its own filter context



PRO TIP: Clicking the **filter icon** will show you the filters currently applied to a selected visual



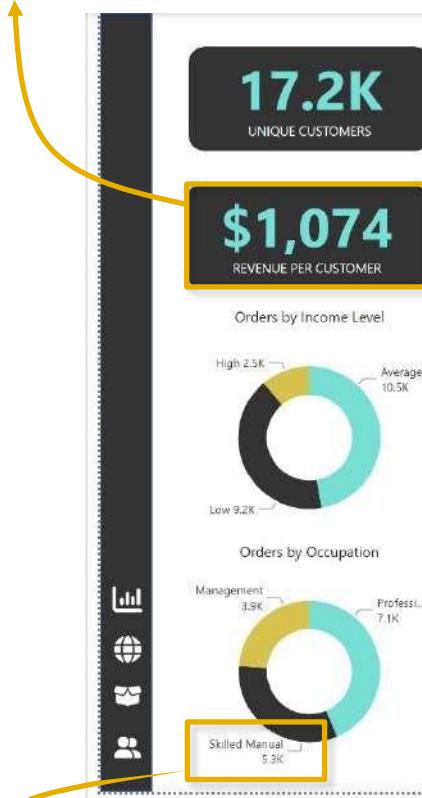


EXAMPLE: FILTER CONTEXT

MEASURE: Revenue Per Customer

FILTER CONTEXT:

- *Calendar[Year]* = 2021 or 2022



MEASURE: Total Orders

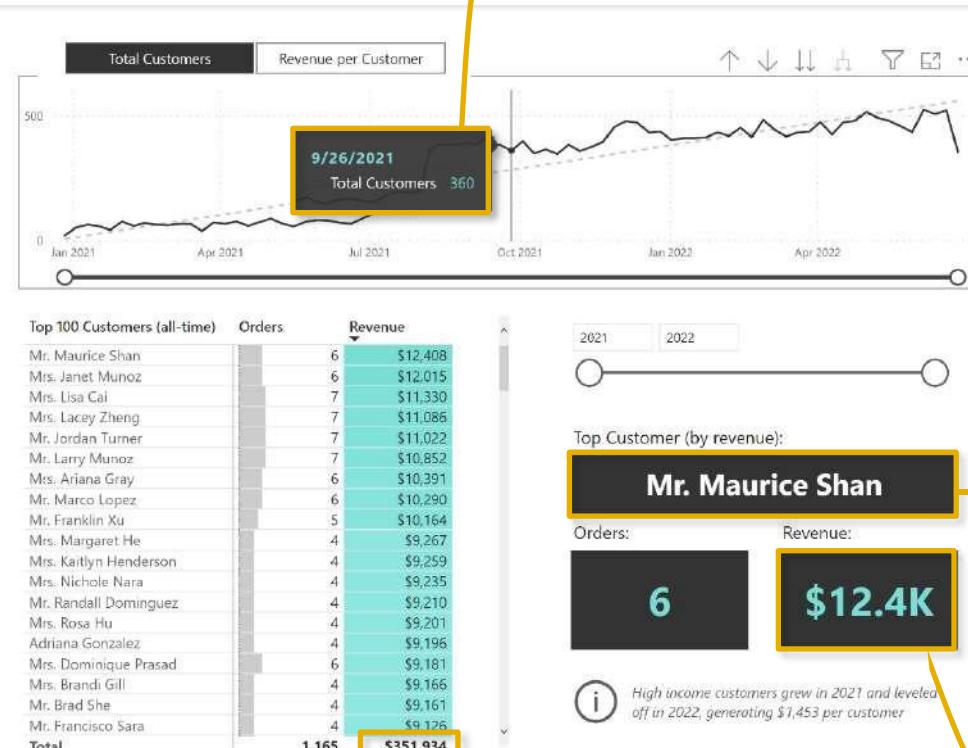
FILTER CONTEXT:

- *Calendar[Year]* = 2021 or 2022
- *Customers[Occupation]* = Skilled Manual

MEASURE: Total Customers

FILTER CONTEXT:

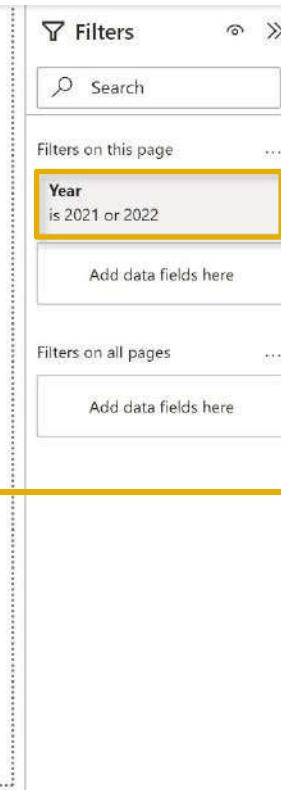
- *Calendar[Date]* = September 26, 2021



MEASURE: Total Revenue

FILTER CONTEXT:

- *Calendar[Year]* = 2021 or 2022
- *Customer[Full Name]* = Top 100 by Total Orders



MEASURE: Total Revenue

FILTER CONTEXT:

- *Calendar[Year]* = 2021 or 2022
- *Customer[Full Name]* = Mr. Maurice Shan

This is a **page-level filter**, which impacts **ALL** visuals on this report page (*more on this later!*)

COLUMN: Customer Full Name

FILTER CONTEXT:

- *Calendar[Year]* = 2021 or 2022
- *Customer[Full Name]* = Top 1 by Total Revenue



STEP-BY-STEP MEASURE CALCULATION

Product Color	Quantity Sold
Black	10,590
Red	4,011
Yellow	4,638

How exactly is this measure value calculated?

- **NOTE:** This all happens *instantly* behind the scenes, every time the filter context changes

STEP 1

Filter context is detected & applied



Product Color	Quantity Sold
Black	10,590
Red	4,011
Yellow	4,638

'Product Lookup'[Product Color] = "Black"

STEP 2

Filters flow “downstream” to related tables



Product Key	Product Color
P00000001	Black
P00000002	Red
P00000003	Yellow

Order Date	Customer ID	Product Key	Order Quantity
Wednesday, June 3, 2020	1047118	Black	360
Wednesday, June 3, 2020	1047118	Black	360
Wednesday, June 3, 2020	1047118	Black	360
Wednesday, June 3, 2020	1047118	Black	360

Product Key	Product Name
Black	Smartphone X
Red	Smartphone Y
Yellow	Smartphone Z

STEP 3

Measure evaluates against the filtered table



```

1 Quantity Sold =
2 SUM(
3   |   'Sales Data'[Order Quantity]
4 )

```

Product Key	Product Name
Black	Smartphone X
Red	Smartphone Y
Yellow	Smartphone Z

Return Date	Customer ID	Product Key	Return Quantity
Wednesday, February 12, 2020	312	Black	1
Wednesday, February 12, 2020	312	Black	1
Wednesday, February 12, 2020	312	Black	1
Wednesday, February 12, 2020	312	Black	1

Sum of values in the **Order Quantity** column of the **Sales Data** table, filtered to rows where the product color is "Black"

= 10,590



DAX SYNTAX

MEASURE NAME

- Measures are always surrounded by brackets (i.e. **[Total Quantity]**) when referenced in formulas, so spaces are OK

Total Quantity: = **SUM(Transactions[quantity])**

FUNCTION NAME

- Calculated columns don't always use functions, but measures do:
 - In a **Calculated Column**, **=Transactions[quantity]** returns the value from the quantity column in each row (*since it evaluates one row at a time*)
 - In a **Measure**, **=Transactions[quantity]** will return an **error** since Power BI doesn't know how to translate that as a single value – you need some sort of aggregation

Referenced
TABLE NAME

Referenced
COLUMN NAME

This is a “**fully qualified**” column, since it’s preceded by the table name.

NOTE: Table names with spaces must be surrounded by **single quotes**:

- Without a space: **Transactions[quantity]**
- With a space: **'Transactions Table'[quantity]**

PRO TIP:

Column references use fully qualified names (i.e. **'Table'[Column]**)

Measure references just use the measure name (i.e. **[Measure]**) and can be called by typing an open square bracket “ [”





DAX OPERATORS

Arithmetic Operator	Meaning	Example
+	Addition	$2 + 7$
-	Subtraction	$5 - 3$
*	Multiplication	$2 * 6$
/	Division	$4 / 2$
\wedge	Exponent	$2 \wedge 5$

Pay attention to these!

Comparison Operator	Meaning	Example
=	Equal to	[City] = "Boston"
>	Greater than	[Quantity] > 10
<	Less than	[Quantity] < 10
\geq	Greater than or equal to	[Unit Price] \geq 2.5
\leq	Less than or equal to	[Unit Price] \leq 2.5
\neq	Not equal to	[Country] \neq "Mexico"

Text/Logical Operator	Meaning	Example
&	Concatenates two values to produce one text string	[City] & " " & [State]
&&	Create an AND condition between two logical expressions	([State] = "MA") && ([Quantity] > 10)
(double pipe)	Create an OR condition between two logical expressions	([State] = "MA") ([State] = "CT")
IN	Creates a logical OR condition based on a given list (using curly brackets)	'Store Lookup'[State] IN { "MA", "CT", "NY" }



COMMON FUNCTION CATEGORIES

MATH & STATS Functions	LOGICAL Functions	TEXT Functions	FILTER Functions	TABLE Functions	DATE & TIME Functions	RELATIONSHIP Functions
<p>Functions used for aggregation or iterative, row-level calculations</p> <p>Common Examples:</p> <ul style="list-style-type: none"> • SUM • AVERAGE • MAX/MIN • DIVIDE • COUNT/COUNTA • COUNTROWS • DISTINCTCOUNT <p>Iterator Functions:</p> <ul style="list-style-type: none"> • SUMX • AVERAGEX • MAXX/MINX • RANKX • COUNTX 	<p>Functions that use conditional expressions (IF/THEN statements)</p> <p>Common Examples:</p> <ul style="list-style-type: none"> • IF • IFERROR • AND • OR • NOT • SWITCH • TRUE • FALSE 	<p>Functions used to manipulate text strings or value formats</p> <p>Common Examples:</p> <ul style="list-style-type: none"> • CONCATENATE • COMBINEVALUES • FORMAT • LEFT/MID/RIGHT • UPPER/LOWER • LEN • SEARCH/FIND • REPLACE • SUBSTITUTE • TRIM 	<p>Functions used to manipulate table and filter contexts</p> <p>Common Examples:</p> <ul style="list-style-type: none"> • CALCULATE • FILTER • ALL • ALLEXCEPT • ALLSELECTED • KEEPFILTERS • REMOVEFILTERS • SELECTEDVALUE 	<p>Functions that create or manipulate tables and output tables vs. scalar values</p> <p>Common Examples:</p> <ul style="list-style-type: none"> • SUMMARIZE • ADDCOLUMNS • GENERATESERIES • DISTINCT • VALUES • UNION • INTERSECT • TOPN 	<p>Functions used to manipulate date & time values or handle time intelligence calculations</p> <p>Common Examples:</p> <ul style="list-style-type: none"> • DATE • DATEDIFF • YEARFRAC • YEAR/MONTH • DAY/HOUR • TODAY/NOW • WEEKDAY • WEEKNUM • NETWORKDAYS <p>Time Intelligence:</p> <ul style="list-style-type: none"> • DATESYTD • DATESMTD • DATEADD • DATESBETWEEN 	<p>Functions used to manage & modify table relationships</p> <p>Common Examples:</p> <ul style="list-style-type: none"> • RELATED • RELATEDTABLE • CROSSFILTER • USERELATIONSHIP

***Note:** This is NOT a comprehensive list. DAX contains more than 250 different functions!



BASIC MATH & STATS FUNCTIONS

SUM

Evaluates the sum of a column

=SUM(Column**Name**)

AVERAGE

Returns the average (arithmetic mean) of all the numbers in a column

=AVERAGE(Column**Name**)

MAX

Returns the largest value in a column or between two scalar expressions

=MAX(Column**NameOrScalar1**, [Scalar2])

MIN

Returns the smallest value in a column or between two scalar expressions

=MIN(Column**NameOrScalar1**, [Scalar2])

DIVIDE

Performs division and returns the alternate result (or blank) if DIV/0

=DIVIDE(Numerator, Denominator, [AlternateResult])



COUNTING FUNCTIONS

COUNT

Counts the number of non-empty cells in a column
(excluding Boolean values)

=**COUNT**(ColumnName)

COUNTA

Counts the number of non-empty cells in a column
(including Boolean values)

=**COUNTA**(ColumnName)

DISTINCTCOUNT

Counts the number of distinct values in a column

=**DISTINCTCOUNT**(ColumnName)

COUNTROWS

Counts the number of rows in the specified table,
or a table defined by an expression

=**COUNTROWS**([Table])



ASSIGNMENT: MATH & STATS

  NEW MESSAGE

From: **Dianne A. Xu** (Senior Analyst)
Subject: Help with a few measures

Hey there, excited to start working with you!
I'll need to pull some high-level metrics from our model to share with leadership, and I could use some help with the calculations.
For now, could you please create one measure to calculate the total number of distinct customers, and a second measure that we can use to calculate return rate (quantity returned / quantity sold)? Thank you!

-Dianne

Reply Forward

Key Objectives

1. Create a measure named **Total Customers**, to calculate the number of distinct AdventureWorks customers who made a transaction
2. Create a measure named **Return Rate**, defined as quantity returned divided by quantity sold



SOLUTION: MATH & STATS

  NEW MESSAGE

From: **Dianne A. Xu** (Senior Analyst)
Subject: Help with a few measures

Hey there, excited to start working with you!
I'll need to pull some high-level metrics from our model to share with leadership, and I could use some help with the calculations.
For now, could you please create one measure to calculate the total number of distinct customers, and a second measure that we can use to calculate return rate (quantity returned / quantity sold)? Thank you!

-Dianne

Reply Forward

Solution Preview

```
1 Total Customers =  
2 DISTINCTCOUNT(  
3     'Sales Data'[Customer Key]  
4 )
```

```
1 Return Rate =  
2 DIVIDE(  
3     [Quantity Returned],  
4     [Quantity Sold],  
5     "No Sales"  
6 )
```



BASIC LOGICAL FUNCTIONS

IF

Checks if a given condition is met and returns one value if the condition is TRUE, and another if the condition is FALSE

=**IF**(LogicalTest, ResultIfTrue, [ResultIfFalse])

IFERROR

Evaluates an expression and returns a specified value if it returns an error, otherwise returns the expression itself

=**IFERROR**(Value, ValueIfError)

SWITCH

Evaluates an expression against a list of values and returns one of multiple possible expressions

=**SWITCH**(Expression, Value1, Result1, ..., [Else])

AND

Checks whether both arguments are TRUE to return TRUE, otherwise returns FALSE

=**AND**(Logical1, Logical2)

OR

Checks whether any argument is TRUE to return TRUE, otherwise returns FALSE

=**OR**(Logical1, Logical2)

Note: Use the **&&** and **||** operators to include more than two conditions



SWITCH

SWITCH

Evaluates an expression against a list of values and returns one of multiple possible expressions

=**SWITCH**(Expression, Value1, Result1, ..., [Else])

Any **DAX expression** that returns a single scalar value, evaluated multiples times

Examples:

- *Calendar[Month ID]*
- *'Product Lookup'[category]*

List of **values** produced by the expression, each paired with a result to return for rows/cases that match

Examples:

```
=SWITCH( Calendar[Month ID],  
        1, "January",  
        2, "February" )
```

Value returned if the expression doesn't match any value argument



PRO TIP

SWITCH(TRUE) is a common DAX pattern to replace multiple nested IF statements

ASSIGNMENT: LOGICAL FUNCTIONS



 NEW MESSAGE

From: **Dianne A. Xu** (Senior Analyst)

Subject: **Customer segmentation fields**

Hey there!

Ethan has been working with the DS team on a customer segmentation analysis, and came back to us with a few requests.

Could you please add some new columns in our customer table to identify “priority” customers, segment customers based on income level, and group some of the education categories?

I've attached the logic to use, but reach out with any questions!

-Dianne

Reply Forward

Key Objectives

1. Create a calculated column in the Customer Lookup table named **Customer Priority**:
 - If the customer is a parent and has an annual income > \$100,000, Customer Priority = **Priority**
 - Otherwise, Customer Priority = **Standard**
2. Create a calculated column in the Customer Lookup table named **Income Level**:
 - If annual income is >= \$150,000, **Very High**
 - If annual income is >= \$100,000, **High**
 - If annual income is >= \$50,000, **Average**
 - Otherwise, Income Level = **Low**

ASSIGNMENT: LOGICAL FUNCTIONS



 NEW MESSAGE

From: **Dianne A. Xu** (Senior Analyst)

Subject: **Customer segmentation fields**

Hey there!

Ethan has been working with the DS team on a customer segmentation analysis, and came back to us with a few requests.

Could you please add some new columns in our customer table to identify “priority” customers, segment customers based on income level, and group some of the education categories?

I've attached the logic to use, but reach out with any questions!

-Dianne

Reply Forward

Key Objectives

BONUS: Use a SWITCH function* to create another column named **Education Category**:

- If EducationLevel is High School or Partial High School, Education Category = **High School**
- If EducationLevel is Bachelors or Partial College, Education Category = **Undergrad**
- If EducationLevel is Graduate Degree, Education Category = **Graduate**

*You can use the “data groups” tool to do this too!

SOLUTION: LOGICAL FUNCTIONS



 NEW MESSAGE

From: **Dianne A. Xu** (Senior Analyst)

Subject: **Customer segmentation fields**

Hey there!

Ethan has been working with the DS team on a customer segmentation analysis, and came back to us with a few requests.

Could you please add some new columns in our customer table to identify "priority" customers, segment customers based on income level, and group some of the education categories?

I've attached the logic to use, but reach out with any questions!

-Dianne

Reply Forward

Solution Preview

```
1 Customer Priority =  
2 IF(  
3     'Customer Lookup'[AnnualIncome] > 100000 &&  
4     'Customer Lookup'[Is Parent?] = "Yes",  
5     "Priority",  
6     "Standard"  
7 )
```

```
1 Income Level =  
2 IF('Customer Lookup'[AnnualIncome] >= 150000, "Very High",  
3 IF('Customer Lookup'[AnnualIncome] >= 100000, "High",  
4 IF('Customer Lookup'[AnnualIncome] >= 50000, "Average",  
5 "Low")))
```

```
1 Education Category =  
2 SWITCH('Customer Lookup'[EducationLevel],  
3 "High School", "High School",  
4 "Partial High School", "High School",  
5 "Bachelors", "Undergrad",  
6 "Partial College", "Undergrad",  
7 "Graduate Degree", "Graduate")
```



TEXT FUNCTIONS

LEN

Returns the number of characters in a string

=**LEN**(Text)

Note: Use the & operator as a shortcut,
or to combine more than two strings

CONCATENATE

Joins two text strings into one

=**CONCATENATE**(Text1, Text2)

UPPER/LOWER

Converts a string to upper or lower case

=**UPPER/LOWER** (Text)

LEFT/RIGHT/MID

Returns a number of characters from the
start/middle/end of a text string

=**LEFT/RIGHT**(Text, [NumChars])

=**MID**(Text, StartPosition, NumChars)

SUBSTITUTE

Replaces an instance of existing text with
new text in a string

=**SUBSTITUTE**(Text, OldText, NewText,
[InstanceNumber])

SEARCH

Returns the position where a specified string
or character is found, reading left to right

=**SEARCH**(FindText, WithinText,
[StartPosition], [NotFoundValue])



ASSIGNMENT: TEXT

  NEW MESSAGE

From: **Dianne A. Xu** (Senior Analyst)
Subject: **Couple random requests**

Good morning!

Hoping you can help with a couple quick updates to the model:

- 1) Ethan wants to make the month abbreviations ALL CAPS to make them more readable in our reports.
- 2) The product team asked us to break out the SKU category into its own field, which we can define as any characters before the first hyphen (“-”) in the ProductSKU column.

Thanks, reach out with any questions!

Reply Forward

Key Objectives

1. Update the **Month Short** column in the Calendar Lookup table to extract and capitalize the first 3 characters of the month name
2. Create a new column in the Product Lookup table named **SKU Category**, to return any number of characters before the first hyphen in the ProductSKU column



SOLUTION: TEXT

 NEW MESSAGE

From: **Dianne A. Xu** (Senior Analyst)

Subject: **Couple random requests**

Good morning!

Hoping you can help with a couple quick updates to the model:

- 1) Ethan wants to make the month abbreviations ALL CAPS to make them more readable in our reports.
- 2) The product team asked us to break out the SKU category into its own field, which we can define as any characters before the first hyphen (“-”) in the ProductSKU column.

Thanks, reach out with any questions!

Reply Forward

Solution Preview

```
1 Month Short =
2 UPPER(
3   LEFT(
4     'Calendar Lookup'[Month Name],
5     3
6   )
7 )
```

```
1 SKU Category =
2 LEFT(
3   'Product Lookup'[Product SKU],
4   SEARCH(
5     "-",
6     'Product Lookup'[Product SKU]
7   )
8   -1
9 )
```



BASIC DATE & TIME FUNCTIONS

TODAY/NOW

Returns the current date or exact time

=**TODAY/NOW()**

DAY/MONTH/YEAR

Returns the day of the month (1-31), month of the year (1-12), or year of a given date

=**DAY/MONTH/YEAR**(Date)

HOUR/MINUTE/SECOND

Returns the hour (0-23), minute (0-59), or second (0-59) of a given datetime value

=**HOUR/MINUTE/SECOND**(Datetime)

WEEKDAY/WEEKNUM

Returns a weekday number from 1 (Sunday) to 7 (Saturday), or the week # of the year

=**WEEKDAY/WEEKNUM**(Date, [ReturnType])

EOMONTH

Returns the date of the last day of the month, +/- a specified number of months

=**EOMONTH**(StartDate, Months)

DATEDIFF

Returns the difference between two dates, based on a given interval (day, hour, year, etc.)

=**DATEDIFF**(Date1, Date2, Interval)



ASSIGNMENT: DATE & TIME

  NEW MESSAGE

From: **Dianne A. Xu** (Senior Analyst)

Subject: **Customer birth years**

Hey there, super easy one for you.

The customer segmentation project got me wondering if there are any interesting patterns or insights based on customer age.

Could you please add a field in our customer table to extract only the year from the birthdate field?

Thanks!
-Dianne

Reply Forward

Key Objectives

1. Create a new column in the Customer Lookup table named **Birth Year**, to extract only the year from the BirthDate column



SOLUTION: DATE & TIME

  NEW MESSAGE

From: **Dianne A. Xu** (Senior Analyst)
Subject: Customer birth years

Hey there, super easy one for you.

The customer segmentation project got me wondering if there are any interesting patterns or insights based on customer age.

Could you please add a field in our customer table to extract only the year from the birthdate field?

Thanks!
-Dianne

Reply Forward

Solution Preview

```
1 Birth Year =  
2 YEAR(  
3     'Customer Lookup'[BirthDate]  
4 )
```



RELATED

RELATED()

Returns related values in each row of a table based on relationships with other tables

=RELATED(Column**Name**)

The **column** from a related table containing the values you want to retrieve

Examples:

- 'Product Lookup'[Product Name]
- 'Territory Lookup'[Country]



HEY THIS IS IMPORTANT!

RELATED works like a **VLOOKUP** function in Excel – it uses the relationship between tables (*defined by primary and foreign keys*) to pull values from one table into a new column of another.

Since this function requires row context, it can only be used as a **calculated column** or as part of an **iterator function** that cycles through all rows in a table (*FILTER, SUMX, MAXX, etc.*)



PRO TIP:

Instead of using RELATED to create extra columns (which increases file size), **nest it within measures like FILTER or SUMX**



CALCULATE

CALCULATE()

Evaluates an expression in a context that is modified by filters

=CALCULATE(Expression, [Filter1], [Filter2],...)

Name of an **existing measure** or a **DAX formula** for a valid measure

Examples:

- [Total Orders]
- SUM('Returns Data'[Return Quantity])

A Boolean (True/False) expression or a table expression that defines a filter

Note: these require fixed values or aggregation functions that return a scalar value (you cannot create filters based on measures)

Examples:

- 'Territory Lookup'[Country] = "USA"
- Calendar[Year] <> MAX(Calendar[Year])



PRO TIP:

Think of CALCULATE as a **filter modifier**; it allows you to overrule existing report filters and “force” new filter context



EXAMPLE: CALCULATE

X ✓ 1 Red Sales = CALCULATE([Quantity Sold], 'Product Lookup'[Product Color] = "Red")

Here we've defined a new measure named **Red Sales**, which evaluates the **Quantity Sold** measure under a filter context where the product color is "**Red**"

Product Color	Quantity Sold	Red Sales
Black	10,590	4,011
Multi	5,756	4,011
Red	4,011	4,011
Silver	3,257	4,011
Total	23,614	4,011

Note how we see *the same repeated values* for each product color, and even the total!



HEY THIS IS IMPORTANT!

The **CALCULATE** function **modifies and overrules any competing filter context!**

In this matrix, the "Black" row has competing filter context: Product Color = **Black** (from the row label) and Product Color= "**Red**" (from the CALCULATE function)

Both can't be true at the same time, so the "**Red**" filter from CALCULATE takes priority



EXAMPLE: CALCULATE

CALCULATE

Filters are modified by CALCULATE

[Product Color] = "Red"

If the measure being evaluated contains a **CALCULATE** function,
filter context is *overwritten* between **Step 1 & Step 2**

STEP 1

Filter context is detected
& applied



Product Color	Quantity Sold	Red Sales
Black	10,590	4,011
Red	4,011	4,011
Silver	3,257	4,011

'Product Lookup'[Product Color] = "Black"

Red

Product Color	Count
Black	10590
Red	4011

Black

STEP 2

Filters flow "downstream"
to related tables



Product Color	Count
Black	1
Red	1

Red

Order Date	Order Number	Product Key
Wednesday, April 8, 2020	504738	360
Wednesday, April 15, 2020	504739	360
Wednesday, April 22, 2020	504740	360
Wednesday, April 29, 2020	504741	360
Wednesday, May 6, 2020	504742	360
Wednesday, May 13, 2020	504743	360
Wednesday, May 20, 2020	504744	360
Wednesday, May 27, 2020	504745	360
Wednesday, June 3, 2020	504746	360
Wednesday, June 10, 2020	504747	360
Wednesday, June 17, 2020	504748	360
Wednesday, June 24, 2020	504749	360
Wednesday, July 1, 2020	504750	360
Wednesday, July 8, 2020	504751	360
Wednesday, July 15, 2020	504752	360
Wednesday, July 22, 2020	504753	360
Wednesday, July 29, 2020	504754	360
Wednesday, August 5, 2020	504755	360
Wednesday, August 12, 2020	504756	360
Wednesday, August 19, 2020	504757	360
Wednesday, August 26, 2020	504758	360
Wednesday, September 2, 2020	504759	360
Wednesday, September 9, 2020	504760	360

Red

Return Date	Return Reason	Product Key	Returns Quantity
Friday, March 13, 2020		3	312
Sunday, March 15, 2020		3	313
Monday, March 16, 2020		3	314
Sunday, March 22, 2020		3	315
Tuesday, April 7, 2020		3	316
Tuesday, April 7, 2020		3	317
Thursday, April 9, 2020		3	318
Friday, April 10, 2020		3	319
Wednesday, April 15, 2020		3	320
Saturday, April 18, 2020		3	321
Wednesday, April 22, 2020		3	322
Wednesday, April 22, 2020		3	323
Wednesday, April 22, 2020		3	324
Wednesday, April 22, 2020		3	325

Red

STEP 3

Measure evaluates against the
filtered table

$$1 \text{ Quantity Sold} = \\ 2 \text{ SUM('Sales Data'[Order Quantity]) }$$

Sum of the Order **Quantity** column in the
Sales Data table, filtered to rows where
the product color is "**Red**"

= 4,011



DAX MEASURE TOTALS

Measure totals may seem incorrect or inconsistent depending on how they are calculated, because they **don't simply add up the visible values in the report**



Total Returns look right, but
shouldn't Total Orders be **37,888??**
-Anonymous confused man

Category Name	Total Returns	Total Orders
Accessories	1,115	16,983
Bikes	427	13,929
Clothing	267	6,976
Total	1,809	25,164



PRO TIP:

Understand EXACTLY how your measures calculate and **what they are designed to measure**

```
1 Total Orders =  
2 DISTINCTCOUNT(  
3 | 'Sales Data'[Order Number]  
4 )
```

[Total Orders] counts **distinct orders** in the Sales Data table

Order Date	Stock Date	Order Number	Product Key
Thursday, June 30, 2022	Thursday, April 07, 2022	S074140	568
Thursday, June 30, 2022	Friday, March 04, 2022	S074140	477
Thursday, June 30, 2022	Monday, May 30, 2022	S074140	223
Thursday, June 30, 2022	Friday, April 29, 2022	S074141	604
Thursday, June 30, 2022	Wednesday, May 04, 2022	S074141	471
Thursday, June 30, 2022	Monday, May 30, 2022	S074142	383
Thursday, June 30, 2022	Friday, March 18, 2022	S074142	490
Thursday, June 30, 2022	Tuesday, March 15, 2022	S074143	479
Thursday, June 30, 2022	Friday, April 08, 2022	S074143	606
Thursday, June 30, 2022	Tuesday, March 22, 2022	S074143	477
Thursday, June 30, 2022	Thursday, June 02, 2022	S074143	462
Thursday, June 30, 2022	Monday, April 25, 2022	S074144	574
Thursday, June 30, 2022	Sunday, April 24, 2022	S074144	220
Thursday, June 30, 2022	Monday, March 14, 2022	S074145	561
Thursday, June 30, 2022	Tuesday, June 14, 2022	S074146	584
Thursday, June 30, 2022	Friday, March 18, 2022	S074147	605
Thursday, June 30, 2022	Sunday, May 29, 2022	S074147	538
Thursday, June 30, 2022	Thursday, March 24, 2022	S074147	490

Order **S074144** included **two products**: a bike and a helmet.

That counts as **1** distinct order for the Total and **1** distinct order for BOTH **Accessories & Bikes**

With no filter context, there are **25,164** total distinct orders



ASSIGNMENT: CALCULATE



NEW MESSAGE

From: **Dianne A. Xu** (Senior Analyst)

Subject: **URGENT: Bike returns**

Hey there,

Apparently George (our Product VP) has been speaking with some of the store managers, and they've raised concerns about the number of bike returns they are seeing recently.

Can you please create a measure to calculate total returns for bikes specifically, and let me know what you see? Volume alone won't tell the full story, so let's calculate the return *rate* for bikes as well, and see how it's trending before responding to George.

Need this ASAP – thank you!

Reply **Forward**

Key Objectives

1. Create a new measure named **Bike Returns** to calculate the total quantity of bikes returned
2. Create a matrix to show **Bike Returns** (values) by **Start of Month** (rows). What do you notice about the volume of bike returns over time?
3. Create a new measure named **Bike Sales** to calculate the total quantity of bikes sold, and add it to the matrix. What do you notice?
4. Create a new measure named **Bike Return Rate** using either CALCULATE or DIVIDE, and add it to the matrix
5. How would you respond to the Product VP's concerns about rising bike returns?



SOLUTION: CALCULATE

  NEW MESSAGE

From: **Dianne A. Xu** (Senior Analyst)
Subject: **URGENT: Bike returns**

Hey there,

Apparently George (our Product VP) has been speaking with some of the store managers, and they've raised concerns about the number of bike returns they are seeing recently.

Can you please create a measure to calculate total returns for bikes specifically, and let me know what you see? Volume alone won't tell the full story, so let's calculate the return *rate* for bikes as well, and see how it's trending before responding to George.

Need this ASAP – thank you!

Reply Forward

Solution Preview

```
1 Bike Returns =  
2 CALCULATE(  
3     [Total Returns],  
4     'Product Categories Lookup'[Category Name] = "Bikes"  
5 )
```

```
1 Bike Sales =  
2 CALCULATE(  
3     [Quantity Sold],  
4     'Product Categories Lookup'[Category Name] = "Bikes"  
5 )
```

```
1 Bike Return Rate =  
2 CALCULATE(  
3     [Return Rate],  
4     'Product Categories Lookup'[Category Name] = "Bikes"  
5 )
```

(Solution continued on next slide)



SOLUTION: CALCULATE

  NEW MESSAGE

From: **Dianne A. Xu** (Senior Analyst)

Subject: **URGENT: Bike returns**

Hey there,

Apparently George (our Product VP) has been speaking with some of the store managers, and they've raised concerns about the number of bike returns they are seeing recently.

Can you please create a measure to calculate total returns for bikes specifically, and let me know what you see? Volume alone won't tell the full story, so let's calculate the return *rate* for bikes as well, and see how it's trending before responding to George.

Need this ASAP – thank you!

Reply Forward

Solution Preview

6/1/2021	8	312	2.564%
7/1/2021	12	506	2.372%
8/1/2021	14	485	2.887%
9/1/2021	22	575	3.826%
10/1/2021	26	612	4.248%
11/1/2021	25	688	3.634%
12/1/2021	26	1038	2.505%
1/1/2022	14	766	1.828%
2/1/2022	22	806	2.730%
3/1/2022	27	888	3.041%
4/1/2022	38	956	3.975%
5/1/2022	36	1116	3.226%
6/1/2022	34	1157	2.939%
Total	429	13929	3.080%

The volume of bike returns has risen over time, but so has the number of bikes being sold.

When we look at the rate of returns as a percent of sales, we don't see a concerning trend.



ALL

ALL

Returns all rows in a table, or all values in a column, ignoring any filters that have been applied

=**ALL**(Table or Column, [Column2], [Column3],...)

The **table** or **column** that you want to clear filters on

Examples:

- Transactions
- Products[Category]

Additional columns that you want to clear filters on (optional)

- Cannot specify columns if your first parameter is a **table**
- All columns must include the **table name** and come from the **same table**

Examples:

- 'Customer Lookup'[City], 'Customer Lookup'[Country]
- Products[Product Name]

PRO TIP:



Instead of adding filter context, **the ALL function removes it**. This is often used in “**% of Total**” calculations, when the denominator needs to remain fixed regardless of filter context.



ASSIGNMENT: CALCULATE & ALL

  NEW MESSAGE

From: Dianne A. Xu (Senior Analyst)
Subject: Return analysis follow-up

Hey again,

Thanks for the quick turnaround on that bike return analysis – crisis averted!

That got me thinking about how we could start analyzing the return data in our reports. Could you please help me create two new measures, one to calculate ALL returns (regardless of filter context), and another that divides Total Returns by All Returns?

That should allow us to see the % of returns by different products and product categories.

Reply Forward

Key Objectives

1. Create a new measure named **All Returns** to calculate the total number of returns, regardless of filter context
2. Create a new measure named **% of All Returns** that divides Total Returns by All Returns
3. Create a matrix to show % of All Returns (values) by product Category Name (rows). Which category accounts for the largest percentage of returns? The smallest?



SOLUTION: CALCULATE & ALL

 NEW MESSAGE

From: Dianne A. Xu (Senior Analyst)

Subject: Return analysis follow-up

Hey again,

Thanks for the quick turnaround on that bike return analysis – crisis averted!

That got me thinking about how we could start analyzing the return data in our reports. Could you please help me create two new measures, one to calculate ALL returns (regardless of filter context), and another that divides Total Returns by All Returns?

That should allow us to see the % of returns by different products and product categories.

Reply Forward

Solution Preview

```
1 All Returns =  
2 CALCULATE(  
3     [Total Returns],  
4     ALL(  
5         'Returns Data'  
6     ))  
7 )
```

```
1 % of All Returns =  
2 DIVIDE(  
3     [Total Returns],  
4     [All Returns]  
5 )
```

Category Name	% of All Returns
Bikes	23.60%
Clothing	14.76%
Accessories	61.64%
Total	100.00%



FILTER

FILTER

Returns a table that represents a subset of another table or expression

=**FILTER**(Table, FilterExpression)

Table to be filtered

Examples:

- Territory Lookup
- Customer Lookup

A Boolean (True/False) filter expression to be evaluated for each row of the table

Examples:

- 'Territory Lookup'[Country] = "USA"
- Calendar[Year] = 1998
- Products[Price] > [Overall Avg Price]

HEY THIS IS IMPORTANT!

FILTER is used to add new filter context, and can handle **more complex filter expressions** than CALCULATE (by referencing measures, for example)

Since FILTER returns an entire table, it's often **nested within other functions**, like CALCULATE or SUMX



PRO TIP:



Since FILTER **iterates through each row in a table**, it can be slow and computationally expensive; only use FILTER if a simple CALCULATE function won't get the job done!



ITERATOR FUNCTIONS

Iterator (or “X”) **functions** allow you to loop through the same expression on each row of a table, then apply some sort of aggregation to the results (SUM, MAX, etc.)

=SUMX(Table, Expression)

Aggregation to apply to calculated rows*

Table in which the expression will be evaluated

Expression to be evaluated for each row of the given table

Examples:

- SUMX
- COUNTX
- AVERAGEX
- RANKX
- MAXX/MINX

Examples:

- Sales
- FILTER(Sales,
RELATED(Products[Category])="Clothing")

Examples:

- [Total Orders]
- Sales[Retail Price] * Sales[Quantity]

PRO TIP:



Imagine that iterator functions **add a temporary new column** to a table, calculate a value in each row based on the given expression, then aggregate the values within that temporary column (similar to **SUMPRODUCT** in Excel)

*In this example we're looking at **SUMX**, but other iterator functions follow a similar syntax



ASSIGNMENT: ITERATORS

 NEW MESSAGE

From: **Dianne A. Xu** (Senior Analyst)
Subject: Profit calculation – HELP!

Hey,
Ethan asked for a quick analysis of company profit over the past few years, but I'm struggling with the calculation.
We need a measure that multiplies order quantity by product cost, but I'd like to do it without adding redundant columns to our Sales table.
Could you take a stab at this please?
-Dianne

Reply Forward

Key Objectives

1. Create a new measure named **Total Cost** that multiplies the order quantities in the Sales Data table by the product cost in the Product Lookup table, then calculates the sum
2. Create a new measure named **Total Profit** (revenue minus cost)
3. Create a matrix to show Total Profit (values) by Year (rows). How much profit has AdventureWorks earned so far in 2022?



SOLUTION: ITERATORS

 NEW MESSAGE

From: **Dianne A. Xu** (Senior Analyst)

Subject: **Profit calculation – HELP!**

Hey,

Ethan asked for a quick analysis of company profit over the past few years, but I'm struggling with the calculation.

We need a measure that multiplies order quantity by product cost, but I'd like to do it without adding redundant columns to our Sales table.

Could you take a stab at this please?

-Dianne

Reply Forward

Solution Preview

```
1 Total Cost =  
2 SUMX(  
3     'Sales Data',  
4     'Sales Data'[Order Quantity]  
5     *  
6     RELATED(  
7         'Product Lookup'[Product Cost]  
8     )  
9 )
```

```
1 Total Profit =  
2 [Total Revenue] - [Total Cost]
```

Year	Total Profit
2020	\$2,601,606
2021	\$3,967,023
2022	\$3,888,952
Total	\$10,457,581



TIME INTELLIGENCE

Time Intelligence patterns are used to calculate common date-based comparisons

Performance
To-Date

=**CALCULATE**(Measure, **DATESYTD**(Calendar[Date]))

Use **DATESYTD** for Years, **DATESQTD** for Quarters, **DATESMTD** for Months

Previous
Period

=**CALCULATE**(Measure, **DATEADD**(Calendar[Date], -1, **MONTH**))

Select an interval (**DAY**, **MONTH**, **QUARTER**, or **YEAR**) and the
of intervals to compare (e.g. previous month, rolling 10-day)

Running
Total

=**CALCULATE**(Measure,
DATESINPERIOD(Calendar[Date], **MAX**(Calendar[Date]), -10, **DAY**))



PRO TIP:

To calculate a **moving average**, use the running total calculation above and **divide by the number of intervals**



ASSIGNMENT: TIME INTELLIGENCE



NEW MESSAGE

From: **Dianne A. Xu** (Senior Analyst)

Subject: **Time Intelligence Measures**

Hey there, need a big favor!

The leadership team has been asking a lot of questions about month-over-month and year-over-year comparisons, and I've been pulling the numbers pretty manually.

Could you please add the following list of measures, to make these metrics easier to track and share with stakeholders?

Thank you!

-Dianne

Reply **Forward**

Key Objectives

Add the following measures to the model:

- 1. Previous Month Returns**
- 2. Previous Month Orders**
- 3. Previous Month Profit**
- 4. Order Target** (10% increase over previous month)
- 5. Profit Target** (10% increase over previous month)
- 6. 90-day Rolling Profit**



SOLUTION: TIME INTELLIGENCE

 NEW MESSAGE

From: **Dianne A. Xu** (Senior Analyst)

Subject: **Time Intelligence Measures**

Hey there, need a big favor!

The leadership team has been asking a lot of questions about month-over-month and year-over-year comparisons, and I've been pulling the numbers pretty manually.

Could you please add the following list of measures, to make these metrics easier to track and share with stakeholders?

Thank you!

-Dianne

Reply **Forward**

Solution Preview

```
1 Previous Month Orders =  
2 CALCULATE(  
3     [Total Orders],  
4     DATEADD(  
5         'Calendar Lookup'[Date],  
6         -1,  
7         MONTH  
8     )  
9 )
```

```
1 90-day Rolling Profit =  
2 CALCULATE(  
3     [Total Profit],  
4     DATESINPERIOD(  
5         'Calendar Lookup'[Date],  
6         LASTDATE(  
7             'Calendar Lookup'[Date]  
8         ),  
9         -90,  
10        DAY  
11    )  
12 )
```

```
1 Order Target =  
2 [Previous Month Orders] * 1.1
```



DAX BEST PRACTICES



Know when to use calculated columns vs. measures

- *Use calculated columns for filtering, and measures for aggregating values*



Use explicit measures, even for simple calculations

- *Explicit measures can be referenced anywhere, and nested within other measures*



Use fully-qualified column references in measures

- *This makes your DAX more readable, and differentiates column references from measure references*



Move column calculations “upstream” when possible

- *Adding calculated columns at the source or in Power Query improves report speed and efficiency*



Minimize the use of “expensive” iterator functions

- *Use iterators with caution, especially if you are working with large tables or complex models*

VISUALIZING DATA



VISUALIZING DATA



In this section we'll **build dynamic interactive reports**, introduce visualization best practices, and explore features like bookmarks, drillthrough filters, parameters, tooltips, and more

TOPICS WE'LL COVER:

Data Viz Best Practices

Formatting & Filtering

Bookmarks

Report Interactions

User Roles

Parameters

Custom Tooltips

Mobile Layouts

GOALS FOR THIS SECTION:

- Review frameworks and best practices for visualizing data and designing effective reports and dashboards
- Explore tools and techniques for inserting, formatting and filtering visuals in the Power BI Report view
- Add interactivity using tools like bookmarks, slicer panels, parameters, tooltips, and report navigation
- Learn how to configure row-level security with user roles
- Optimize reports for mobile viewing using custom layouts



THREE KEY QUESTIONS

1

What **TYPE OF DATA** are you working with?

- Geospatial? Time-series? Hierarchical? Financial?
-

2

What do you want to **COMMUNICATE**?

- Comparison? Composition? Relationship? Distribution?
-

3

Who is the **END USER** and what do they need?

- Analyst? Manager? Executive? General public?



THREE KEY QUESTIONS

1 What **TYPE OF DATA** are you working with?

 Time-series

 Financial

 Geospatial

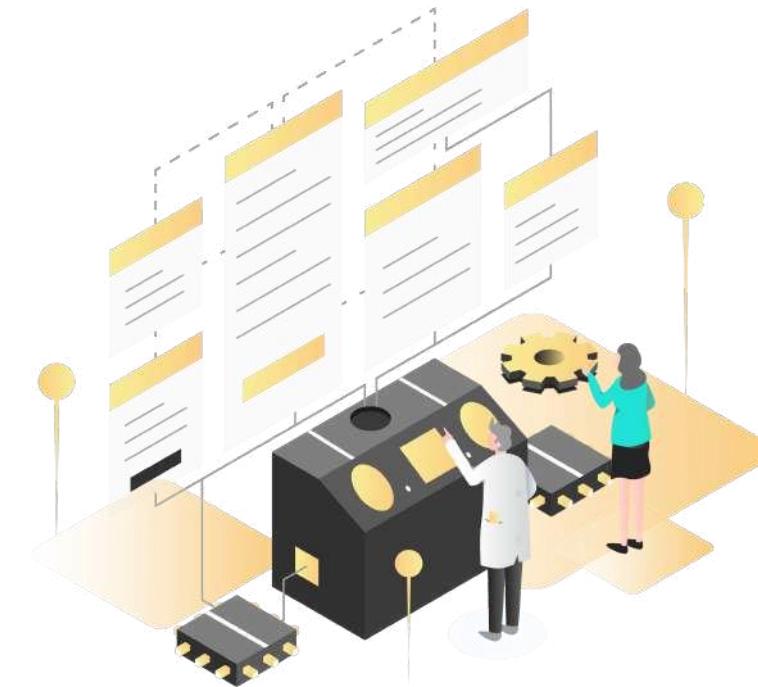
 Textual

 Categorical

 Funnel

 Hierarchical

 Survey



The type of data you're working with often determines **which type of visual will best represent it**; for example, using maps to represent geospatial data, line charts for time-series data, or tree maps for hierarchical data.

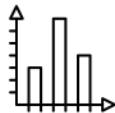


THREE KEY QUESTIONS

2

What do you want to **COMMUNICATE?**

COMPARISON



Used to **compare values over time or across categories**

Common visuals:

- Column/Bar Chart
- Clustered Column/Bar
- Data Table/Heat Map
- Radar Chart
- Line Chart (*time series*)
- Area Chart (*time series*)

COMPOSITION



Used to **break down the component parts of a whole**

Common visuals:

- Stacked Bar/Column Chart
- Pie/Donut Chart
- Stacked Area (*time series*)
- Waterfall Chart (*gains/losses*)
- Funnel Chart (*stages*)
- Tree Map/sunburst (*hierarchies*)

DISTRIBUTION

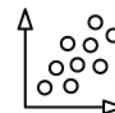


Used to **show the frequency of values within a series**

Common visuals:

- Histogram
- Density Plot
- Box & Whisker
- Scatter Plot
- Data Table/Heat Map
- Map/Choropleth (*geospatial*)

RELATIONSHIP



Used to **show correlation between multiple variables**

Common visuals:

- Scatter Plot
- Bubble Chart
- Data Table/Heat Map
- Correlation Matrix

Keep it simple! While there are *hundreds* of charts to choose from, basic options like bars and columns, line charts, histograms and scatterplots often tell the simplest and clearest story



THREE KEY QUESTIONS

3 Who is the **END USER** and what do they need?

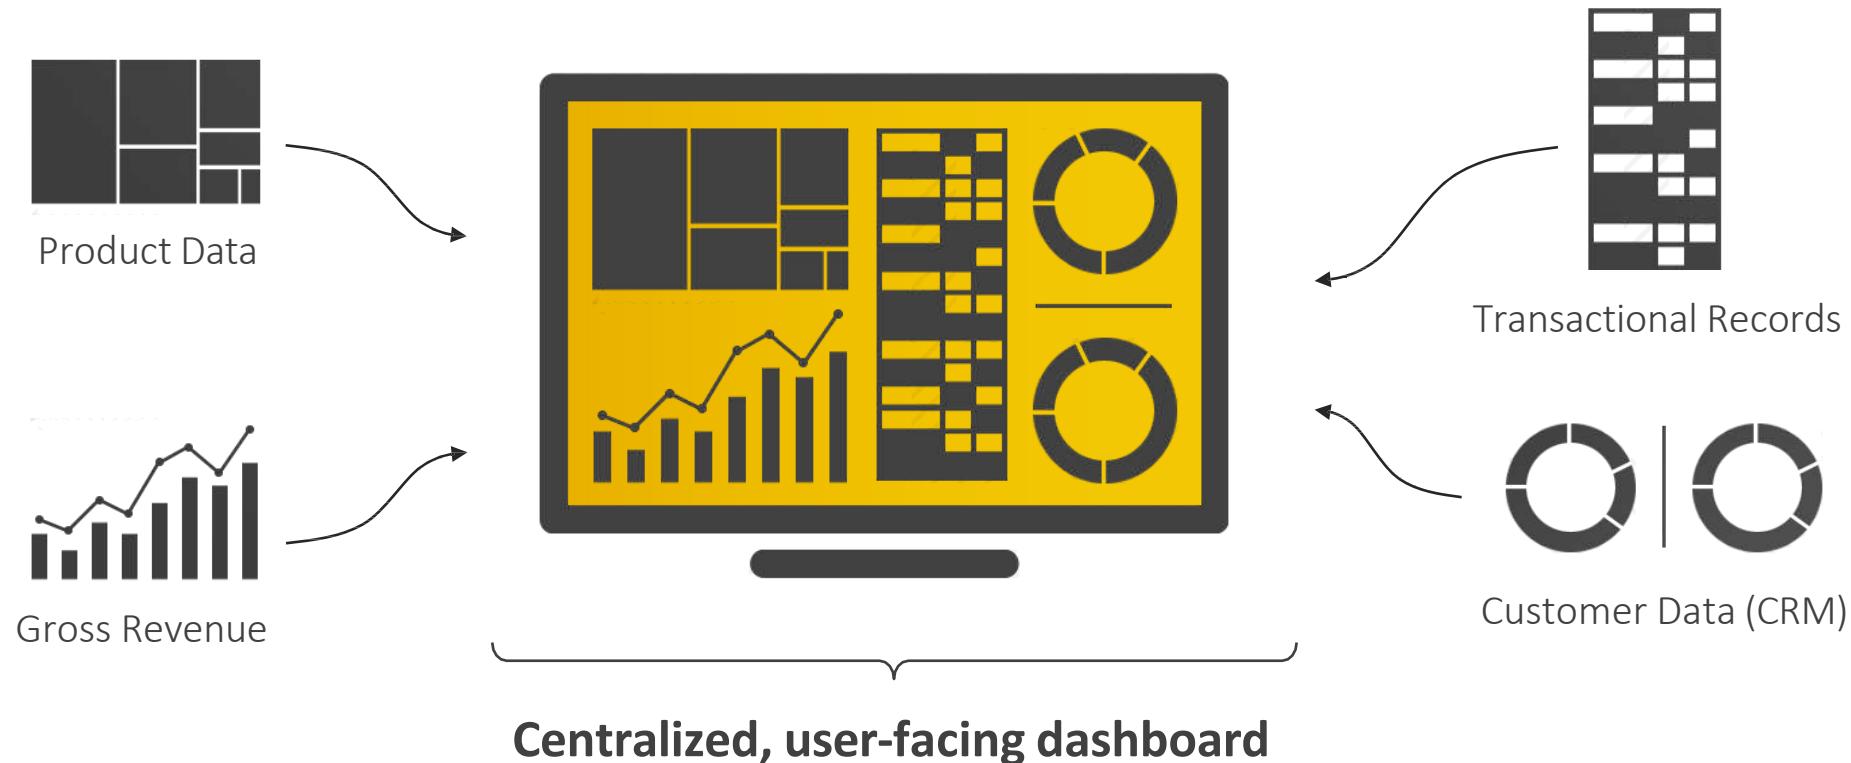


How you visualize and present your data is a function of **who will be consuming it**; a fellow analyst may want to see granular details, while managers and executives often prefer topline KPIs and clear, data-driven insights



ANALYTICS DASHBOARDS

Dashboards are analytics tools designed to consolidate data from multiple sources, track key metrics at a glance, and facilitate data-driven storytelling and decision making



DASHBOARD DESIGN FRAMEWORK



1

Define the purpose

2

Choose the right metrics

3

Present the data effectively

4

Eliminate clutter & noise

5

Use layout to focus attention

6

Tell a clear story

A well-designed dashboard should **serve a distinct purpose for a distinct audience**, use **clear and effective metrics and visuals**, and **provide a simple, intuitive user experience**.



Key questions to consider:

- Who are the **end-users** of your dashboard?
- What are their **key business goals** and objectives?
- What are the **most important questions** they need answers to?
- How can I present information **as clearly as possible**?

“

Perfection is achieved not when there is nothing more to add, **but when there is nothing left to take away**

Antoine de Saint-Exupery

”



THE REPORT VIEW

Insert Menu (Add pages, visuals, buttons, shapes, images, etc.)

Report View

Report Canvas

Report Pages (each tab is a blank report canvas)

View Options (Zoom, fit to page)

Panes (Data, Format, Bookmarks, Selection)

Filter Pane (Page-level, report-level, visual-level filters)

The screenshot illustrates the Microsoft Power BI Report View interface. The main workspace is the 'Report Canvas', displaying a dashboard with several key performance indicators (KPIs) and data visualizations. The top navigation bar includes the 'File' tab and the 'Insert' tab, which is currently active. The 'Visuals' section of the Insert ribbon is highlighted, showing options like 'New page', 'New visual', and 'More visual'. To the right of the canvas is the 'Filters' pane, which contains sections for 'Filters on this page' and 'Filters on all pages', each with a placeholder 'Add data fields here'. Below the filters is a large empty space. Further down is the 'Panes' pane, which includes sections for 'Data', 'Format', 'Bookmarks', and 'Selection', also with 'Add data fields here' placeholders. At the bottom of the interface is the 'View Options' bar, featuring zoom controls and a 'Fit to page' button. The left side of the interface features a vertical legend with labels: 'Report View' (pointing to the navigation bar), 'Report Canvas' (pointing to the main dashboard area), 'Report Pages' (pointing to the tabs at the bottom), 'Insert Menu' (pointing to the Insert ribbon), 'Panes' (pointing to the pane section), 'Filter Pane' (pointing to the filter section), and 'View Options' (pointing to the bottom bar). The bottom of the legend also indicates 'Page 1 of 8'.



ASSIGNMENT: CARDS



NEW MESSAGE

From: **Victor Ignatius Zabel (BI Analyst)**

Subject: **Let's get visual!**

Hey there!

We've kicked off the visualization work for our Power BI dashboard, and I'm hoping you can help.

For now I'd love for you to focus on building out the **Customer Detail** report. Can you start by adding some KPIs to show total customers and revenue per customer?

-Vic

Reply **Forward**

Key Objectives

1. Insert a **card** in the **Customer Detail** report page to show **Total Customers**, and rename the field "UNIQUE CUSTOMERS"
2. Add a background shape and match the formatting of the cards in the **Exec Dashboard** tab
3. Copy and paste to create a second card showing **Average Revenue per Customer**, and rename the field "REVENUE PER CUSTOMER"



SOLUTION: CARDS

  NEW MESSAGE

From: **Victor Ignatius Zabel** (BI Analyst)
Subject: Let's get visual!

Hey there!
We've kicked off the visualization work for our Power BI dashboard, and I'm hoping you can help.
For now I'd love for you to focus on building out the **Customer Detail** report. Can you start by adding some KPIs to show total customers and revenue per customer?

-Vic

Reply Forward

Solution Preview





BUILDING & FORMATTING CHARTS

The screenshot shows the Power BI interface. On the left, there is a bar chart titled "Orders by Category" with three bars: Accessories (17.0K), Bikes (13.9K), and Clothing (7.0K). Below the chart is a "Select data" pane listing various data sources. On the right, a "Build a visual" contextual menu is open, showing options for "Visual types", "Y-axis", "X-axis", "Legend", "Small multiples", and "Tooltips". A yellow arrow points from the "Select data" pane towards the "X-axis" section of the build menu.

Orders by Category

Category	Value
Accessories	17.0K
Bikes	13.9K
Clothing	7.0K

Select data

- Measure Table
- Calendar Lookup
- Customer Lookup
- Customer Metric Selection
- Metric Selection
- Price Adjustment (%)
- Product Categories Lookup
- Product Lookup
- Product Subcategories Lookup
- Returns Data
- Rolling Calendar
- Sales Data
- Territory Lookup

Build a visual

Visual types

- Bar
- Line
- Stacked Bar
- Stacked Line
- Table
- Matrix

Y-axis

Category Name

+Add data

X-axis

Total Orders

+Add data

Legend

+Add data

Small multiples

+Add data

Tooltips

Total Revenue

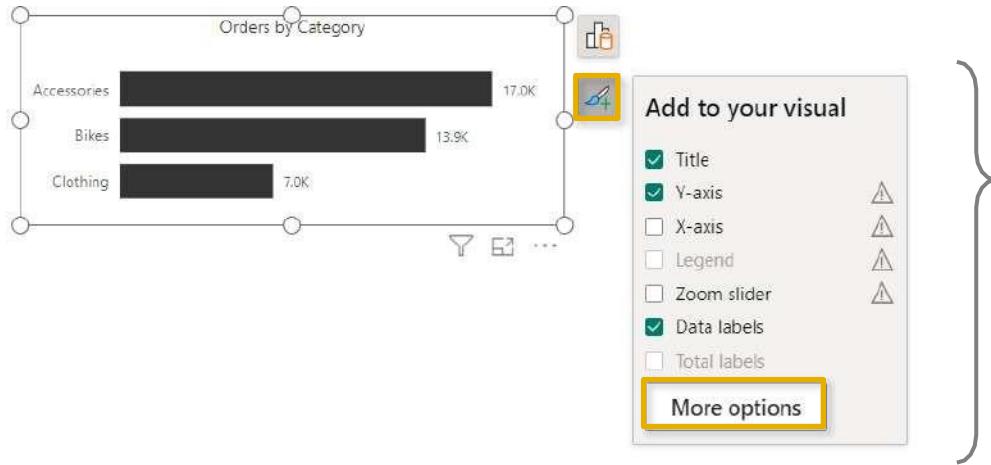
+Add data

The **Build** menu allows you to change the visual type, auto-suggest visuals, and add data to customize chart components (*x-axis, y-axis, legend, tooltips, etc.*)

- This is a **contextual menu**, so you will only see options which are relevant to the selected visual
- You can build visuals by either inserting a specific chart type and adding data, or by dragging a field from the Data pane onto the canvas

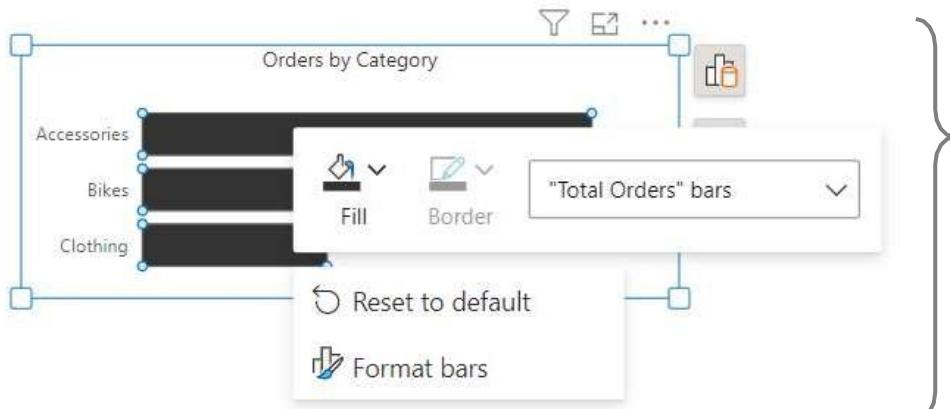


BUILDING & FORMATTING CHARTS



The **Format** menu allows you to quickly add common chart elements (*title, axis labels, data labels, legends, etc.*) and access additional options and properties in the Format pane

- This is a **contextual menu**, so you will only see options which are relevant to the selected visual

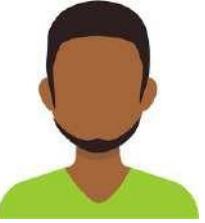


Enable **on-object formatting** by double-clicking the chart object (*or right-click > format*), which allows you to select and edit individual chart elements

- On-object formatting is only available for certain visuals (bar, column, line, area, combo & scatter)



ASSIGNMENT: LINE CHARTS



NEW MESSAGE

From: **Victor Ignatius Zabel (BI Analyst)**

Subject: **Customer count by week**

Nice work on those cards!

Next up let's add a weekly line chart to show how our customer base is trending over time.

Please add a zoom bar to make it interactive, and format the tooltips to match the line chart in the Exec Dashboard.

Thanks!

-Vic

Reply **Forward**

Key Objectives

1. Add a **line chart** to the **Customer Detail** report showing **Total Customers** by week
2. Add a **trend line** and a **zoom slider** to the x-axis
3. Enable **tooltips**, and format to match line chart in the **Exec Dashboard** tab



SOLUTION: LINE CHARTS

  NEW MESSAGE

From: **Victor Ignatius Zabel (BI Analyst)**
Subject: **Customer count by week**

Nice work on those cards!

Next up let's add a weekly line chart to show how our customer base is trending over time.

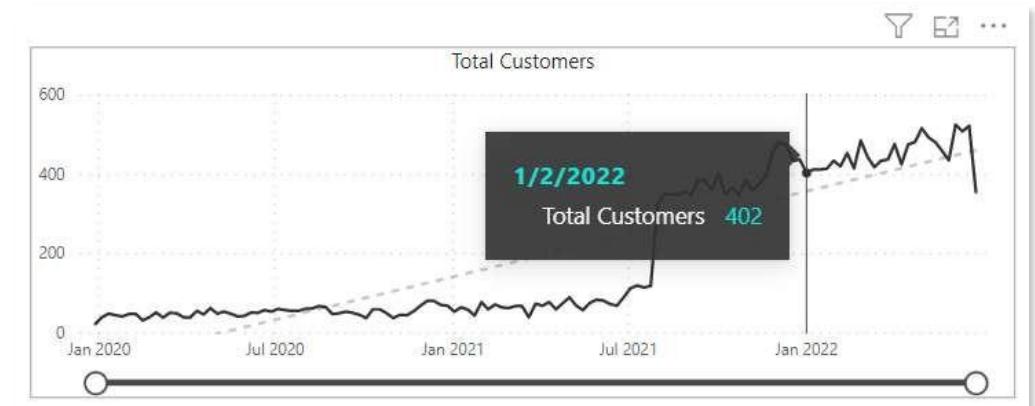
Please add a zoom bar to make it interactive, and format the tooltips to match the line chart in the Exec Dashboard.

Thanks!

-Vic

Reply Forward

Solution Preview





FILTERING OPTIONS

Filters ⟳ >

Search

Filters on this visual ...

Sum of Order Quantity
is (All)

Add data fields here

Filters on this page ...

Add data fields here

Filters on all pages ...

Add data fields here

There are **3 types of filters** accessible from the **Filters** pane*:

1. **Visual-level** filters apply to specific visuals
2. **Page-level** filters apply to all visuals on the report page
3. **Report-level** filters apply to all visuals across all report pages

**Drillthrough* filters can be configured in the page formatting pane – more on that later!

Filter type ⓘ
Basic filtering

Search

Select all
 Accessories 1
 Bikes 1
 Clothing 1
 Components 1

Basic Options

Filter type ⓘ
Top N

Show items
Top 2

By value
Total Orders

Apply filter

Top N Options

Show items when the value

is greater than
is less than
is less than or equal to
is greater than
is greater than or equal to
is
is not

Advanced (Values)

Filter type ⓘ
Advanced filtering

Show items when the value

contains
contains
does not contain
starts with
does not start with
is

Advanced (Text)

Filters can be configured using basic **selections**, **logical operators**, or **Top N** conditions



ASSIGNMENT: DONUT CHARTS



NEW MESSAGE

From: **Victor Ignatius Zabel (BI Analyst)**

Subject: **Customer demographics**

Good morning!

Just got a note from Ethan to see if we can build some demographic info into the customer report.

Let's add a couple donut charts to show the composition of customers by income level and occupation. We'll want to limit to just a few segments (maybe 3?) and do some formatting to match the rest of the dashboard.

Thanks, you rock!

-Vic

Reply **Forward**

Key Objectives

1. Add a **donut chart** to the **Customer Detail** report showing **Total Orders** by **Income Level**
2. Add a **chart title**, turn off the **legend**, and update the **data labels** to show the category and value (font size 8, 1 decimal place)
3. Update the colors of the slices to match the screenshot in the solution preview
4. Add a **visual-level filter** to exclude customers with a "Very High" income level
5. Copy the chart to show **Total Orders** by **Occupation**, and add a **visual-level filter** to display the three occupations with the most orders (*bonus points if you use a Top N filter!*)



SOLUTION: DONUT CHARTS

  NEW MESSAGE

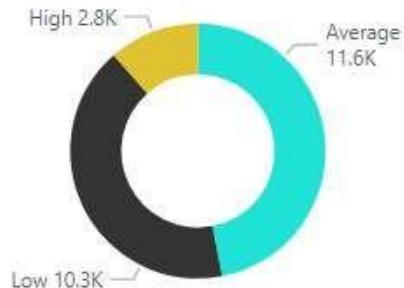
From: **Victor Ignatius Zabel (BI Analyst)**
Subject: **Customer demographics**

Good morning!
Just got a note from Ethan to see if we can build some demographic info into the customer report.
Let's add a couple donut charts to show the composition of customers by income level and occupation. We'll want to limit to just a few segments (maybe 3?) and do some formatting to match the rest of the dashboard.
Thanks, you rock!
-Vic

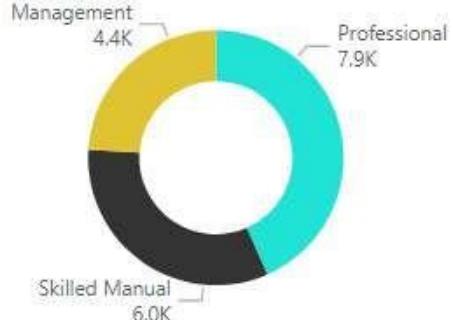
[Reply](#) [Forward](#)

Solution Preview

Orders by Income Level



Orders by Occupation





CONDITIONAL FORMATTING

Top 10 Products	Orders	Revenue	Return %
Water Bottle - 30 oz.	3,983	\$39,755	1.95%
Patch Kit/8 Patches	2,952	\$13,506	1.61%
Mountain Tire Tube	2,846	\$28,333	1.64%
Road Tire Tube	2,173	\$17,265	1.55%
Sport-100 Helmet, Red	2,099	\$73,444	3.33%
AWC Logo Cap	2,062	\$35,865	1.11%
Sport-100 Helmet, Blue	1,995	\$67,112	3.31%
Fender Set - Mountain	1,975	\$87,041	1.36%
Sport-100 Helmet, Black	1,940	\$65,262	2.68%
Mountain Bottle Cage	1,896	\$38,062	2.02%



Conditional formatting allows you to dynamically format Table or Matrix visuals based on cell values

- Conditionally formatting options can be found in the **Format** pane, under **Cell elements**
- Options include background color, font color, data bars, icons, or Web URL

Format X

Search

Visual Properties ...

Cell elements

Apply settings to Series Orders

Background color Off On

Font color Off On

Data bars On Off

Icons Off On

Web URL Off On

Format X

Search

Visual Properties ...

Cell elements

Apply settings to Series Revenue

Background color Off On

Font color Off On

Data bars Off On

Icons Off On

Web URL Off On

Format X

Search

Visual Properties ...

Cell elements

Apply settings to Series Return %

Background color On Off

Font color Off On

Data bars Off On

Icons Off On

Web URL Off On

ASSIGNMENT: TABLES





NEW MESSAGE

From: **Victor Ignatius Zabel (BI Analyst)**

Subject: **Top customer table**

Hey there, this customer report is really coming together! Since the management team needs a way to identify high-value customers, let's add a table to our report showing customer keys, full names, orders, and revenue. Probably makes sense to add some conditional formatting and limit to the top 100 customers for now.

Thanks!

-Vic

Reply **Forward**

Key Objectives

1. Add a **table** to the **Customer Detail** report to show **Customer Key**, **Full Name**, **Total Orders** (as “Orders”) and **Total Revenue** (as “Revenue”)
2. Use conditional formatting to add light gray **data bars** to the orders column and a white > blue **color scale** to the revenue column
3. Add a **visual-level filter (Top N)** to show the 100 customers with the most orders, and add a **chart title** (“Top 100 Customers”)
4. **Sort** the table descending by orders



SOLUTION: TABLES

  NEW MESSAGE

From: **Victor Ignatius Zabel (BI Analyst)**

Subject: **Top customer table**

Hey there, this customer report is really coming together!

Since the management team needs a way to identify high-value customers, let's add a table to our report showing customer keys, full names, orders, and revenue.

Probably makes sense to add some conditional formatting and limit to the top 100 customers for now.

Thanks!

-Vic

Reply Forward

Solution Preview

Top 100 Customers				
Customer Key	Full Name	Orders	Revenue	
11091	Mr. Dalton Perez	26	\$1,513	
11223	Mrs. Hailey Patterson	26	\$1,616	
11300	Mr. Fernando Barnes	26	\$1,839	
11330	Mr. Ryan Thompson	26	\$1,597	
11331	Mrs. Samantha Jenkins	26	\$1,740	
11185	Mrs. Ashley Henderson	25	\$1,717	
11200	Mr. Jason Griffin	25	\$1,614	
11176	Mr. Mason Roberts	24	\$1,526	
11262	Mrs. Jennifer Simmons	24	\$1,465	
11277	Mr. Charles Jackson	24	\$1,777	
11287	Mr. Henry Garcia	24	\$1,443	
11566	Ms. April Shan	24	\$1,424	
11711	Mr. Daniel Davis	24	\$1,404	
11276	Mrs. Nancy Chapman	23	\$1,111	
11203	Mr. Luis Diaz	17	\$1,002	
11215	Mrs. Ana Perry	17	\$1,336	
11078	Ms. Gina Martin	16	\$991	
Total			1,272	\$615,328

ASSIGNMENT: TOP N TEXT CARDS





NEW MESSAGE

From: **Victor Ignatius Zabel (BI Analyst)**

Subject: **Top customers by revenue**

Hey,

Ethan is loving the customer report so far – great job!

He mentioned that he'd like to highlight top customers based on *revenue* as well, so I'm thinking we could add some text cards to show the top customer name, along with total revenue and the number of orders placed.

We'll be offering some coupons based on how much customers have spent in the past, so accuracy is critical here!

-Vic

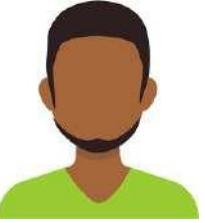
Reply **Forward**

Key Objectives

1. Add a **card** to the **Customer Detail** report to show **Full Name**
 - What do you notice when you filter the report for low income customers?
(Hint: check your value against the table)
2. Add a **visual-level filter (Top N)** to show the top customer (Full Name) in terms of **Total Revenue**
 - How could you modify the Top N filter to correct this?
2. Copy and paste the card (x2) to show **Total Orders** and **Total Revenue** for the top customer
3. Add **text boxes** for titles and adjust formatting to match the solution preview



SOLUTION: TOP N TEXT CARDS

  NEW MESSAGE

From: **Victor Ignatius Zabel (BI Analyst)**
Subject: **Top customers by revenue**

Hey,
Ethan is loving the customer report so far – great job!
He mentioned that he'd like to highlight top customers based on *revenue* as well, so I'm thinking we could add some text cards to show the top customer name, along with total revenue and the number of orders placed.
We'll be offering some coupons based on how much customers have spent in the past, so accuracy is critical here!
-Vic

Reply Forward

Solution Preview

Top Customer (by revenue):

Mr. Maurice Shan

Orders:

6

Revenue:

\$12.4K



MAP VISUALS

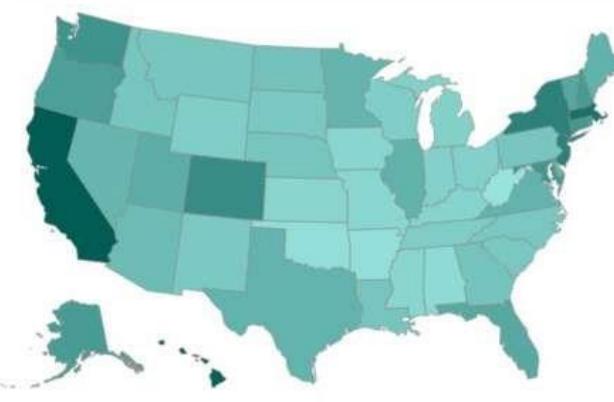
Map



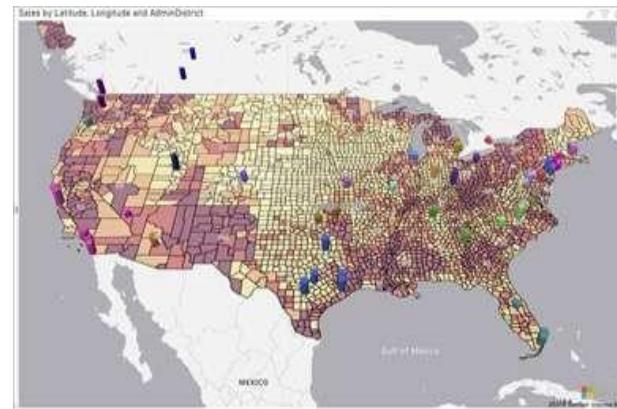
Filled map



Shape map



Azure map



Power BI includes several types of **map visuals** powered by Bing Maps

Tips for creating accurate maps:

1. Assign **categories** to geospatial fields
2. Add **multiple location** fields
3. Use **latitude/longitude** when possible



HEY THIS IS IMPORTANT!

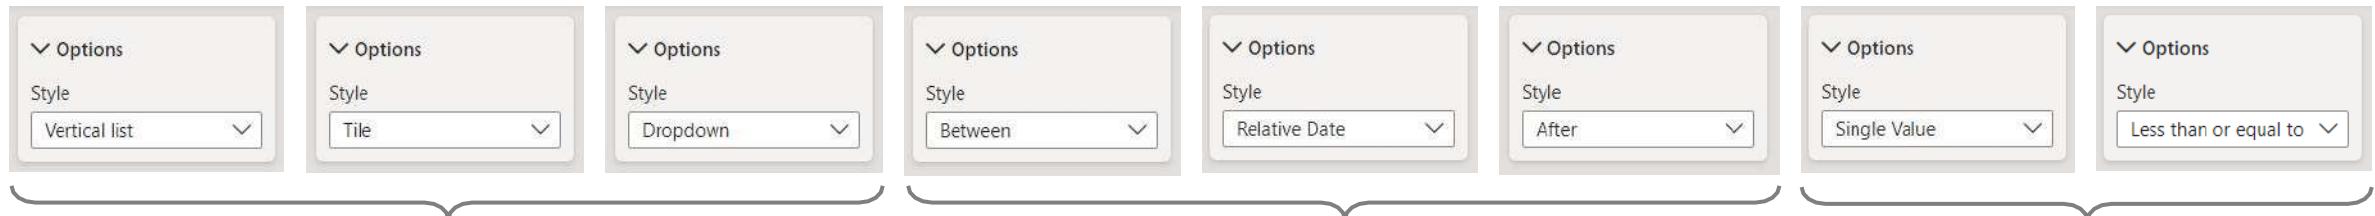
An administrator may need to **enable maps in your tenant settings** in order to use them in Power BI Service



SLICERS

Slicers are visual filters which affect all other visuals on a report page (by default)

- Slicers can take many formats depending on the data type, including **lists**, **dropdowns**, **tiles**, **ranges**, and more



Categorical/Text options

IncomeLevel

- Average
- High
- Low
- Very High

Continent, Country, Region

- Europe
 - France
 - Germany
 - United Kingdom
- North America
- Pacific

Date/Time options

Date

1/1/2020 6/30/2022

Date

Last 11 Months

3/18/2022 - 2/17/2023

Date

1/1/2020 6/30/2022

Numeric Range options

Price Adjustment (%)

0.20

Price Adjustment (%)

-1.00 0.50

PRO TIP:
Use **Apply/Clear All Slicers** buttons for more filtering control



ASSIGNMENT: SLICERS



NEW MESSAGE

From: **Victor Ignatius Zabel (BI Analyst)**

Subject: **Year slicer for customer report**

Hey there, quick request when you get a sec...
Could you please add a slicer to the customer report, so that users can filter the entire page by year?
No preference for which specific type of slicer you use, as long as managers can filter customers for a specific year or across multiple years.
Thanks!
-Vic

Reply **Forward**

Key Objectives

1. Add a **slicer** to filter the **Customer Detail** report page by **Year**
2. Add a **visual-level filter** to exclude blanks
3. Choose any **slicer style** that allows users to filter individual years or across multiple years



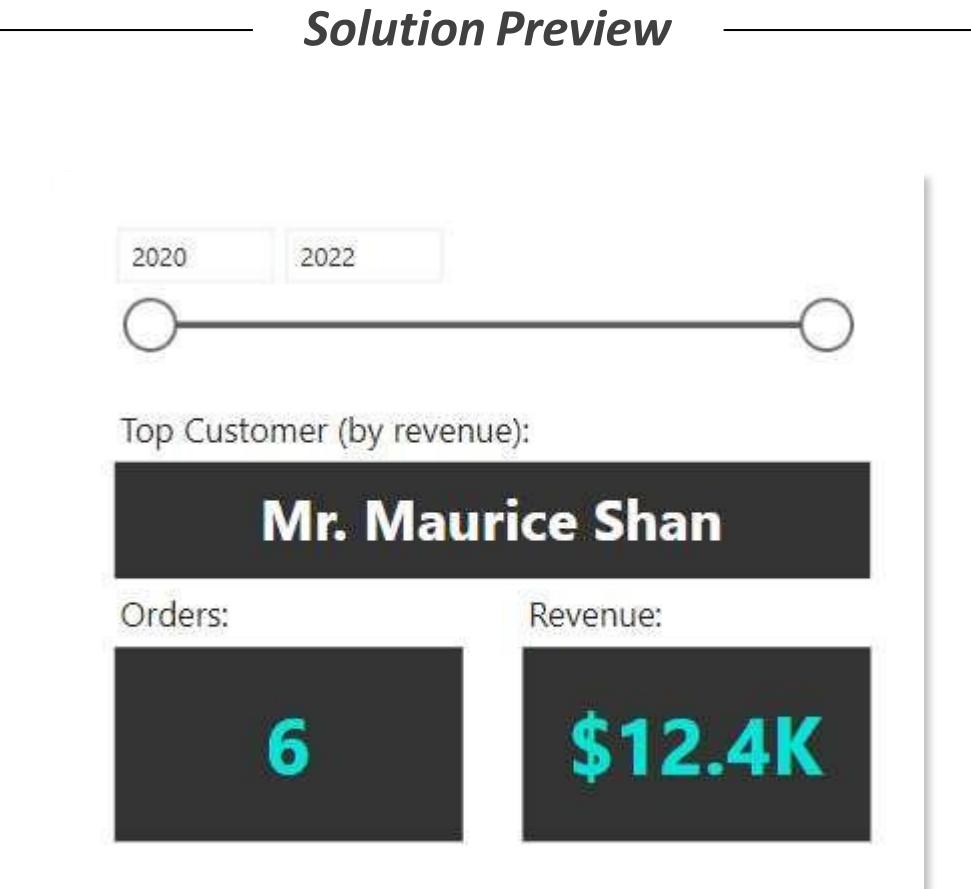
SOLUTION: SLICERS

  NEW MESSAGE

From: **Victor Ignatius Zabel (BI Analyst)**
Subject: **Year slicer for customer report**

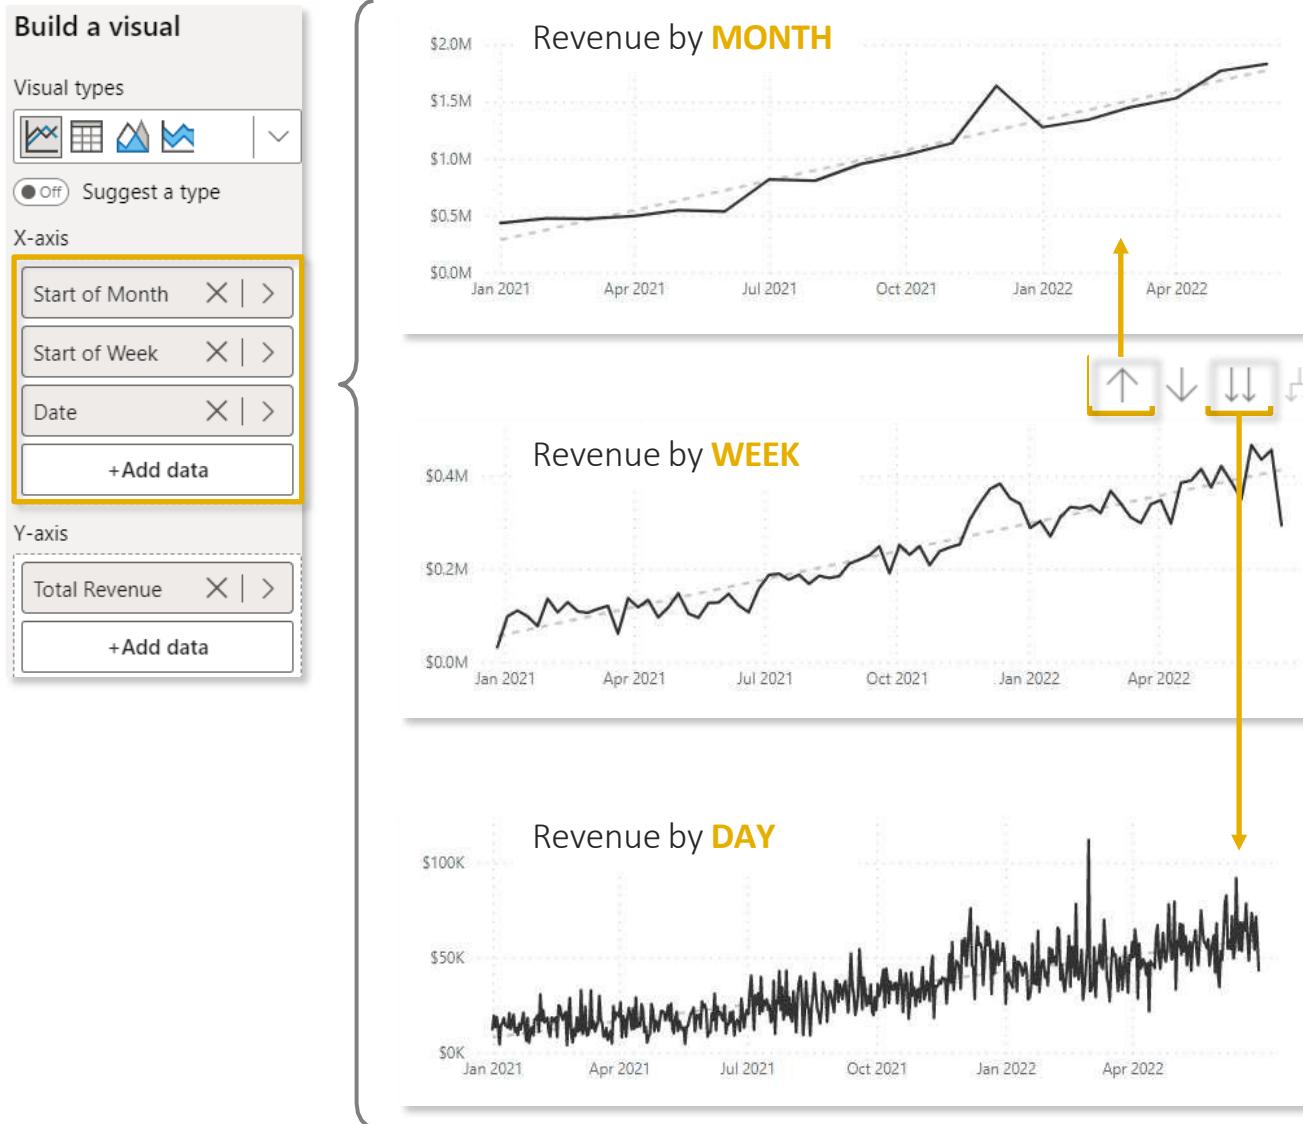
Hey there, quick request when you get a sec...
Could you please add a slicer to the customer report, so that users can filter the entire page by year?
No preference for which specific type of slicer you use, as long as managers can filter customers for a specific year or across multiple years.
Thanks!
-Vic

[Reply](#) [Forward](#)





DRILL UP & DRILL DOWN



Drill Up and **Drill Down** tools allow you to switch between different levels of granularity

- In this example users can “drill up” from **weekly** to **monthly**, or “drill down” to **daily**
- The single down arrow activates **drill mode**, allowing users to drill by clicking data points
- The forked down arrow **expands each level** of the hierarchy (used in matrix visuals)

PRO TIP:

Use **location hierarchies** and enable drill mode to create interactive map visuals



ASSIGNMENT: DRILL DOWN



NEW MESSAGE

From: **Victor Ignatius Zabel (BI Analyst)**

Subject: **Dynamic time periods**

Hey again, just got some feedback from the managers about our customer report.

Chad loves the weekly trending chart, but Thad wants to see the data by *day* and Vlad was hoping for an *annual* breakdown.

Instead of building multiple versions of the same line chart, could you please make it interactive so that Chad, Thad and Vlad get the views they want?

Thanks!

-Vic

Reply **Forward**

Key Objectives

1. In the **Customer Detail** report, update the X-axis of the line chart to pull in **Date Hierarchy**
2. Use the chart header to **drill up** and **drill down** to explore trends at each level of granularity
3. Test **drill mode** to change the granularity by selecting individual data points in the chart
 - Why do some weeks look very low?
4. Turn off drill mode and show the chart at a weekly level of granularity by default



SOLUTION: DRILL DOWN



NEW MESSAGE

From: **Victor Ignatius Zabel (BI Analyst)**

Subject: **Dynamic time periods**

Hey again, just got some feedback from the managers about our customer report.

Chad loves the weekly trending chart, but Thad wants to see the data by *day* and Vlad was hoping for an *annual* breakdown.

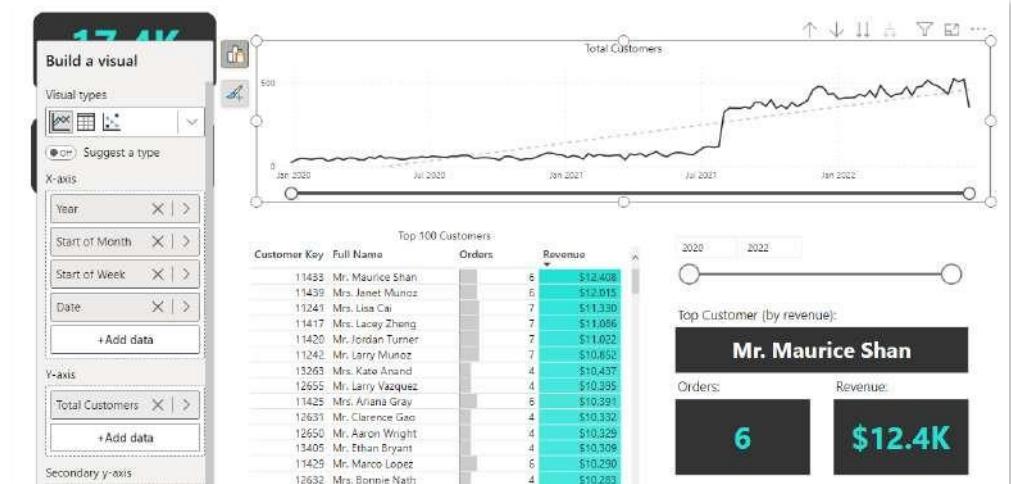
Instead of building multiple versions of the same line chart, could you please make it interactive so that Chad, Thad and Vlad get the views they want?

Thanks!

-Vic

Reply **Forward**

Solution Preview





DRILL THROUGH FILTERS

Drill through filters allow users to navigate to a specific report page, pre-filtered on the item selected

- Here we've created a **Product Detail** page, set the type to **Drillthrough**, and configured drill through from **Product Name**
- This means that users can right-click any instance of product name (i.e. in a matrix visual) and use the Drill through option to navigate straight to the Product Detail report filtered on that product (in this case "Mountain Tire Tube")

The screenshot illustrates the configuration of a drill-through filter between a matrix visual and a product detail page.

Matrix Visual (Left): A "Top 10 Products" matrix with columns: Orders, Revenue, and Return %. The "Mountain Tire Tube" row is highlighted. A context menu is open at this row, with the "Drill through" option selected and its target "Product Detail" highlighted.

Product Detail Page (Center): A dashboard for "Mountain Tire Tube" showing various metrics like monthly orders, revenue, profit, and price adjustment over time. It includes a line chart for total orders and a legend for metric selection.

Format Panel (Right): The "Format" pane is open, showing the "Page type" set to "Drillthrough". The "Drill through from" section is expanded, showing "Product Name" selected and "+Add data" available.



REPORT INTERACTIONS

Edit **report interactions** to customize how filters applied to one visual impact other visuals on the page

- Cross-filter options include **filter** (), **highlight** () and **none** (), depending on the visual type

Format > **Edit Interactions**



In this example, selecting a product in the matrix visual:

- **Filters** the line chart & KPIs
- **Highlights** the bar chart
- **Doesn't impact** the text cards

ASSIGNMENT: REPORT INTERACTIONS



  NEW MESSAGE

From: **Victor Ignatius Zabel (BI Analyst)**

Subject: **Weird report interactions**

Hey there,

I was playing with the customer report this morning and noticed some odd visual interactions. For example, selecting a specific customer shouldn't filter the line chart, and probably shouldn't filter the donut charts either.

Could you please take a pass through the report interactions and update any that seem off?

Thanks!

-Vic

Reply Forward

Key Objectives

1. On the **Customer Detail** tab, edit the **report interactions** based on the following logic:
 - When a filter is applied to the line chart, the donut charts should **filter** (not highlight)
 - When a filter is applied to the table, the line chart and donuts should **not filter**
 - The slicer should **filter all visuals** on the report page



SOLUTION: REPORT INTERACTIONS



NEW MESSAGE

From: **Victor Ignatius Zabel** (BI Analyst)

Subject: **Weird report interactions**

Hey there,

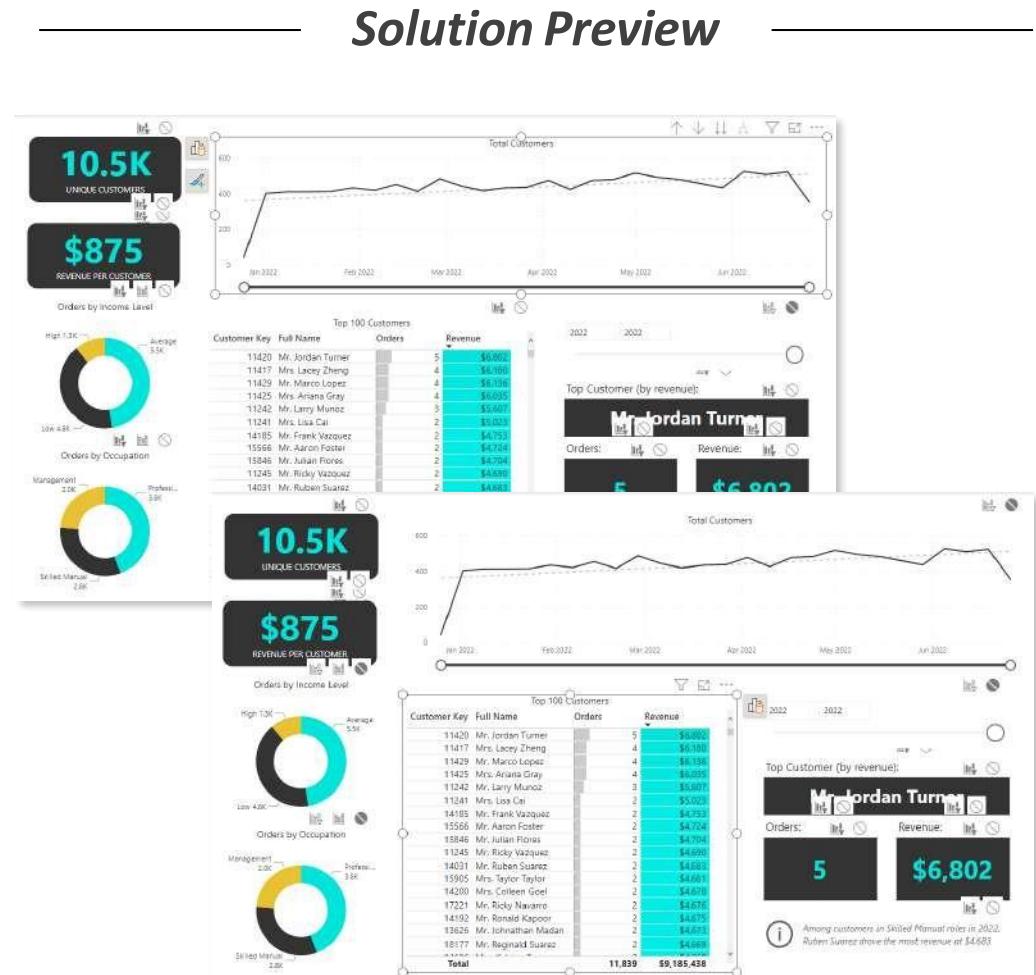
I was playing with the customer report this morning and noticed some odd visual interactions. For example, selecting a specific customer shouldn't filter the line chart, and probably shouldn't filter the donut charts either.

Could you please take a pass through the report interactions and update any that seem off?

Thanks!

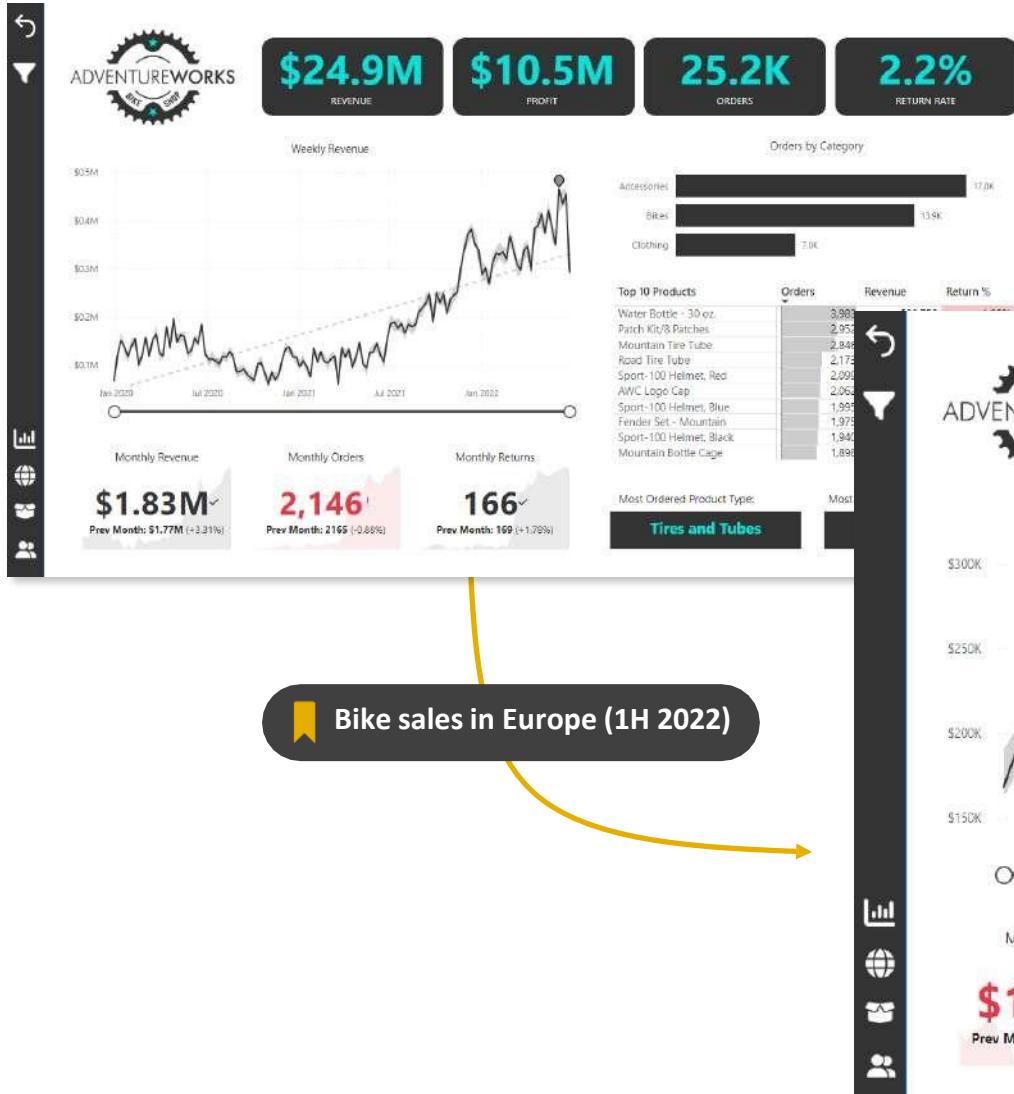
-Vic

Reply **Forward**



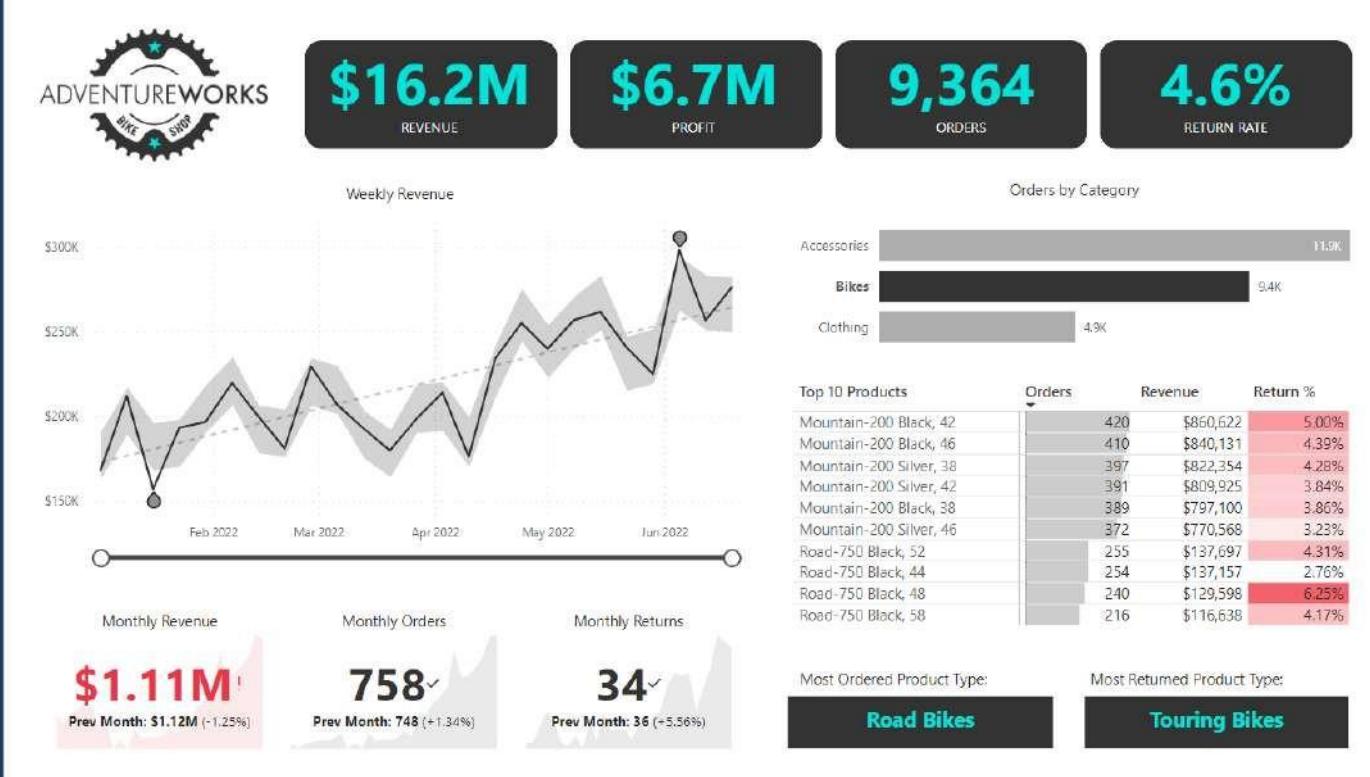


BOOKMARKS



Bookmarks capture the current state of a page, and allow users to return to that state using report actions

- Bookmarks are commonly used for clearing filters, highlighting specific insights, navigating reports, etc.





ASSIGNMENT: BOOKMARKS



NEW MESSAGE

From: **Victor Ignatius Zabel (BI Analyst)**

Subject: **Finding anything interesting?**

Hey,

Now that you're getting pretty familiar with our customer data, are you noticing any interesting insights or trends that might be worth explicitly calling out in the report?

This could be a great way for us to use bookmarks to draw attention to some specific stories in the dashboard. While we're at it, let's add another one to clear all filters from the page.

Let me know what you think!

-Vic

Reply **Forward**

Key Objectives

1. Explore the **Customer Detail** report by adjusting filters until you find an interesting insight or trend (*this can be anything you choose!*)
2. Add a new **bookmark** to capture the current state of the report, and name it “Customer Insight”
3. Insert an **Information button** and add text to the button style to summarize what you’ve found
4. Assign a **bookmark action** to the button, and link to the Customer Insight bookmark you created
5. Create a second bookmark named “Clear all Customer Filters” which returns the page to an unfiltered state, and link it to a **Reset button**
6. Test both bookmarks using **CTRL-click**



SOLUTION: BOOKMARKS



NEW MESSAGE

From: **Victor Ignatius Zabel** (BI Analyst)

Subject: **Finding anything interesting?**

Hey,

Now that you're getting pretty familiar with our customer data, are you noticing any interesting insights or trends that might be worth explicitly calling out in the report?

This could be a great way for us to use bookmarks to draw attention to some specific stories in the dashboard. While we're at it, let's add another one to clear all filters from the page.

Let me know what you think!

-Vic

Reply **Forward**

Solution Preview

Example: Among customers in Skilled Manual roles in 2022, Ruben Suarez drove the most revenue at \$4,683





PARAMETERS

Parameters allow you to create variables which can be referenced in measures and controlled via slicers

Numeric range parameters

Typically used for scenario testing, where users adjust numerical inputs to see the impact on a given output

Fields parameters

Typically used to allow users to dynamically change the metrics or dimensions displayed in a report visual

The image displays two side-by-side screenshots of the Microsoft Power BI 'New parameter' dialog box. Both screenshots show the 'Parameters' tab with the following configuration:

- What will your variable adjust?**:
 - Left screenshot: Numeric range (selected)
 - Right screenshot: Fields (selected)
- Name**:
 - Left screenshot: Price Adjustment (%)
 - Right screenshot: Y-Axis Dynamic Value
- Data type**:
 - Left screenshot: Decimal number
 - Right screenshot: (not explicitly shown)
- Minimum**: -1
- Maximum**: 1
- Increment**: 0.1
- Default**: 0

Both screenshots include a 'Fields' section on the right side, listing various data sources and tables:

- Left screenshot fields: Total Cost, Total Revenue, Total Profit.
- Right screenshot fields: Measure Table, Calendar Lookup, Customer Lookup, Price Adjustment, Product Categories Lookup, Product Lookup, Product Subcategories Lookup, Returns Data, Rolling Calendar, Sales Data, Territory Lookup.

At the bottom of both dialogs, there is a checked checkbox for 'Add slicer to this page' and two buttons: 'Create' and 'Cancel'.



EXAMPLE: NUMERIC RANGE PARAMETER

Parameters

Add parameters to visuals and DAX expressions so people can use slicers to adjust the inputs and see different outcomes. [Learn more](#)

What will your variable adjust?

Numeric range

Name: Price Adjustment (%)

Data type: Decimal number

Minimum: -1

Maximum: 1

Increment: 0.1

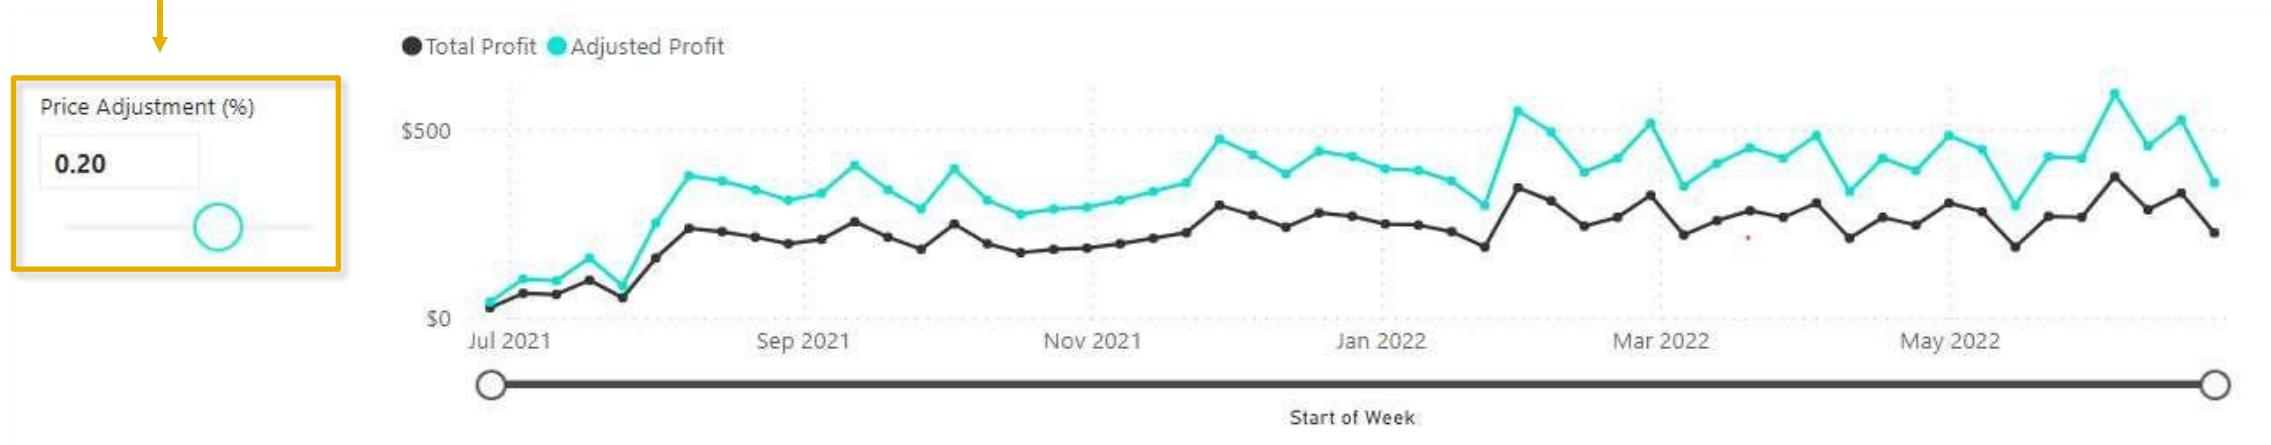
Default: 0

When you create a numeric parameter, Power BI generates **two new measures**: one to define the parameter and another to capture the selected value:

Parameter = `GENERATESERIES(-1, 1, 0.1)`

Parameter Value = `SELECTEDVALUE(Parameter[Parameter], 0)`

Here we've created a parameter named **Price Adjustment %**, added it as a slicer, and created measures to calculate **Adjusted Profit** based on the parameter value





EXAMPLE: FIELDS PARAMETER

Parameters

Add parameters to visuals and DAX expressions so people can use slicers to adjust the inputs and see different outcomes. [Learn more](#)

What will your variable adjust?

Fields

Name

Metric Selection

Add and reorder fields

- Total Orders
- Total Revenue
- Total Profit
- Total Returns
- Return Rate

Add slicer to this page

Create Cancel

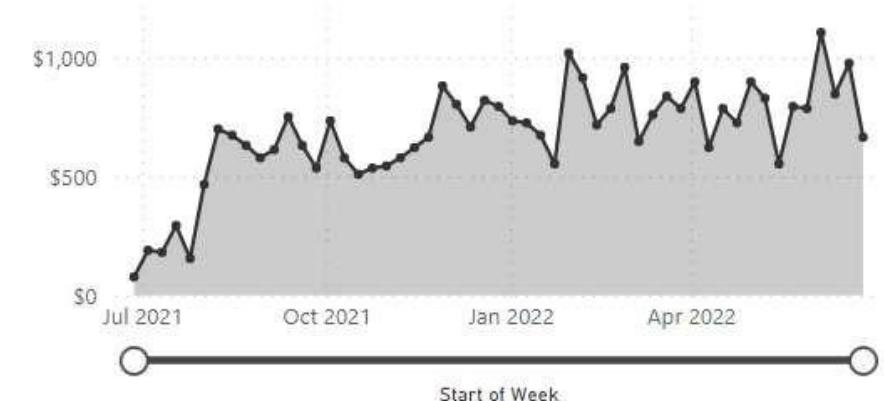
When you create a fields parameter, Power BI **adds a report slicer** and **generates a new measure** to capture the selected value:

```
1 Metric Selection = {  
2     ("Total Orders", NAMEOF('Measure Table'[Total Orders]), 0),  
3     ("Total Revenue", NAMEOF('Measure Table'[Total Revenue]), 1),  
4     ("Total Profit", NAMEOF('Measure Table'[Total Profit]), 2),  
5     ("Total Returns", NAMEOF('Measure Table'[Total Returns]), 3),  
6     ("Return Rate", NAMEOF('Measure Table'[Return Rate]), 4)  
7 }
```

Here we've created a parameter named **Metric Selection** and added it to the Y-axis to let users dynamically change the metric shown

Metric Selection

- Total Orders
- Total Revenue
- Total Profit
- Total Returns
- Return Rate



ASSIGNMENT: FIELDS PARAMETERS





NEW MESSAGE

From: **Victor Ignatius Zabel (BI Analyst)**

Subject: **More line chart updates**

Good news and bad news...

The good news is that Chad, Thad and Vlad LOVE the drill options in the line chart – nice work!

The bad news is that now they can't align on what's the best metric to show. Chad likes seeing total customers, but Vlad is pushing for revenue per customer.

What do you think we should do?

-Vic

Reply **Forward**

Key Objectives

1. Add a new **Fields parameter** named “Customer Metric Selection”, which includes **Total Customers** and **Average Revenue per Customer**
2. Add the parameter as a slicer to the **Customer Detail** report, change the slicer style to **Tile**, turn off the **header**, update to **single select**, and resize to create a horizontal layout
3. Select the DAX measure automatically created, and update the text from “Average Revenue per Customer” to “Revenue per Customer”
4. Update the line chart Y-Axis to use the **Customer Metric Selection** parameter, remove the chart title, and update the line colors to match the solution preview



SOLUTION: FIELDS PARAMETERS

  NEW MESSAGE

From: **Victor Ignatius Zabel (BI Analyst)**
Subject: **More line chart updates**

Good news and bad news...

The good news is that Chad, Thad and Vlad LOVE the drill options in the line chart – nice work!

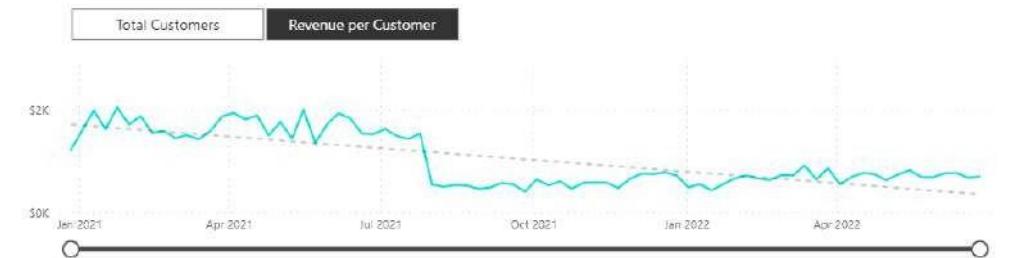
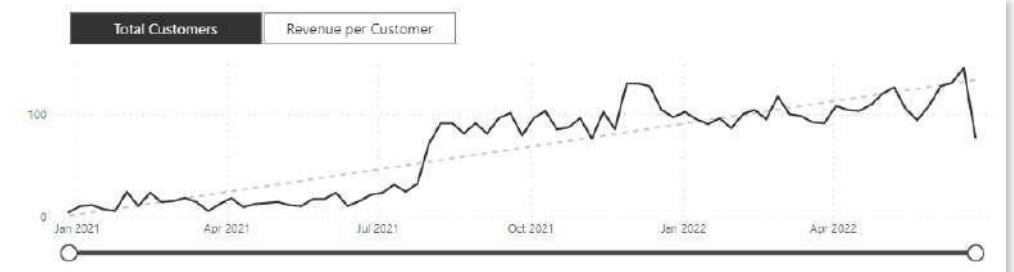
The bad news is that now they can't align on what's the best metric to show. Chad likes seeing total customers, but Vlad is pushing for revenue per customer.

What do you think we should do?

-Vic

Reply Forward

Solution Preview





PRO TIP: CUSTOM TOOLTIPS

Create **custom tooltips** by designing a new report page, setting the page type to **Tooltip**, and configuring a visual to use the “Report page” tooltip type

Category Tooltip report page

Category Tooltip report page

Weekly Orders

Total Revenue: \$24,914,567
Total Profit: \$10,457,581
Total Orders: 25,164
Total Returns: 1,809
Return Rate: 2.17%

Revenue **Return %**

Revenue	Return %
\$20,694	3.74%
\$17,849	2.60%
\$6,962	3.13%
\$56,533	2.10%
\$21,998	3.50%
\$31,083	7.14%
\$30,949	5.65%
\$32,051	7.64%
\$6,879	3.89%
\$30,345	5.65%

Format

Type: Custom
Height: 225 px
Width: 425 px
Vertical alignment: Top

Name: Category Tooltip
Page type: Tooltip
Keep all filters: On
Show tooltip on: +Add data

Orders by Category

Accessories

Weekly Orders

Revenue **Return %**

Revenue	Return %
\$5.3K	8.03%

Monthly Returns

Sport-100 Helmet, Black
Sport-100 Helmet, Red
Road Tire Tube
HL Mountain Tire

Format

Search

Visual Properties

Header icons: On

Tooltips: On

Options

Type: Report page
Page: Category Tooltip

PRO TIP:
Keep your published reports clean by **hiding your tooltip pages**



IMPORTING CUSTOM VISUALS

Power BI offers a library of **custom visuals** (via **AppSource**) from Microsoft-certified partners and developers, which can be imported into the visualizations pane

The screenshot illustrates the process of importing custom visuals into Power BI. It shows the 'Insert' tab in the ribbon, the 'Visuals' section, and the 'From AppSource' option being selected. Below, a detailed view of the AppSource listing for the 'Supermetrics Charts – Tile grid map' visual is shown, including its description, preview, and download options.

Power BI desktop interface:

- File
- Home
- Insert**
- Modeling
- View
- Optimize

Visuals section:

- New page ▾
- New visual
- More visuals ▾

From AppSource (highlighted)

From my files

Power BI visuals dialog:

- All visuals
- Organizational visuals
- AppSource visuals** (highlighted)

Enhance your Power BI experience using tested and approved visuals. [Learn more](#)

Filter by: All

Sort by: Popularity

Visuals listed:

- Gantt Chart by MA... (by MAQ LLC)
- Tachometer (by Annik Inc.)
- Zebra BI Charts (by Zebra BI)
- Supermetrics Chart... (by Supermetrics)** (highlighted)
- Balance Sheet Visu... (by K Team Solutions GmbH)

AppSource - Apps for Power BI visuals

Supermetrics Charts – Tile grid map by Supermetrics

Overview Plans + Pricing Ratings + reviews

Create (heat)maps with tiles of the same shape and size

Build your own tile grid map and compare your metrics across regions or countries. Use pre-defined map configurations or create a custom map. The tile grid map isn't applicable just to geographical data. It can also be used as a custom heatmap for many data sources. A tile grid map is a map in which each region is represented by tiles of equal size. Each map consists of square tiles for each state, region, or borough, regardless of the geographic area. They can be useful when you want to visualize data where the importance of a region does not depend on the region's geographic area. Regions with large areas won't dominate the map, and small regions will be as prominent. The Supermetrics tile grid map allows you also to recreate any other country or create any shape based on a tile grid and map it to your data.

Starts at Add for free or buy an upgraded plan

Products Power BI visuals

Publisher Supermetrics

Acquire Using Work or school account

Version 1.0.0.3

Updated 3/29/2023

Visual capabilities When this visual is used, it

- Can access external services or resources

HEY THIS IS IMPORTANT!

You need a **Power BI account** to browse or import custom visuals from the AppSource marketplace



MANAGING & VIEWING ROLES

Manage security roles

Create new security roles and use filters to define row-level data restrictions.

Roles

+ New

Select tables

Customer Met...
Measure Table
Price Adjustme...
Product Categ...
Product Lookup
Product Metric...
Product Subca...
Returns Data
Rolling Calendar
Sales Data
Territory Lookup

Filter data

Select all + Add Delete Group Ungroup

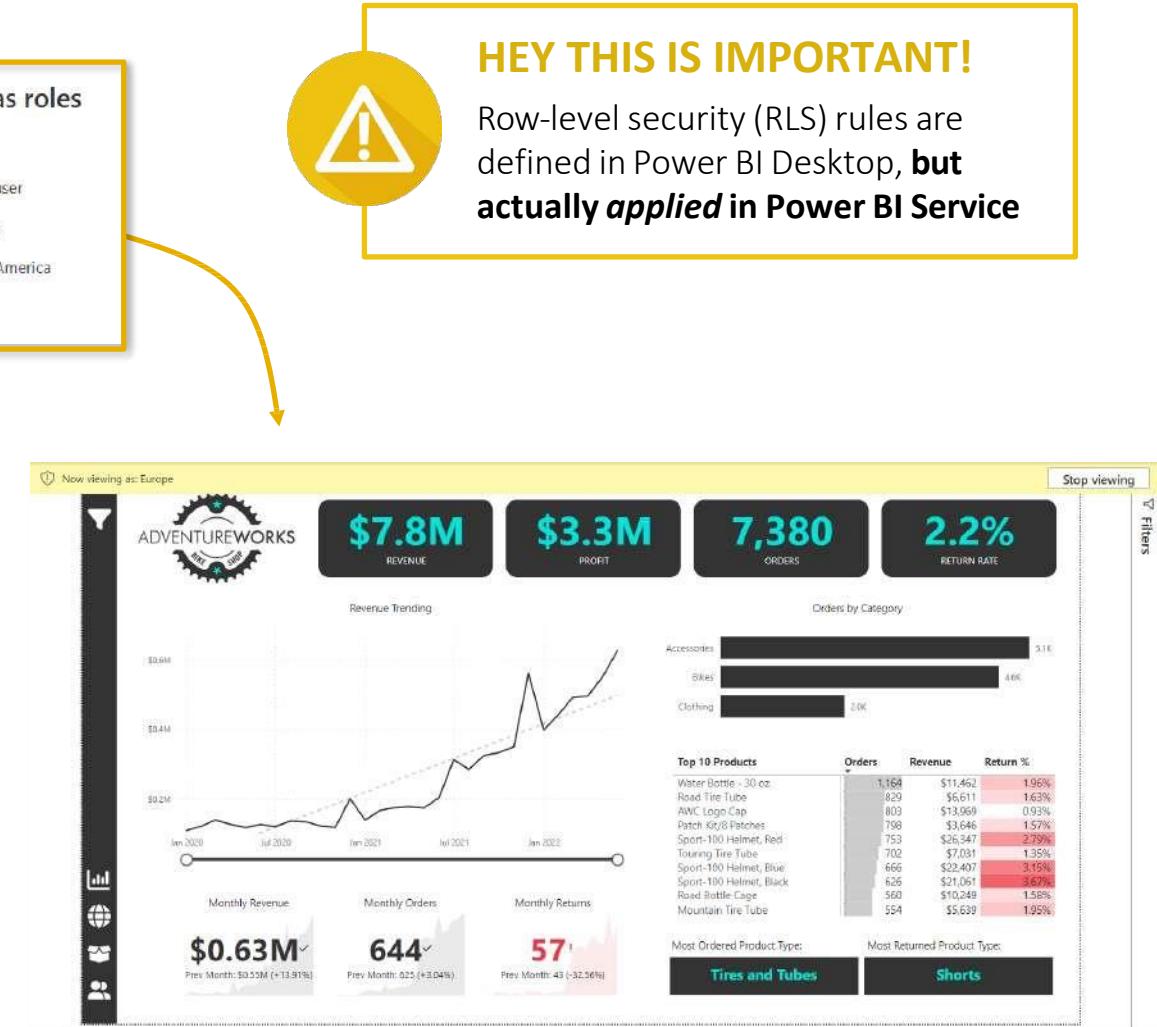
Show data when...
All of these rules are true
Continent Equals North America

Switch to DAX editor



Roles allow you to define row-level security rules, and create filtered views to restrict access for specific audiences

- Here we've created views for territory managers (**Europe**, **N. America**, **Pacific**), which filters records in the model

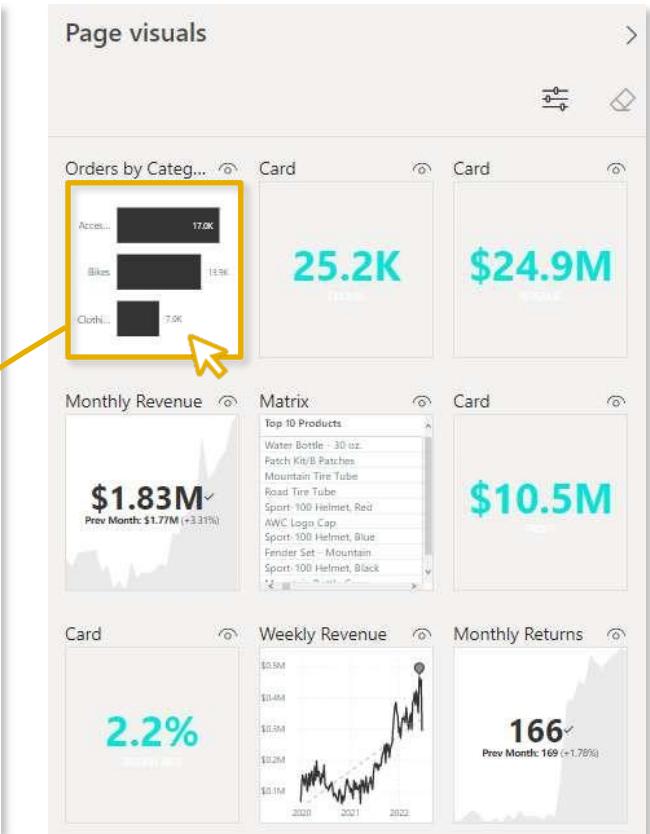
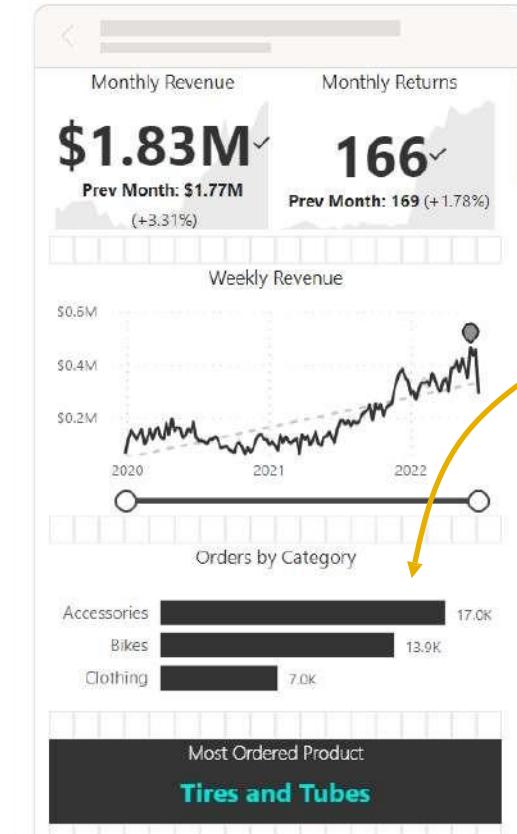
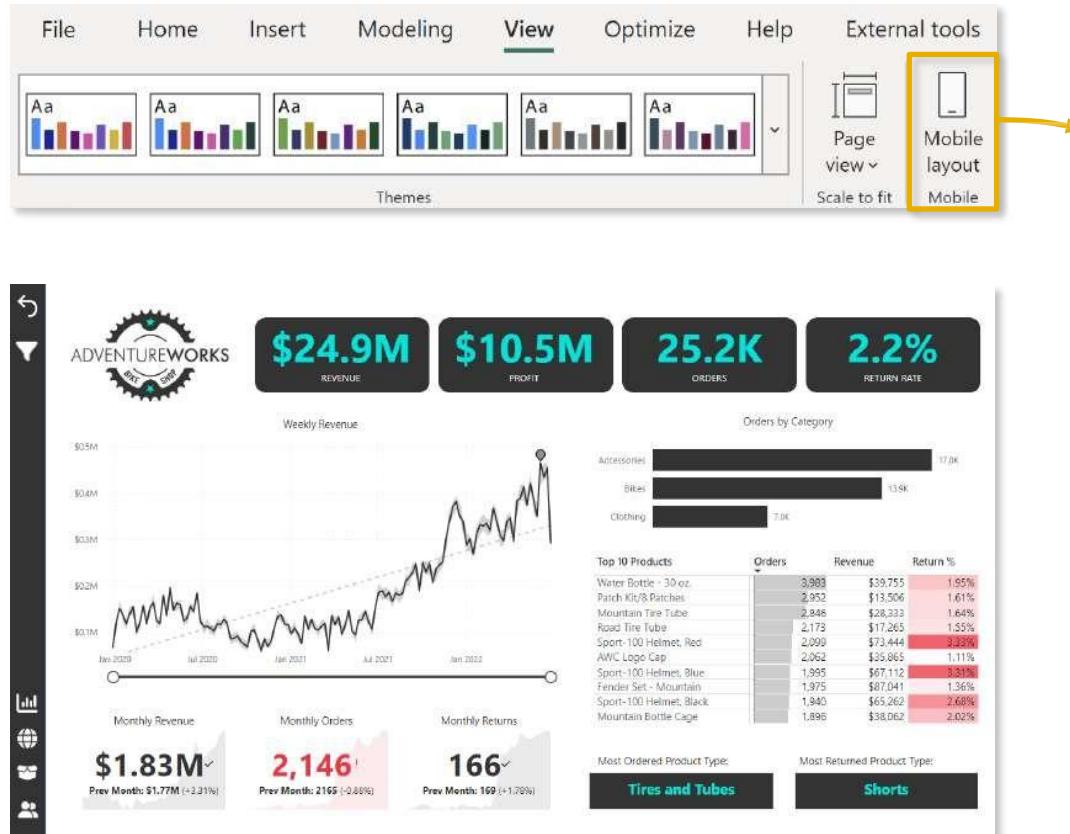




MOBILE LAYOUT

Mobile layout allows you to design mobile-specific versions of report pages by assembling visuals into new layouts

- **NOTE:** This is designed to optimize reports for viewing on the Power BI mobile app (after publishing to Power BI Service)



DATA VISUALIZATION BEST PRACTICES



★ Always ask yourself the three key questions

- *What type of data are you visualizing, what are you communicating, and who is the end user?*

★ Strive for clarity and simplicity above all else

- *"Perfection is achieved not when there's nothing more to add, but when there's nothing left to take away"*

★ Focus on creating clear narratives and intuitive user experiences

- *Use bookmarks, drillthroughs, tooltips and navigation buttons to seamlessly guide users through reports*

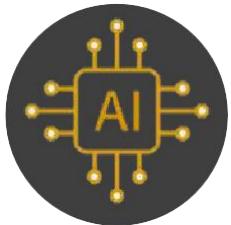
★ Create optimized layouts for mobile viewers

- *Create custom mobile layouts if you plan to publish reports to Power BI Service or use the Power BI app*

ARTIFICIAL INTELLIGENCE



ARTIFICIAL INTELLIGENCE



In this section we'll explore Power BI's artificial intelligence features, including anomaly detection, smart narratives, natural language Q&A, decomposition trees, and more

TOPICS WE'LL COVER:

Anomaly Detection

Smart Narrative

Q&A Visual

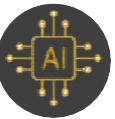
Decomposition Tree

Key Influencers

Top Segments

GOALS FOR THIS SECTION:

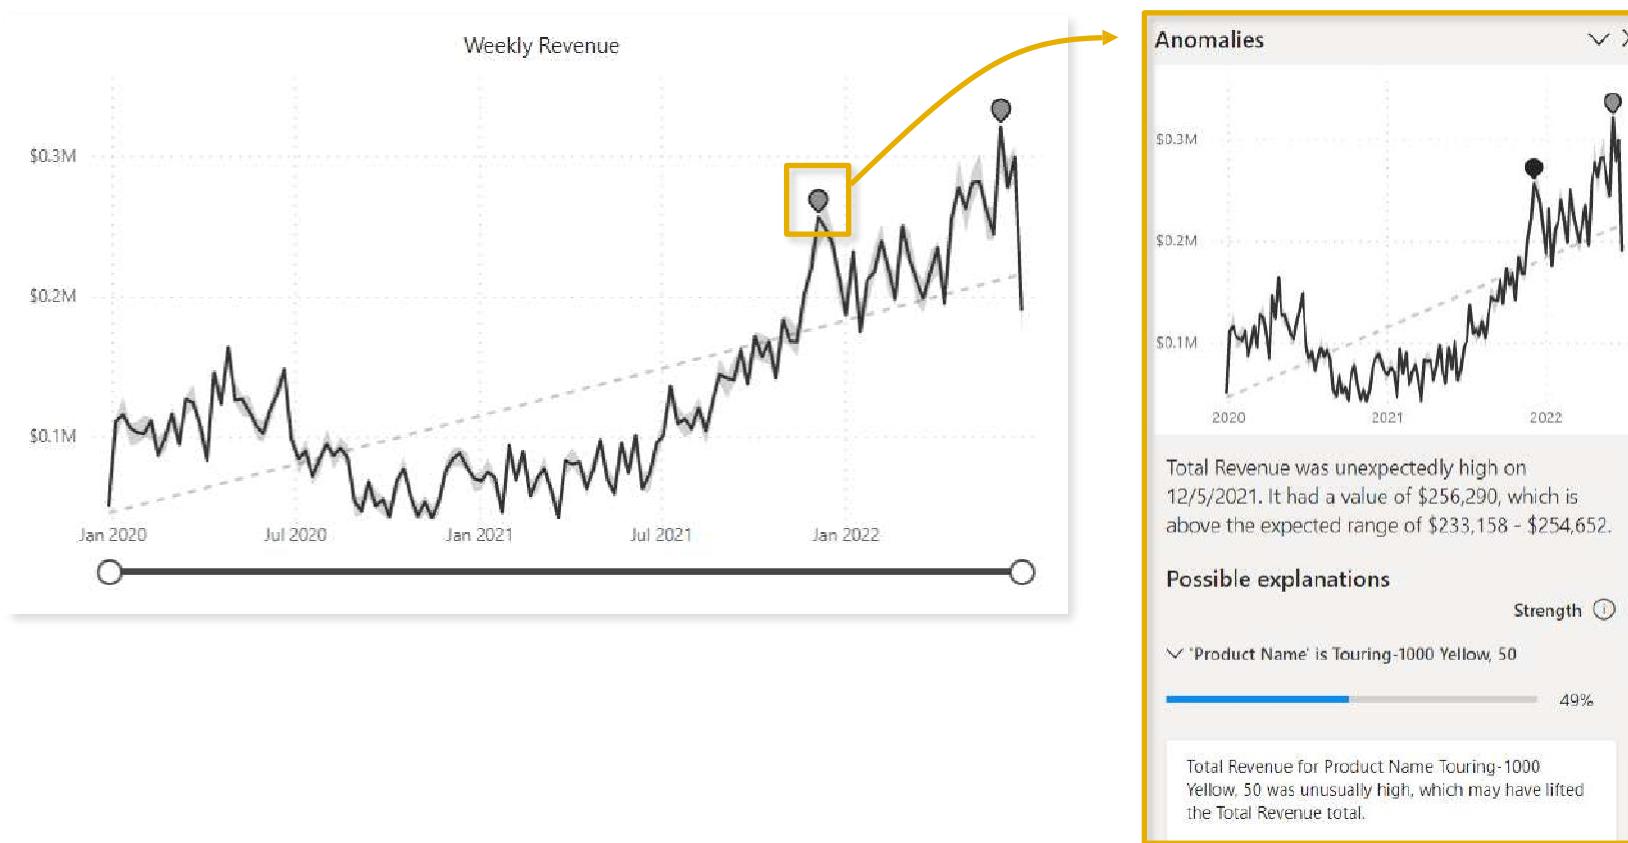
- Explore AI-generated insights using smart narratives and anomaly detection
- Build and train Q&A visuals to allow users to explore Power BI models using natural language queries
- Learn how to use decomposition trees for data exploration and root cause analysis
- Use key influencer visuals to identify the underlying factors that drive specific outcomes for the business



ANOMALY DETECTION

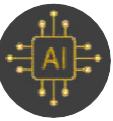
Anomaly detection is used to automatically detect and explain anomalies in time series data

- The anomaly detection feature adds “flags” to existing line charts, which link to AI-generated explanations and summaries



Limitations:

- Only supported for line charts with a time-series field on the X-axis
- Does not support charts with legends, multiple values, or a secondary axis
- Cannot be applied at the same time as forecasts
- Not compatible with drill up/drill down
- Requires at least four data points



SMART NARRATIVES

Smart narratives create customizable, AI-generated text summaries based on report pages or visuals

- Smart narratives react to report filters like any other visual, and can be updated with custom, dynamic values

Selected Product:
Patch Kit/8 Patches

Monthly Orders vs. Target: 265

Monthly Revenue vs. Target: \$1,225

Monthly Profit vs. Target: \$765

Total Profit: ● Total Profit ● Adjusted Profit

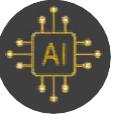
Price Adjustment (%): 0.00

Return Rate

Metric Selection: Total Orders Total Revenue Total Profit Total Returns Return Rate

Report Summary: Total orders for Patch Kit/8 Patches were 265 this month. All metrics trended up between Sunday, June 27, 2021 and Sunday, June 26, 2022, each increasing by 3,200.00%. Return Rate had two high anomalies on Sunday, July 4, 2021 (8.00%) and Sunday, July 25, 2021 (9.52%).

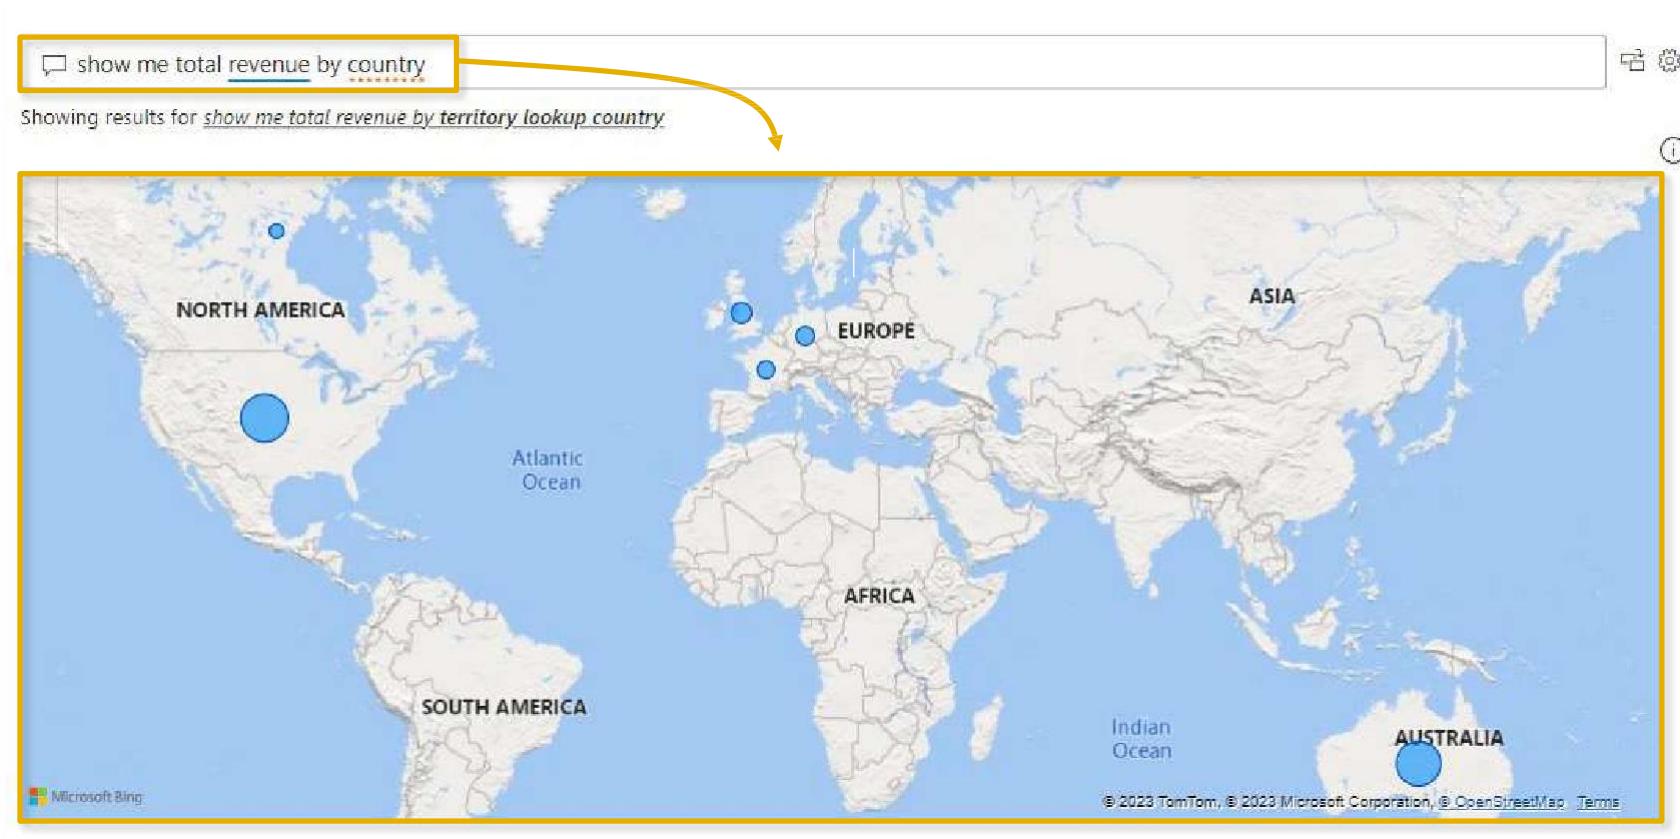
Create a dynamic value that updates with your data
How would you calculate this value
current product
Result
Patch Kit/8 Patches
\$%
Name your value
Product
Save Cancel

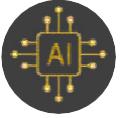


Q&A VISUALS

Q&A visuals allow users to explore and visualize data using intuitive, natural language prompts

- Q&A visuals are only as useful as the data model behind them, and typically require significant “training” to be effective

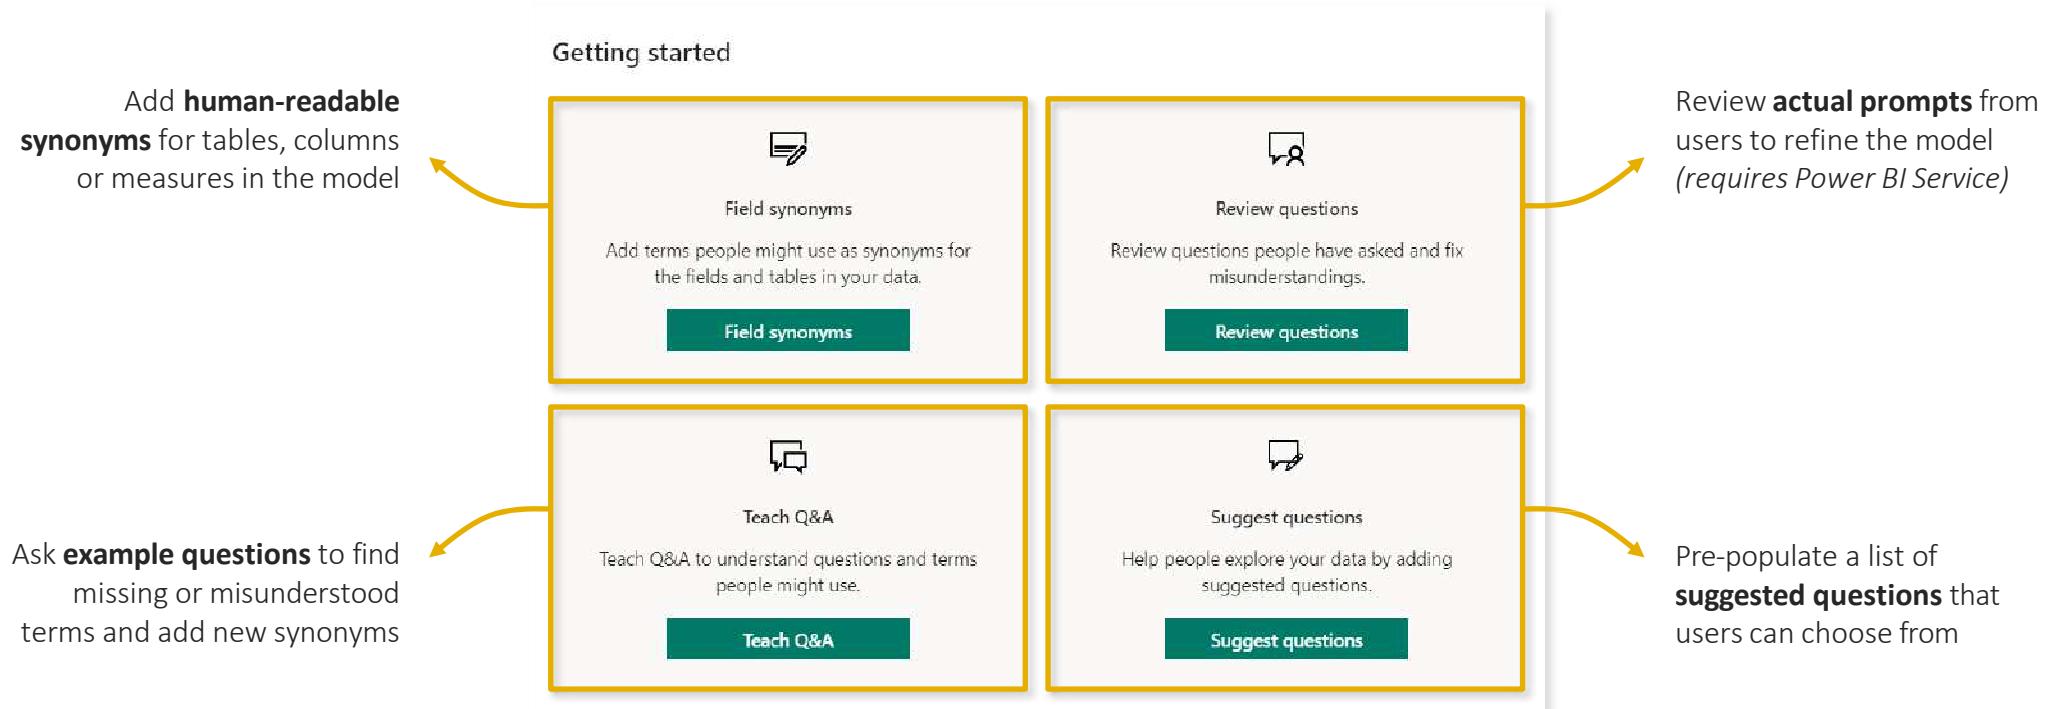




Q&A TRAINING

Q&A visuals allow users to explore and visualize data using intuitive, natural language prompts

- Q&A visuals are only as useful as the data model behind them, and typically require significant “training” to be effective

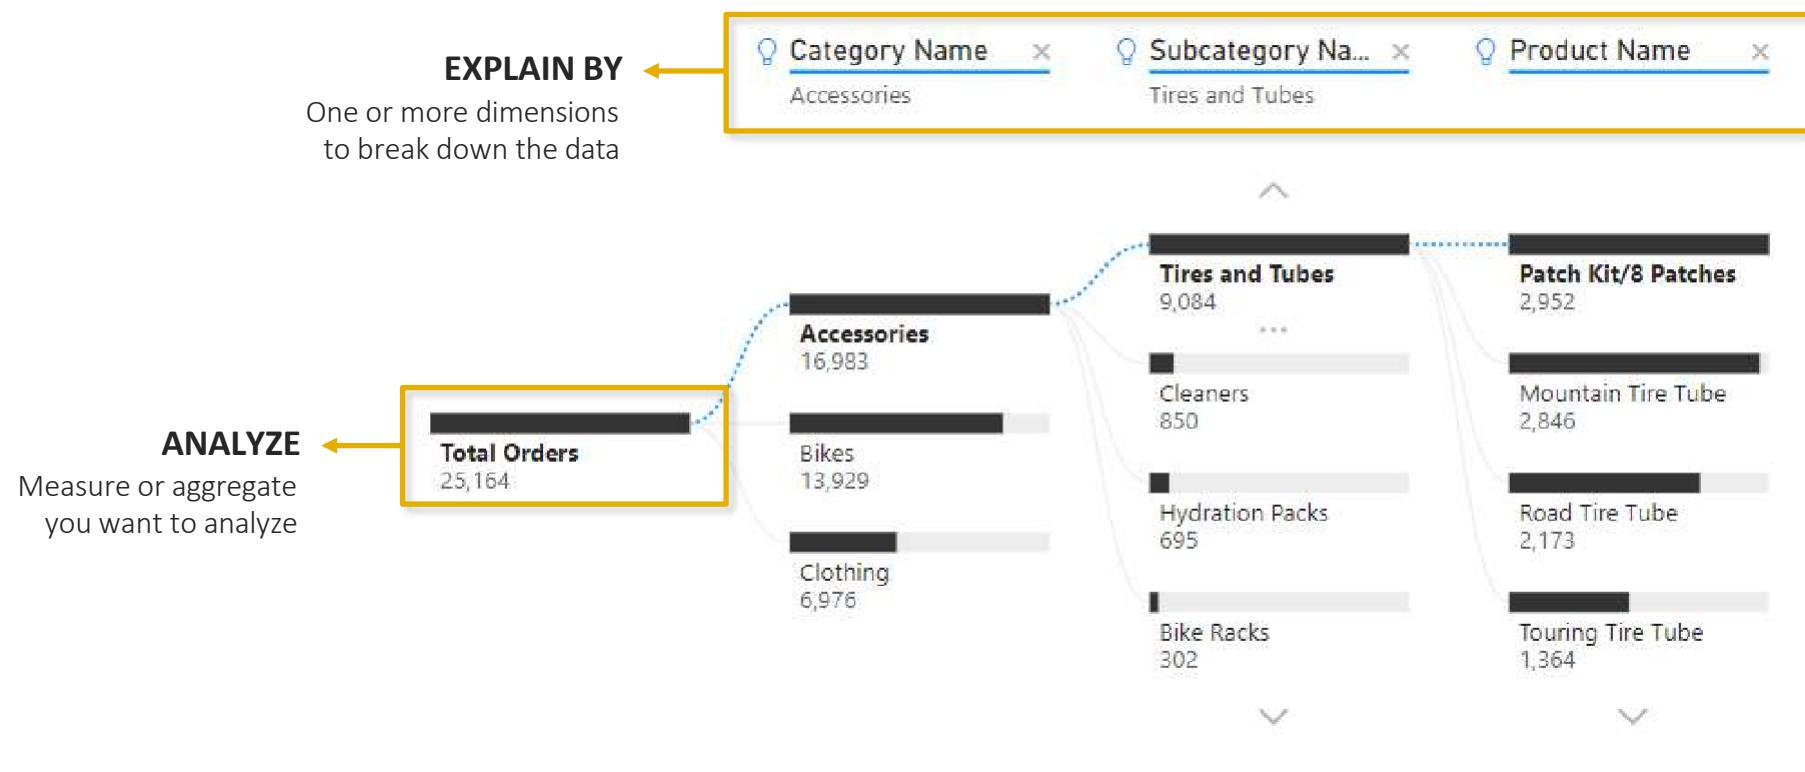




DECOMPOSITION TREES

Decomposition trees allow you to visualize how data is distributed across multiple dimensions

- Decomposition trees can be configured manually for data exploration, or leverage AI to support root cause analysis

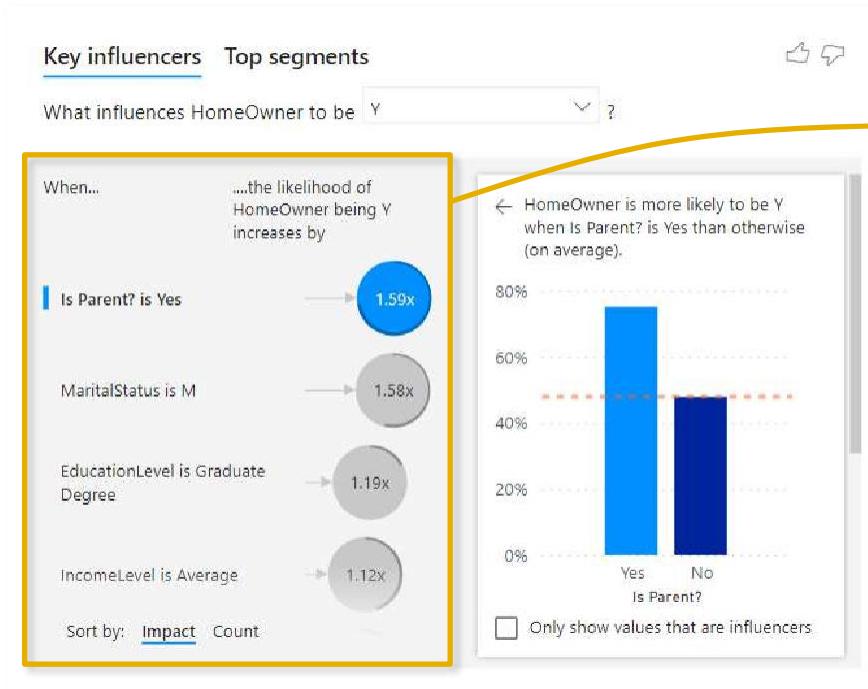




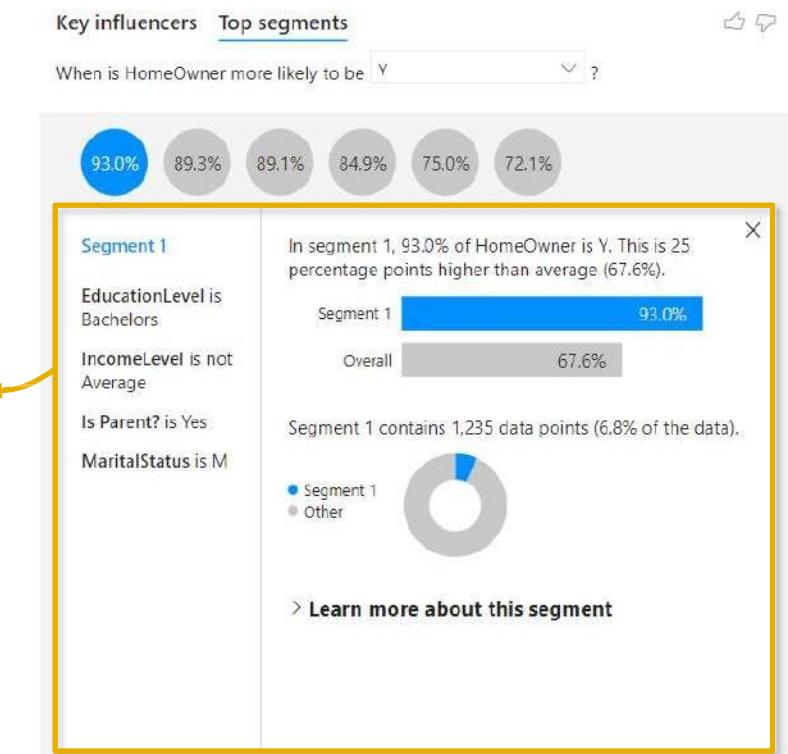
KEY INFLUENCERS

The **key influencer** visual helps you understand the factors that drive specific metrics or outcomes

- This can be used to analyze categorical or continuous outcomes, or identify top segments based on combinations of factors



Here we're identifying factors that are highly correlated with owning a home; for example, **parents are 1.59X more likely to be homeowners**, all else equal



We can also identify customer segments where this outcome is likely; for example, **93% of married customers with children and a Bachelors degree own a home** (vs. 67.6% overall)

OPTIMIZATION TOOLS



PREVIEW: POWER BI OPTIMIZATION



In this section, we'll investigate several native and external tools that can be used to optimize and enhance your Power BI reports

TOPICS WE'LL COVER:

Optimize Ribbon

Pause Visuals

Optimization Presets

Apply all Slicers

Performance Analyzer

External Tools

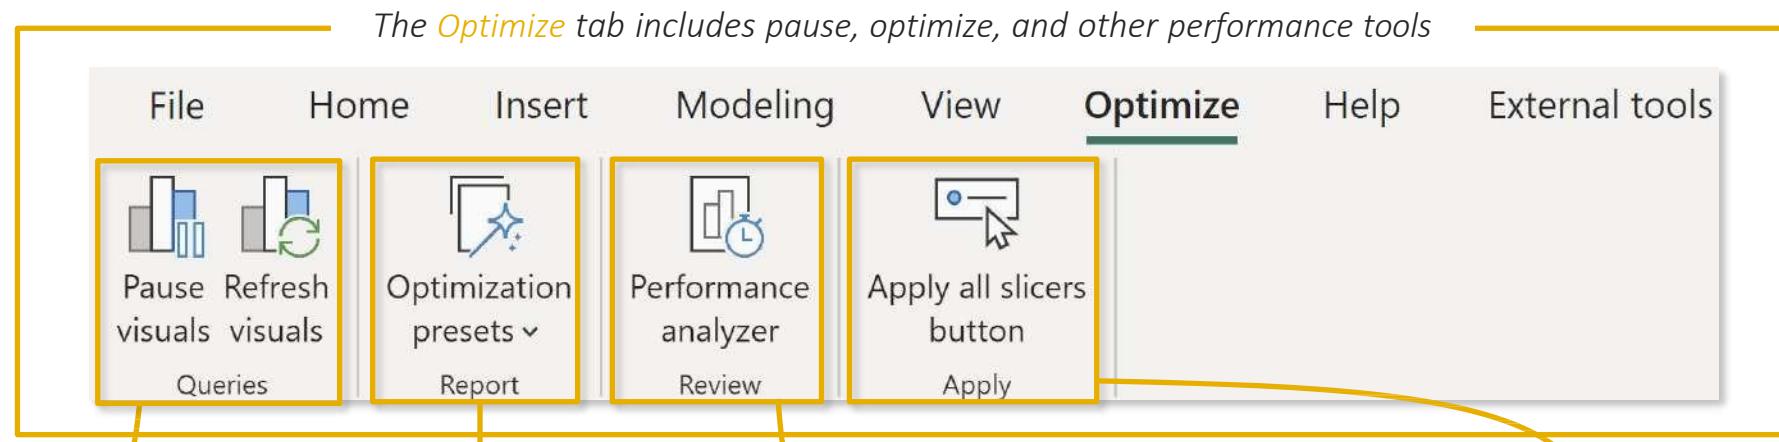
GOALS FOR THIS SECTION:

- Explore the optimize ribbon tools, features, and use cases
- Understand how and when pausing visuals can aid in report development and creation
- Use Performance Analyzer to measure and compare the impact of report elements on speed and performance
- Explore external tools that can aid in report development, learning, and optimization



OPTIMIZE RIBBON

The **Optimize ribbon** helps report authoring by allowing developers to pause data source queries, apply preset settings, and view logs that measure report element performance



Pause or refresh queries to make updates without processing changes

Predefined optimization presets that can be applied based on your reporting scenario

Show **record logs** that measure each elements performance within a report

Apply and **clear multiple slicer selections** on a report page at once



PAUSE VISUALS

Pause visuals stops queries from running and is used when you don't want to immediately apply additions or changes made to a report page or visual

The screenshot shows the Power BI desktop interface with the 'Optimize' tab selected in the ribbon. The 'Paused visuals' button is highlighted with a yellow box. A banner at the top of the main area says 'Visuals are paused. Some edits won't be applied until you refresh or resume visual queries.' Below the banner are several visualizations: a line chart titled 'Weekly Revenue' showing revenue from Jan 2020 to Jul 2023; a bar chart titled 'Orders by Category' showing Accessories, Bikes, and Clothing; and a table titled 'Top 10 Products' listing items like Water Bottle - 30 oz., Patch Kit/8 Patches, Mountain Tire Tube, etc. At the bottom, there are cards for 'Monthly Revenue' (\$1.83M), 'Monthly Orders' (2,146), and 'Monthly Returns' (166).

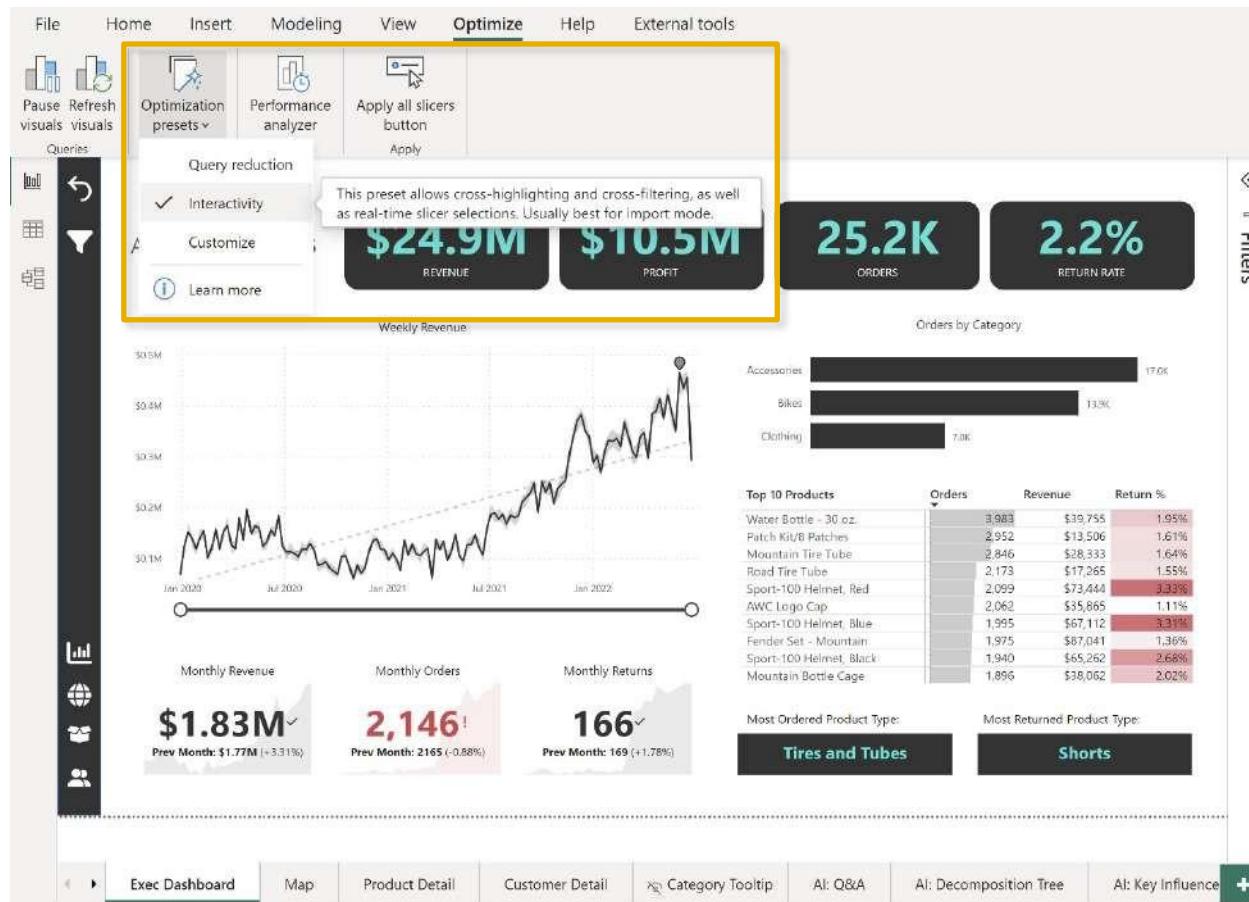
When paused, the report:

- Holds all changes & updates and sets them to a **"visual has pending changes"** state
- Shows a banner with **refresh** & **resume visual queries**
- Adds a **refresh button to individual visuals** allowing you to only refresh that visual
- Allows you to **add, move, and remove columns** and measures without having to wait for visuals to refresh
- **Blocks formatting actions**



OPTIMIZATION PRESETS

Optimization presets allow you to apply different predefined query optimization settings like query reduction, interactive, and custom



Query Reduction

- Is *best for DirectQuery connections* because it follows the best practices for DirectQuery optimization, turns off cross-highlighting, cross-filtering, and adds an Apply button to the filters pane

Interactivity

- Is the default setting and *best used for Import mode* because it allows cross-highlighting, cross-filtering, and real-time changes to slicers and filters

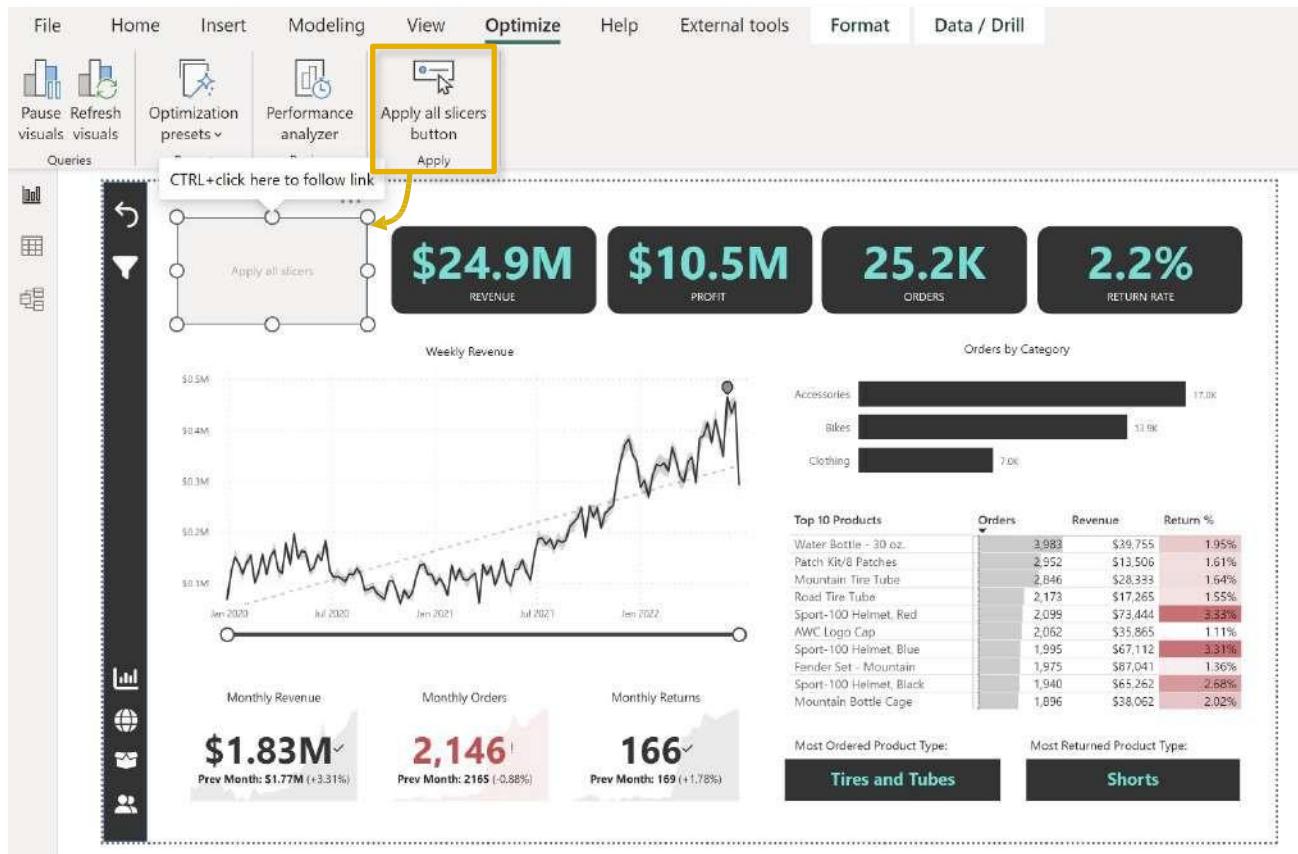
Customize

- Is best when you want to *choose which query reduction features to use*



APPLY ALL SLICERS BUTTON

Adding an **apply all slicers** button to your report page tracks all slicer selections and can be used to either apply or clear all slicers at once



Common scenarios & considerations:

- **Apply multiple slicers** on a report page at once
- **Clear all slicers** on a report page at once
- Apply & clear all slicer buttons **impact all slicers** on the report page (*you can't pick and choose!*)
- You can have as **many of these buttons** as you'd like
- Button can be **added and formatted** just like other buttons in Power BI



PERFORMANCE ANALYZER

Performance Analyzer records user actions (*like Excel's macro recorder*), and tracks the load time (*in milliseconds*) for each step in the process

The screenshot shows a Power BI desktop interface. At the top, the ribbon has tabs: File, Home, Insert, Modeling, View, **Optimize**, Help, and External tools. The 'Optimize' tab is highlighted. Below the ribbon is a dashboard with several visual elements: a circular logo for ADVENTUREWORKS, three large green boxes showing revenue (\$24.9M), profit (\$10.5M), and orders (25.2K), a line chart titled 'Weekly Revenue' showing monthly trends from Jan 2020 to Jun 2022, and three smaller cards for 'Monthly Revenue' (\$1.83M), 'Monthly Orders' (2,146), and 'Monthly Returns' (166). On the right side of the dashboard, there is a 'Filters' pane and a 'Top 10 Products' list. At the bottom, there are navigation buttons for 'Exec Dashboard', 'Map', 'Product Detail', 'Customer Detail', and 'Category Tools', along with a '+' button.

The screenshot shows the 'Performance analyzer' tool window. It has a header with 'Start recording', 'Refresh visuals', 'Stop', 'Clear', and 'Export' buttons. Below is a table with two columns: 'Name' and 'Duration (ms)'. The table lists various components and their execution times:

Name	Duration (ms)
Shape	124
Image	124
Dashboard Icon	123
Map Icon	123
Product Icon	122
Customer Icon	122
Card	329
DAX query	5
Visual display	6
Other	318
Copy query	
Card	356
Card	462
Card	384

At the bottom of the window, there is a note: 'Learn more about optimizing your report's performance on our [support site](#). Find help tuning your report from specialist Power BI partners on [AppSource](#).

DAX Query

- Shows the amount of time it takes for the visual to send the query to the engines, and for the engines to return the result (**Note:** DAX Studio can only help optimize this)

Visual Display

- Shows the amount of time it takes for the visual to populate, or “draw”, on the screen. Includes time to retrieve web-based and geocoded images

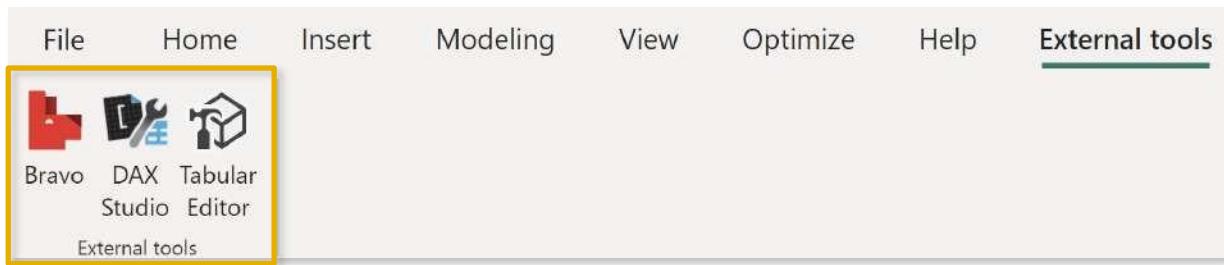
Other

- Shows the amount of time required by the visual to prepare the query, wait for other visuals to complete their queries and perform other processing tasks



EXTERNAL TOOLS

External tools allows quick access to third-party built tools that are *locally installed* on your computer and *registered* with Power BI Desktop



External tools generally fall into one of the following categories:

Semantic Modeling

These tools extend Power BI's functionality for specific data modeling scenarios like DAX optimization, ALM, and metadata translation

- *DAX Studio*
- *ALM Toolkit*
- *Tabular Editor*
- *Bravo*

Data Analysis

Includes tools for connecting a PBI data model to a client application, in read-only mode, to query data and perform analysis tasks

- *Python*
- *Excel*
- *Power BI Report Builder*

Miscellaneous

Some tools are used to make Power BI more useful and accessible but don't connect to the data model

- *PBI.tips tutorials*
- *DAX Guide*
- *PowerBI.tips*